

Statistical Audit Sampling with R

Koen Derks

12/29/22

Contents

Welcome	1
Preface	3
I. The Basics	5
1. Introduction to R	9
1.1. Calculations	9
1.2. Packages	13
1.3. Loading Data	15
2. Background Information	17
2.1. Auditing Standards	17
2.2. Important Concepts	18
2.3. Classical Inference	22
2.4. Bayesian Inference	24
II. Standard Workflow	29
3. Planning	33
3.1. Required Information	33
3.2. The Hypergeometric Likelihood	34
3.3. The Binomial Likelihood	43
3.4. The Poisson Likelihood	52
3.5. Practical Examples	62
4. Selection	79
4.1. Sampling Units	79
4.2. Sampling Methods	80
4.3. Ordering or Randomizing the Population	93
5. Evaluation	95
5.1. Classical Evaluation	95
5.2. Bayesian Evaluation	97
5.3. Evaluation using Data	98

Contents

6. Stratified Evaluation	101
6.1. No pooling	102
6.2. Complete pooling	105
6.3. Partial pooling	108
6.4. Evaluation using data	112
 III. Other Software	 119
7. JASP for Audit	123
7.1. Planning	124
7.2. Selection	124
7.3. Evaluation	125
8. R Packages	127
8.1. MUS	127
8.2. audit	127
8.3. samplingbook	128
 References	 129

Welcome

This book, **Statistical Audit Sampling with R**, is intended as a practical guide for auditors who wish to employ probability theory in their audit sampling activities. While the focus of this book is exclusively on audit sampling, it aims to examine the topic from both the classical (frequentist) perspective and the Bayesian perspective. By examining the subject through these two lenses, the book explains the statistical theory behind commonly used audit sampling procedures and demonstrates how to perform these procedures in accordance with international auditing standards, using the **jfa** R package, in a statistically sound manner.

Preface

This book is intended for auditors that want to obtain the knowledge and skill to utilize statistical audit sampling in their practice using the R programming language. It covers an array of traditional and innovative statistical tools that are available to auditors, explaining their function, the underlying assumptions, and when they are best utilized. In addition, it offers practical guidance on integrating advanced statistical sampling methodology into audit practice and demonstrates its value through real-world case studies. It is our hope that this book will serve as a valuable resource for auditors looking to effectively and efficiently utilize statistical methods in their practice.

The aim of this book is to address the need for a clear and transparent explanation of the use of statistical sampling methodology in audit practice. In my opinion, most guidance about audit sampling (e.g., American Institute of Certified Public Accountants (AICPA) (2016a); American Institute of Certified Public Accountants (AICPA) (2016b)) lacks sufficient detail to allow for full transparency or a deep understanding (one notable exception is Stewart (2012)). Additionally, the implementation of statistical sampling methodology in practice is often even less transparent, as theory and calculations are hidden from auditors in commercial closed-source tools or in Excel sheets from audit guides used internally by audit firms. Thus, while attempting to comprehend the theoretical aspects of statistical audit sampling, auditors may encounter numerous relevant questions that are left unanswered by these tools. This book aims to clarify the statistical methodology utilized in practice, thereby empowering auditors through a comprehensive explanation.

This book discusses two philosophies to statistical audit sampling: the classical (frequentist) philosophy and the Bayesian philosophy. By utilizing these two approaches, the book elucidates the statistical theory that underlies commonly used audit sampling techniques and illustrates how to utilize these techniques in accordance with international auditing standards. Additionally, the book demonstrates the use of Bayesian statistical methods in auditing practice and highlights the practical advantages that these methods can offer for auditors.

The structure of this book is as follows: Chapter 1 introduces the basics of the R programming language. Chapter 2 discusses the fundamental statistical theory relevant to classical and Bayesian audit sampling. Subsequent chapters (3, 4, 5 and 6) provide a more in-depth exploration of the use of these methods for statistical planning, selection, and evaluation of audit samples. These chapters illustrate the practical differences and similarities between the classical and the Bayesian approach through code and examples,

Preface

thus requiring minimal programming knowledge to follow along. Finally, Chapter 7 and 8 discuss other open-source software implementations of audit sampling using R.



Figure 1.: The Bayesian and the classical approach can both lead to an appropriate audit opinion, however, taking a Bayesian approach can often help you get there faster. Image available under a CC-BY-NC 4.0 license.

Part I.

The Basics

In order to effectively use statistical audit sampling techniques, it is important to have a basic understanding of statistical concepts in the context of international auditing standards.

In this part of the book, we go over the basic knowledge required for this book. Chapter 1 introduces the basics of R, a programming language commonly used for statistical analysis. In Chapter 2, we will cover the fundamental statistical concepts related to audit sampling.

1. Introduction to R

i You are reading the work-in-progress first edition of Statistical Auditing with R. This chapter is currently a dumping ground of ideas, and it is incomplete.

R (R Core Team, 2022) is a programming language and software environment for statistical computing and graphics. It is widely used among statisticians and data scientists for data analysis and data visualization. R has a large and active community of users and, as a result, there are many community-made resources available for learning and using R.

In an audit context, R can be used to analyze and visualize large datasets, allowing auditors to identify trends and anomalies in the data. R is particularly useful for performing statistical analysis and testing hypotheses, which can be employed in verifying the accuracy and reliability of financial statements. R can also be used to automate certain audit procedures, reducing the time and effort required to manually review and analyze large amounts of data. Additionally, R allows auditors to easily share their work with others through the use of code and reproducible reports, enabling more efficient and collaborative audit processes. A useful resource that discusses relevant applications of R in the audit can be found in Lin (2021).

This chapter provides a short introduction to working with R but it is not a full programming course in R. If you want to dive deeper into the nooks and crannies of R, the author recommends reading Chang (2022), Wickam & Grolemund (2017), Grolemund (2014) and Wilke (2022). These books are free to read online and discuss various best practices for using R, including examples in code that can be easily reproduced.

1.1. Calculations

One of the basic features of R is its ability to perform calculations. In R, basic calculations work by using the standard arithmetic operators such as + for addition, - for subtraction, * for multiplication, and / for division. For example, if you want to calculate $2 + 3$, you would type in `2 + 3` and R will return the result of 5.

```
2 + 3  
#> [1] 5
```

1. Introduction to R



Figure 1.1.: R can be a powerful tool in a business context. Image available under a CC-BY-NC 4.0 license.

R also allows for more advanced calculations such as exponentiation using the `^` operator, and square roots using the `sqrt()` function. For example, to calculate the square root of 9, you would type in `sqrt(9)` and R will return the result of 3.

```
sqrt(9)
#> [1] 3
```

You can also use parentheses to specify the order of operations in your calculations. For example, if you want to calculate $(2 + 3)^4$, you would type in `(2 + 3)^4` to get the result of 20.

```
(2 + 3)^4
#> [1] 20
```

Overall, basic calculations in R are similar to those in other programming languages and follow the standard order of operations.

1.1.1. Vectors

In R, vectors are one-dimensional arrays of data that can hold numeric, character, or logical values. Vectors can be created using the `c()` function, which stands for concatenate. For example, to create a numeric vector, you can use the following code:

```
x <- c(1, 2, 3, 4, 5)
```

To create a character vector, you can use quotes around the values:

```
y <- c("apple", "banana", "orange")
```

To create a logical vector, you can use the logical values TRUE and FALSE:

```
z <- c(TRUE, FALSE, TRUE, TRUE, FALSE)
```

Vectors can be indexed using square brackets and a numeric value. For example, to access the second element of the vector `x`, you can use the following code:

```
x[2]
#> [1] 2
```

Vectors can also be subsetted using a logical vector. For example, to get all elements of the vector `x` that are greater than 3, you can use the following code:

```
x[x > 3]
#> [1] 4 5
```

Vectors can also be modified using indexing and assignment. For example, to change the third element of the vector `x` to 6, you can use the following code:

```
x[3]
#> [1] 3
x[3] <- 6
x[3]
#> [1] 6
```

R has many built-in functions for performing mathematical operations on vectors. For example, you can use the `mean()` function to calculate the average of a vector of numbers, or we can use the `length()` function to calculate the number of elements in a vector:

```
mean(x)
#> [1] 3.6
length(y)
#> [1] 3
```

Overall, vectors are a useful data structure in R for storing and manipulating data.

1. Introduction to R

1.1.2. Matrices

In R, a matrix is a two-dimensional collection of values that are arranged in rows and columns. You can create a matrix using the `matrix()` function. For example:

```
m <- matrix(1:9, nrow = 3, ncol = 3)
m
#>      [,1] [,2] [,3]
#> [1,]    1    4    7
#> [2,]    2    5    8
#> [3,]    3    6    9
```

This creates a 3x3 matrix with the values 1, 2, 3 in the first column, 4, 5, 6 in the second column, and 7, 8, 9 in the third column.

You can also create a matrix by combining several vectors using the `cbind()` or `rbind()` functions. For example:

```
v1 <- c(1, 2, 3)
v2 <- c(4, 5, 6)
v3 <- c(7, 8, 9)
m <- cbind(v1, v2, v3)
m
#>      v1 v2 v3
#> [1,] 1  4  7
#> [2,] 2  5  8
#> [3,] 3  6  9
```

This creates a matrix with the same values as before, but the columns are created by binding the vectors together.

You can access the elements of a matrix using the square bracket notation. For example, to access the element in the second row and third column of `m`, you would use the following code:

```
m[2, 3]
#> v3
#> 8
```

You can also use the `dim()` function to get the dimensions of a matrix, and the `colnames()` and `rownames()` functions to get the names of the columns and rows, respectively.

There are many other functions and operations available for working with matrices in R, including mathematical operations such as matrix multiplication and inversion.

1.1.3. Data Frames

In R, a data frame is a two-dimensional table of data with rows and columns. Each row represents a single observation or record, and each column represents a particular variable or attribute. Data frames are similar to a spreadsheet in Excel or a table in a database. Each column in a data frame can have a different data type, such as numerical, character, or logical. The data in each row must match the data type of the corresponding column.

To create a data frame in R, you can use the `data.frame()` function and pass in the data you want to include in the data frame as arguments. For example:

```
df <- data.frame(x = c(1, 2, 3), y = c(4, 5, 6))
```

This will create a data frame with two columns, `x` and `y`, and three rows of data. You can access the data in a data frame using indexing and subsetting. For example, to access the first row of the data frame, you can use the following command:

```
df[1, ]
#>   x y
#> 1 1 4
```

To access a specific column, you can use the `$` operator (or the index):

```
df$x
#> [1] 1 2 3
df[, 1]
#> [1] 1 2 3
```

You can also use functions like `head()` and `tail()` to view the first or last few rows of a data frame. Data frames also have several built-in functions that allow you to manipulate and analyze the data. For example, you can use the `summarize()` function to calculate summary statistics for each column, or the `group_by()` function to group the data by a specific variable and apply a function to each group.

1.2. Packages

In addition to these basic features, R has many packages and libraries that extend its capabilities and provide additional functions and tools for data analysis and visualization. A package is a small bundle of code that an R user (or group of users) wrote and uploaded to a central server for everybody to access, download and use. Some popular packages include `dplyr` for data manipulation, `ggplot2` for data visualization, and `caret` for

1. Introduction to R

machine learning. With its versatility and robust community, R is a valuable tool for anyone working with data.

Note

If you want to know more about creating your own R package, Wickam & Brian (2022) provides an excellent first-hand description of the principles and practices of creating R packages.

1.2.1. Installing a Package

In many cases, you will download an R package from the CRAN server. This can be done via the `install.packages()` function by providing the package name in quotes. For example, an important package for data visualization is `ggplot2`. To install this package, you can simply run:

```
install.packages("ggplot2")
```

To illustrate its concepts and ideas, this book heavily draws from the `jfa` package, an R package for statistical auditing, which can also be downloaded from CRAN. Before running the examples in this book, you should install this package by running the following command in R:

```
install.packages("jfa")
```

1.2.2. Loading a Package

Once you have installed a package, you must load it into every R session. To load a package into your R session, call `library()` and provide the name of the package (without quotes) that you want to load. For example, before running the examples in this book, you can load the `jfa` package with:

```
library(jfa)
```

Note

If you want to look at the source code of the `jfa` package, see the package website.

1.2.3. Updating a Package

R packages are updated regularly. To update a package in your R library you should call `update.packages()` and provide the name of the package that you want to update in quotes. For example, each time there is a new release of the `jfa` package, you can update it by running:

```
update.packages("jfa")
```

1.3. Loading Data

When working with data, you will need to load the data file into your R session. How this is done depends on the type of data file that you want to read.

1.3.1. Loading Data from a CSV File

A commonly used data type is a `.csv` file. You can load this type of files via the `read.csv()` function. For example, if the file `example.csv` is in the current working directory, you can load it by running:

```
read.csv("example.csv")
```

1.3.2. Loading Data from an Excel File

Another commonly used data type are Excel files. You can load this type of files via the `read_excel()` function from the `readxl` package. For this to work, you should first install this package using the `install.packages()` command and load it into the R session using a call to `library()`. For example, if the file `example.xlsx` is in the current working directory, and the data you want to load is on the first worksheet, you can load it by running:

```
install.packages("readxl")
library(readxl)
read_excel("example.csv", sheet = 1)
```


2. Background Information

i You are reading the work-in-progress first edition of Statistical Auditing with R. This chapter is currently a dumping ground of ideas, and it is incomplete.

Auditors use audit sampling as a technique to assess a selection of transactions or items within a population in order to form conclusions about the population as a whole. It is a cost-efficient method for examining the accuracy of financial information as it allows auditors to test a representative sample of the population rather than the entire population.

In the field of auditing, sampling becomes necessary when the truth about a population is not readily accessible or discernible through other means. With the advent of modern technology, auditors often have access to an abundance of information about a population, which sometimes enables them to perform integral testing. Nonetheless, there are situations where a sample is still necessary due to the unavailability of complete data. For instance, an auditor may utilize analytical procedures to verify the consistency of payments with payment orders, but then must subsequently confirm the validity of these orders through detailed testing.

In auditing, there are two primary methods of sampling: statistical and nonstatistical. Statistical sampling involves using probability theory to select a sample from the population and draw conclusions about the population based on the sample. Nonstatistical sampling, on the other hand, is based on the auditor's professional judgment and does not use statistical inference to come to a conclusion. This book does not cover nonstatistical sampling.

International auditing standards prescribe the manner in which statistical sampling should be conducted in an audit. The following section discusses what these standards say about statistical sampling.

2.1. Auditing Standards

There are three auditing standards related to staistical sampling in the audit:

- ISA 530: Auditing standard for international firms published by the International Auditing and Assurance Standards Board (IAASB).

2. Background Information

- AU-C 530: Auditing standard for private firms published by the American Institute of Certified Public Accountants (AICPA).
- AS 2315: Auditing standard for public firms published by the Public Company Accounting Oversight Board (PCAOB).

All three standards present a similar explanation of statistical sampling. For instance, ISA 530 (International Auditing and Assurance Standards Board (IAASB), 2018) defines statistical audit sampling as a method that at minimum exhibits the following two characteristics:

- Random selection of sample items,
- The use of an appropriate statistical technique to evaluate sample results, including measurement of sampling risk

According to auditing standards, any sampling approach that lacks these two characteristics is considered nonstatistical sampling.

In order to effectively utilize statistical sampling during an audit, it is necessary to tailor the approach to the specific circumstances of the audit. This may involve considering factors such as the size and complexity of the population, the materiality of the items being tested, and the level of inherent risk in the audit area. It is also essential for the auditor to document the sampling process in order to demonstrate compliance with auditing standards. The following sections will delve further into these crucial concepts in the context of statistical audit sampling.

2.2. Important Concepts

This section aims to delve into several theoretical concepts that are integral to statistical audit sampling.

2.2.1. Materiality

In an audit, materiality is the maximum amount of misstatement that can be present in the financial statements of the auditee before the auditor concludes that the financial statements are materially misstated, meaning that they contain misstatements that would influence the decisions of stakeholders relying on those statements.

The term performance materiality refers to the maximum amount of misstatement that can be present in a given population that is part of the financial statements before the auditor concludes that the population is materially misstated. Performance materiality is used by auditors to determine the appropriate level of testing to be performed on a population.

i Note

The performance materiality is usually defined to be lower than the materiality because an individual population that is subject to audit sampling is often only a (small) part of the financial statements.

For example, consider an audit of a company's financial statements for the year ended December 31, 2021. The auditor determines that the company's accounts receivable balance is a large part of the financial statements and decides to test a sample of the accounts receivable transactions to assess the accuracy of the balance. The auditor calculates the performance materiality for the accounts receivable balance by considering the materiality for the financial statements as a whole. If the auditor finds misstatements in the sample such that their estimate of the misstatement exceeds the performance materiality, the auditor would need to express an unqualified opinion on the population or would need to perform additional testing on the population. If the auditor finds misstatements in the sample such that their estimate of the misstatement does not exceed the performance materiality, the auditor would express a positive opinion on the financial statements.

2.2.2. Audit risk

After completing an audit and making any necessary corrections, an auditor will issue a written report stating whether the financial statements are accurate and free of material misstatement. The potential for this opinion to be incorrect is known as audit risk, and it is the auditor's job to minimize this risk as much as possible.

For example, during an audit of a company's financial statements, the auditor may carefully review documentation, perform tests of details via audit sampling, and speak with management in order to reduce the audit risk and provide a reliable opinion on the accuracy of the financial statements as a whole.

2.2.3. Population

In statistical inference, the term population refers to the entire group of individuals or items that have some common characteristic or interest, and about which we want to gather data or make inferences. A population can be as large as all the people in a country or, as is more sensible in auditing, as small as a group of employees in a specific department of a company.

For example, consider an audit of a company's payroll records. The population in this case would be all the employees of the company, and the goal of the audit would be to gather data on their salaries, benefits, and other payroll-related information. The audit team would collect this data from a representative group of employees (i.e., a

2. Background Information

sample) of the population and use statistical methods to draw conclusions about the entire population.

2.2.4. Sampling risk

There is a possibility that the results of an audit based on a sample may differ from the results if the entire population were examined using the same procedures. This is known as sampling risk. Sampling risk can result in two types of incorrect conclusions:

1. The first type is when, in a test of controls, the controls are perceived to be more effective than they actually are, or in a test of details, a material misstatement is believed to not exist when it actually does. This type of erroneous conclusion is particularly concerning for auditors because it can compromise the effectiveness of the audit and may lead to an inappropriate audit opinion.
2. The second type of incorrect conclusion is when, in a test of controls, the controls are perceived to be less effective than they actually are, or in a test of details, a material misstatement is believed to exist when it actually does not. This type of erroneous conclusion impacts the efficiency of the audit as it may require additional work to determine that the initial conclusions were incorrect.

Many audits are performed according to the *audit risk model (ARM)*, which determines that the uncertainty about the auditor's statement as a whole is a factor of three terms: the inherent risk, the control risk, and the detection risk (i.e., the sampling risk). Inherent risk is the risk posed by a misstatement in the auditee's financial statements that could be material, before consideration of any related control systems (e.g., computer systems). Control risk is the risk that a material misstatement is not prevented or detected by the auditee's internal control systems. Detection risk is the risk that the auditor will fail to find material misstatements that exist in the auditee's financial statements. The *ARM* is practically useful because, for a given level of audit risk, the tolerable detection risk bears an inverse relation to the other two risks.

$$\text{Audit risk} = \text{Inherent risk} \times \text{Control risk} \times \text{Detection risk} \quad (2.1)$$

Usually the auditor judges inherent risk and control risk on a three-point scale consisting of low, medium, and high. Different audit firms handle different standard percentages for these categories. Given an assessment of the inherent risk and the control risk, the detection risk can be calculated as:

$$\text{Detection risk} = \frac{\text{Audit risk}}{\text{Inherent risk} \times \text{Control risk}} \quad (2.2)$$

i Note

The ARM is commonly used in practice, but is not a proper model of audit risk. For example, it is not possible to set one of the risks to zero, as that would result in an infinite detection risk (e.g., $\frac{0.05}{0 \times 1} = \infty$).

Let's consider an example. Suppose that, in their audit guide, an audit firm associates the following percentages with the categories high, medium and low:

- High: 100%
- Medium: 60%
- Low: 50%

If an auditor is working with an audit risk of 5%, and judges inherent and control risk to both be medium, the sampling risk can be calculated as:

$$\frac{0.05}{0.6 \times 0.6} = 0.139 \quad (2.3)$$

2.2.5. Sample Size

The sample size is an important consideration in the context of audit sampling, as it determines the number of items that will be selected for testing during the audit process. This factor has an impact on both effectiveness and efficiency. In general, a larger sample size can provide a higher level of assurance, but it requires more audit effort to obtain and inspect. On the other hand, a smaller sample size offers a lower level of assurance, but it is less costly.

2.2.6. Notation

The table below summarizes the notation used in this book (middle column) and in the **jfa** R package (right column).

Meaning	Symbol	jfa
Probability of misstatement	θ	
Performance materiality	θ_{max}	materiality
Expected deviation rate	θ_{exp}	expected
Type-I error probability	α	1 - conf.level
Type-II error probability	β	
Population size	N	N.units
Population misstatements	K	
Sample size	n	n

2. Background Information

Meaning	Symbol	jfa
Observed misstatements	k	x

2.3. Classical Inference

Frequentist statistics, also known as classical statistics, is a statistical framework that is based on the concept of probability as a long-term frequency of events. This approach assumes that statistical models are purely objective and that data is generated by a well-defined process, which can be described by a set of probabilistic assumptions. The philosophy behind frequentist statistics is that statistical estimates should be based on the frequency of events in a population, rather than on subjective or personal beliefs. This approach is particularly useful for making predictions or making decisions based on data, as it allows for the calculation of confidence intervals and statistical tests, which provide a measure of the reliability of the estimates. Overall, frequentist statistics is a rigorous and reliable approach that is widely used in the scientific community for making informed decisions based on data.

2.3.1. Parameter Estimation

The philosophy behind frequentist parameter estimation is based on the idea that statistical parameters are fixed, but unknown, quantities that can be estimated through the process of repeated sampling. This approach assumes that the sample data represent a random sample from a larger population, and that the sample statistics (i.e., the sample proportion of misstatements) can be used to estimate the corresponding population parameters (i.e., the population misstatement). The key principle of frequentist estimation is that the estimated parameter values should be unbiased and have a certain level of precision, which can be quantified through statistical measures such as confidence bounds or intervals.

2.3.2. Hypothesis Testing

Frequentist hypothesis testing is a statistical method that involves evaluating the probability of obtaining a certain sample outcome or more extreme, given a certain assumption or hypothesis. This probability, known as the p value, is used to determine the likelihood of the hypothesis being true.

For example, in a typical audit sampling hypothesis test using the binomial distribution, we may be interested in testing the hypothesis that the misstatement is higher or lower than the performance materiality. We would inspect a sample and calculate the p value based on the observed frequency of misstatements versus the expected frequency under

the assumption of material misstatement. If the p value is below the sampling risk α , we reject the hypothesis that the population is materially misstated and conclude that it is not materially misstated.

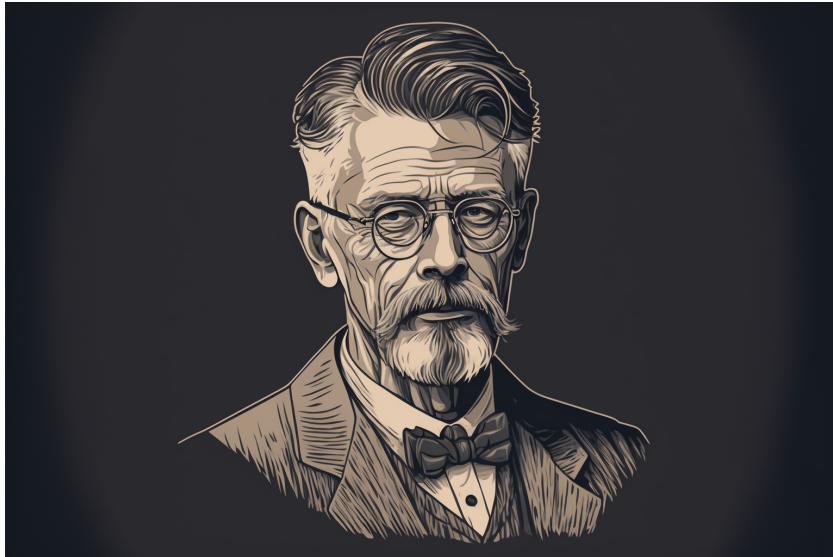


Figure 2.1.: Innovating statistical methods and shaping the field of genetics, Sir Ronald Fisher was a pioneer in the world of science. Image available under a CC-BY-NC 4.0 license.

2.3.3. Example

As an example, the `binom.test()` function in R can be used to test if a population contains less than 3 percent misstatements. Suppose an auditor obtained a sample of $n = 100$ items containing $k = 0$ misstatements. To use the `binom.test()` function, the auditor must input the number of items in the sample `n = 100`, the number of misstatements in the sample `x = 0`, and the hypothesized proportion of misstatement in the population (i.e., the performance materiality) `p = 0.03`. The sampling risk is set to 5%, which the auditor can provide to the function with `conf.level = 1 - 0.05`. Finally, the auditor can specify the alternative hypothesis as `alternative = "less"` to test if the proportion of misstatements in the sample is less than the hypothesized proportion.

```
binom.test(x = 0, n = 100, p = 0.03, alternative = "less",
            conf.level = 0.95)
#>
#>  Exact binomial test
#>
```

2. Background Information

```
#> data: 0 and 100
#> number of successes = 0, number of trials = 100, p-value = 0.04755
#> alternative hypothesis: true probability of success is less than 0.03
#> 95 percent confidence interval:
#> 0.00000000 0.02951305
#> sample estimates:
#> probability of success
#> 0
```

The most likely misstatement in the population is displayed under `sample estimates` and is 0%. The 95% upper confidence bound for the estimate of the population misstatement is displayed under `95 percent confidence interval` and is 2.95%. The p value is shown to be 0.04755. Since the p value is lower than the specified sampling risk α , the auditor can reject the hypothesis of material misstatement.

2.4. Bayesian Inference

Bayesian inference is based on the idea that the parameters in a statistical model are not fixed but uncertain. In this approach, the parameter is considered to be a random variable with a certain distribution, and the goal is to use the data and any prior knowledge about the parameter to update our belief about its value. This is typically done using Bayes' theorem, which states that the posterior probability (i.e., the updated belief about the parameter after seeing the data) is equal to the prior probability (i.e., the belief about the parameter before seeing the data) times the likelihood (i.e., the probability of the data given the parameter).

$$\text{Posterior} \propto \text{Likelihood} \times \text{Prior} \quad (2.4)$$

Bayesian statistics is a more nuanced approach that allows for more efficiency in statistical audit sampling, but it requires the specification of prior distributions that can be difficult to quantify. That is because, especially in an audit, all information that is incorporated into the statistical analysis should be based on audit evidence and should be properly justified.

2.4.1. Parameter Estimation

One major difference between classical and Bayesian statistics is the way they handle uncertainty. In classical statistics, uncertainty is represented by the standard error of an estimate, which is a measure of the precision of an estimate. In Bayesian statistics, uncertainty is represented by the posterior distribution, which is a distribution of the possible values of the population parameter given the sample data and our prior beliefs.

Bayesian inferences uses Bayes' theorem to update the prior beliefs about the population parameter with the new information from the sample data. Bayes' theorem is given by the following formula:

$$p(\theta|y) = \frac{p(y|\theta)p(\theta)}{p(y)} \quad (2.5)$$

where $p(\theta|y)$ is the posterior probability of the population parameter θ given the sample data y , $p(y|\theta)$ is the likelihood of the sample data given θ , $p(\theta)$ is the prior probability of θ , and $p(y)$ is the total probability of the sample data occurring. Because with a fixed sample $p(y)$ is a constant, Bayes' theorem is often given as follows:

$$p(\theta|y) \propto p(y|\theta) \times p(\theta) \quad (2.6)$$

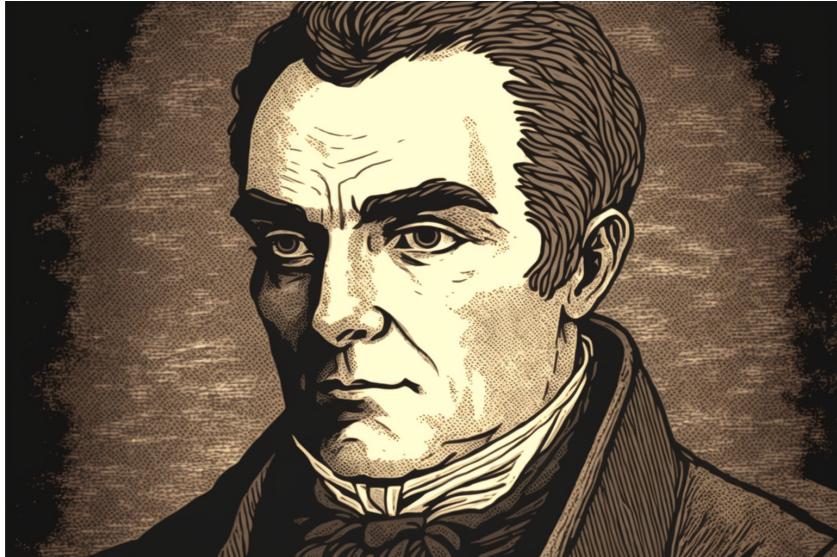


Figure 2.2.: Thomas Bayes revolutionized the world of statistics and probability with his groundbreaking work on Bayes' Theorem. His contributions continue to shape the way we understand and analyze data today. Image available under a CC-BY-NC 4.0 license.

2.4.2. Hypothesis Testing

The Bayes factor is a measure used in Bayesian inference to compare the relative strength of evidence between two competing hypotheses. The Bayes factor is calculated by comparing the probability of the observed data given each of the two competing hypotheses. This probability is known as the likelihood of the data. The Bayes factor is then the ratio of the likelihood of the data under one hypothesis to the likelihood of the data under

2. Background Information

the other hypothesis. The Bayes factor can be used in the context of an audit, where the auditor is trying to determine the likelihood that a particular financial statement is represented fairly or not.

For example, an auditor might be evaluating the fairness of a company's financial statements for the year. They have two hypotheses: the first is that the statements are accurate, and the second is that the statements are not accurate. The auditor gathers data from a statistical audit sample and uses this data to calculate the Bayes factor.

The Bayes factor is calculated by taking the ratio of the probability of the first hypothesis (that the statements are accurate) given the observed data, to the probability of the second hypothesis (that the statements are not accurate) given the observed data. The higher the Bayes factor, the more likely it is that the first hypothesis is true.

The Bayes factor can be used to assess the strength of evidence for one hypothesis over the other and to determine which hypothesis is more likely to be true given the observed data. It is often used in scientific research to help evaluate the validity of different hypotheses and to make informed decisions based on the available evidence. For auditors, the Bayes factor can be a useful tool to determine the likelihood of different hypotheses being true based on the data they have collected, and it can help them make informed decisions about the fairness of the financial statements.

For example, if the Bayes factor is 5, this means that the probability of the statements being accurate given the observed data is 5 times higher than the probability of them being not accurate. In this case, the auditor would be more likely to conclude that the financial statements are accurate.

The **jfa** package can help us to calculate Bayes factors for and against the hypothesis that the financial statements are represented fairly.

2.4.3. Example

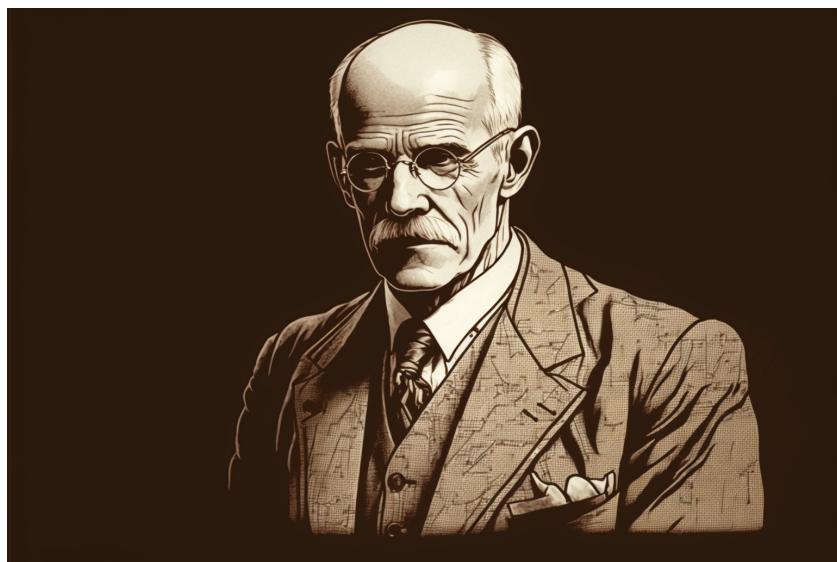


Figure 2.3.: Sir Harold Jeffreys innovated statistical hypothesis testing with his Bayes factor approach, helping us make better decision in the face of uncertainty. Image available under a CC-BY-NC 4.0 license.

Part II.

Standard Workflow

The audit sampling workflow is an essential component of statistical audit sampling with R. It involves a series of steps that help auditors select a representative sample from a population of transactions or items, and then use statistical analysis to draw conclusions about the entire population.

In this part of the book, we will delve into the details of the standard audit sampling workflow and how it can be implemented using the powerful tools and techniques available in R.



Figure 2.4.: The standard audit sampling workflow consists of four stages: planning, selection, execution and evaluation.

Stage 1: Planning

Proper planning of a sample can lead to an increase in audit efficiency. Auditing a small sample does not demand a great deal of audit effort, but it also does not provide a high level of assurance. In contrast, auditing a larger sample provides a higher level of assurance, but it also requires more audit effort. Therefore, it is often advantageous to determine the sample size beforehand in order to find the optimal balance between these two considerations. In Chapter 3, we will examine the various factors that should be taken into account when planning an audit sample and how to use R to aid in the sample planning process.



Figure 2.5.: The planning stage is the first stage in the audit sampling workflow.

Stage 2: Selection

Selecting a sample is the second step in the audit sampling workflow. In Chapter 4, we will explore the various approaches and techniques that can be used to select an audit sample using R. We will also provide examples to help you understand how to apply these techniques in real-world audit scenarios.



Figure 2.6.: The selection stage is the second stage in the audit sampling workflow.

Stage 3: Execution

The third stage of the audit sampling process is known as the execution stage. During this stage, the auditor examines a selection of items from the population and assesses their accuracy or fairness. This stage involves manual labor and does not involve statistical inference.

Stage 4: Evaluation

Evaluating an audit sample is a crucial stage in the audit sampling workflow. In Chapter 5, we will examine the methods and techniques utilized to evaluate an audit sample from a classical and Bayesian point of view. We will delve into the diverse statistical tests and metrics that can be applied to estimate the population misstatement and quantify the strength of evidence for or against a hypothesis provided by the sample data. By the end of this chapter, you will possess a thorough comprehension of how to effectively and efficiently evaluate an audit sample using R and make informed decisions based on the results. Chapter 6 discusses how these methods can be generalized and applied when evaluating stratified samples.



Figure 2.7.: The evaluation stage is the final stage in the audit sampling workflow.

3. Planning

i You are reading the work-in-progress first edition of Statistical Auditing with R. This chapter is currently a dumping ground of ideas, and it is incomplete.

One of the key considerations in audit sampling is determining the appropriate sample size to achieve a desired level of assurance or precision. This chapter will focus on three standard likelihoods commonly used in audit sampling: the hypergeometric likelihood, the binomial likelihood, and the Poisson likelihood.

Note that, when you are using the Bayesian approach to audit sampling, it is not required to plan a specific sample size in advance (Derks et al., 2022). That is because in Bayesian inference, the posterior distribution after seeing each item is used as the prior distribution for the next item. That means that you can simply start sampling and monitor the evidence in the data over time. However, to get an idea of how many samples are required to achieve a certain amount of assurance, planning a sample using a Bayesian approach can still be good practice.

3.1. Required Information

First, planning a minimum sample requires knowledge of the conditions that lead to acceptance or rejection of the population (i.e., the sampling objectives). Typically, sampling objectives can be classified into one or both of the following:

- **Hypothesis testing:** The goal of the sample is to obtain evidence for or against the claim that the misstatement in the population is lower than a given value (i.e., the performance materiality).
- **Estimation:** The goal of the sample is to obtain an accurate estimate of the misstatement in the population (with a minimum precision).

Second, it is advised to specify the expected (or tolerable) misstatements in the sample. The expected misstatements are the misstatements that you allow in the sample, while still retaining the desired amount of assurance about the population. It is strongly recommended to set the value for the expected misstatements in the sample conservatively to minimize the chance of the observed misstatements in the sample exceeding the expected misstatements, which would imply that insufficient work has been done in the end.

3. Planning

Finally, next to determining the sampling objective(s) and the expected misstatements, it is important to determine the statistical distribution linking the sample outcomes to the population misstatement. This distribution is called the likelihood (i.e., `poisson`, `binomial`, or `hypergeometric`). All three aforementioned likelihoods are commonly used in an audit sampling context, however, `poisson` is the default likelihood in `jfa` because it is the most conservative of the three. In the subsections below, we elaborate on the three standard likelihoods for audit sampling and demonstrate how they can be used to obtain a minimum sample size.

In `jfa`, determining an appropriate sample size is achieved via the `planning()` function.



Figure 3.1.: When the sampling objectives, the expected misstatements and the likelihood are set, there exists only one sample size that minimizes the amount of audit effort. Image available under a CC-BY-NC 4.0 license.

3.2. The Hypergeometric Likelihood

The hypergeometric distribution is a discrete probability distribution that is commonly used to model the number of events occurring in a fixed number of trials when the population size is known. It assumes that samples are drawn from the population without replacement, and is therefore the likelihood that most closely resembles the audit practice. For our purpose, we can use the hypergeometric distribution as a likelihood to model the number of misstatements that are expected to be found in the sample.

The probability mass function (PMF) of the hypergeometric distribution is given by:

3.2. The Hypergeometric Likelihood

$$p(X = k) = \frac{\binom{K}{k} \binom{N-K}{n-k}}{\binom{N}{n}}, \quad (3.1)$$

where k is the number of misstatements in the sample, n is the sample size, N is the population size and K is the total number of misstatements assumed in the population. The assumed misstatements K is a linear extrapolation of the assumed misstatement rate in the population θ_{max} to the total population:

$$K = \theta_{max} N. \quad (3.2)$$

Let's consider how to use the hypergeometric likelihood to calculate the minimum sample size needed to achieve a certain level of assurance.

3.2.1. Classical planning

In classical planning using the hypergeometric likelihood, the following statistical model is specified:

$$k \sim \text{Hypergeometric}(n, N, K) \quad (3.3)$$

Given the performance materiality θ_{max} , we can solve for the minimum sample size n needed to achieve this assurance level. This sample size is dependent on the number of misstatements that the auditor expects, or tolerates, in the sample.

3.2.1.1. No Expected Misstatements

If the auditor does not expect any misstatements in the sample, they can set $k = 0$, which consequently determines how the sample size can be calculated. For example, if we want to achieve an assurance level of 95% ($\alpha = 0.05$) for a performance materiality of $\theta_{max} = 0.03$ in a population of $N = 1000$ items, then the required sample size under the assumption of no expected misstatements in the sample is $n = 94$.

```
plan <- planning(materiality = 0.03, expected = 0, conf.level = 0.95,
                   likelihood = "hypergeometric", N.units = 1000)
plan
#>
##> Classical Audit Sample Planning
#>
##> minimum sample size = 94
##> sample size obtained in 95 iterations via method 'hypergeometric'
```

3. Planning

💡 Additional arguments

The `planning()` function has two additional arguments that are not shown in the call above: `by` and `max`. The argument `by` sets the increment between possible sample sizes for consideration. For example, `by = 5` considers only samples of size 10, 20, 30, etc:

```
planning(materiality = 0.03, expected = 0, conf.level = 0.95,
          likelihood = "hypergeometric", N.units = 1000, by = 10)
#>
#> Classical Audit Sample Planning
#>
#> minimum sample size = 100
#> sample size obtained in 11 iterations via method 'hypergeometric'
```

The argument `max` sets the sample size at which the algorithm terminates. This can be used to avoid too many iterations of the algorithm at very low values of the performance materiality. For instance, `max = 50` throws an error if more than 100 samples are required.

```
planning(materiality = 0.03, expected = 0, conf.level = 0.95,
          likelihood = "hypergeometric", N.units = 1000, max = 50)
#> Error in planning(materiality = 0.03, expected = 0, conf.level = 0.95, : the
```

The sample size of 94 can be confirmed by checking that 94 is the minimum integer that results in less than 5% probability of finding 0 misstatements, given the assumption that the population misstatement is truly 3%. The `dhyper()` function calculates the probability of observing k misstatements in a sample of n items given the assumed hypergeometric distribution with N items and K assumed misstatements in the population. By calculating this probability for $n = 93$, we can show that this sample size is insufficient as the relevant probability is higher than the sampling risk α .

```
K <- ceiling(0.03 * 1000)
dhyper(x = 0, m = K, n = 1000 - K, k = 93) < 0.05
#> [1] FALSE
```

However, for $n = 94$ the relevant probability is lower than the sampling risk α and thus the sample size is considered to be sufficient.

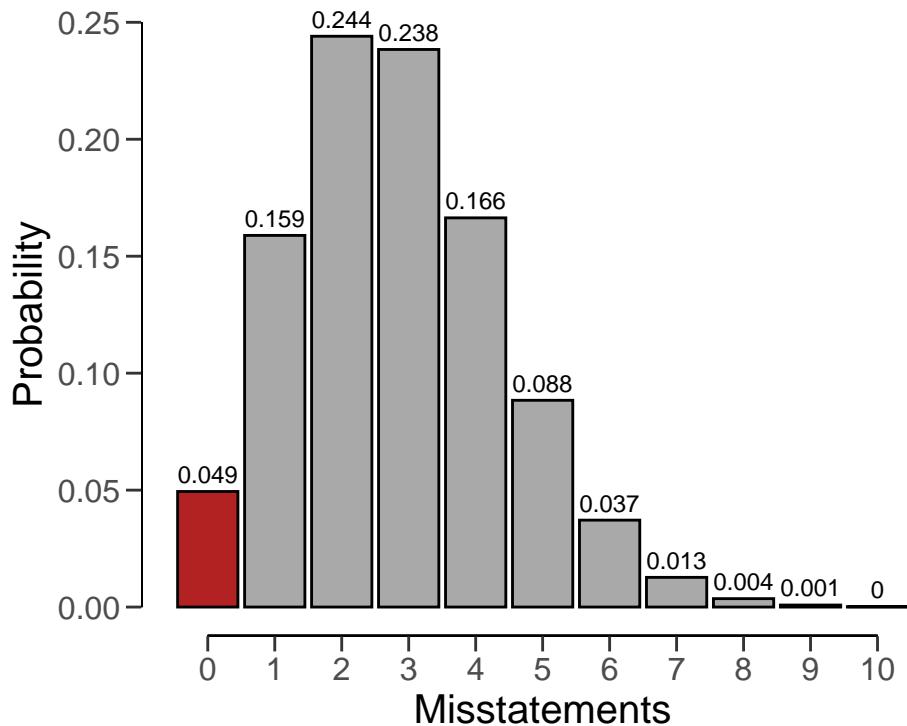
```
dhyper(x = 0, m = K, n = 1000 - K, k = 94) < 0.05
#> [1] TRUE
```

We can make this sample size visually intuitive by showing the $\text{hypergeometric}(k \mid 94,$

3.2. The Hypergeometric Likelihood

1000, 30) distribution and highlighting the probability for $k = 0$. This probability is lower than the required sampling risk $\alpha = 0.05$.

```
plot(plan)
```



3.2.1.2. Expected Misstatements

If the auditor expects misstatements in the sample, they can set k to any integer value, which consequently determines how the sample size can be calculated. As another example, if we want to achieve an assurance level of 95% ($\alpha = 0.05$) for a performance materiality of $\theta_{max} = 0.03$ in a population of $N = 1000$ items, then the required sample size under the assumption of one expected misstatement in the sample is $n = 147$.

```
plan <- planning(materiality = 0.03, expected = 1, conf.level = 0.95,
                   likelihood = "hypergeometric", N.units = 1000)
plan
#>
#> Classical Audit Sample Planning
#>
```

3. Planning

```
#> minimum sample size = 147  
#> sample size obtained in 146 iterations via method 'hypergeometric'
```

Once again, the sample size of 147 can be confirmed by checking that 147 is the minimum integer that results in less than 5% probability of finding 0 or 1 misstatements, given the assumption that the population misstatement is truly 3%. By calculating this probability for $n = 146$, we can show that this sample size is insufficient as the relevant probability is higher than the sampling risk α .

```
sum(dhyper(x = 0:1, m = K, n = 1000 - K, k = 146)) < 0.05  
#> [1] FALSE
```

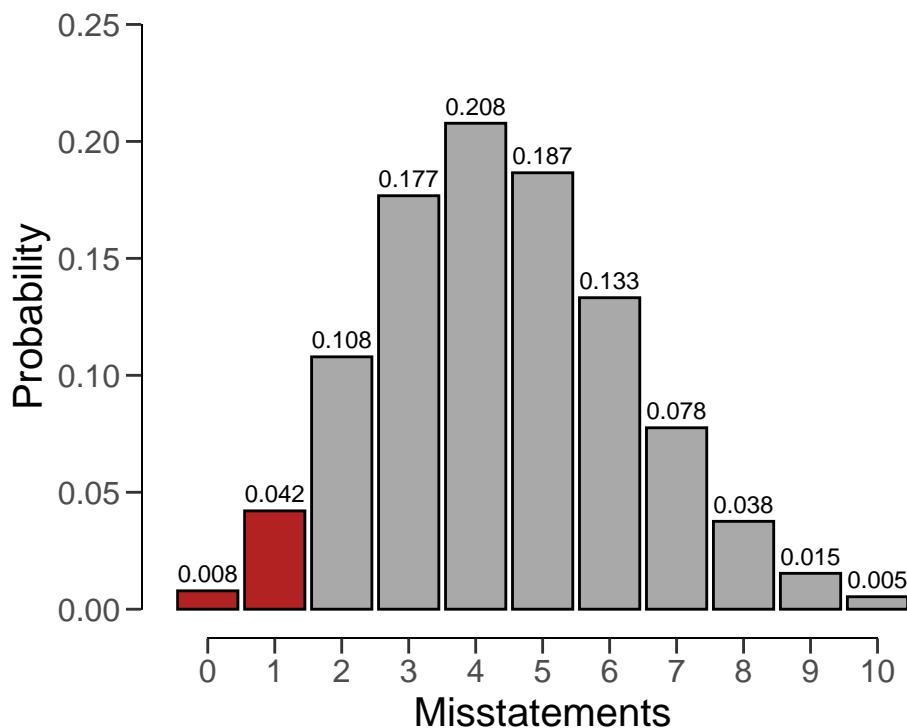
However, for $n = 147$ the relevant probability is lower than the sampling risk α and thus the sample size is considered to be sufficient.

```
sum(dhyper(x = 0:1, m = K, n = 1000 - K, k = 147)) < 0.05  
#> [1] TRUE
```

Like before, we can make this sample size visually intuitive by showing the $\text{hypergeometric}(k \mid 147, 1000, 30)$ distribution and highlighting the probabilities for $k = 0$ and $k = 1$. The sum of these probabilities is lower than the required sampling risk $\alpha = 0.05$.

```
plot(plan)
```

3.2. The Hypergeometric Likelihood



3.2.2. Bayesian Planning

Performing Bayesian planning with the hypergeometric likelihood (Dyer & Pierce, 1993) requires that you specify a prior distribution for the total misstatements K . Practically, this means that you should provide an input for the `prior` argument in the `planning()` function.

Setting `prior = TRUE` performs Bayesian planning using a default prior conjugate to the specified `likelihood` (i.e., a beta-binomial prior). Because this is a Bayesian analysis, the following statistical model is specified:

$$k \sim \text{Hypergeometric}(n, N, K) \quad (3.4)$$

$$K \sim \text{Beta-binomial}(N, \alpha, \beta) \quad (3.5)$$

Note

The beta-binomial prior distribution is the conjugate prior for to the hypergeometric likelihood (see this list of conjugate priors), which means that the posterior distribution of K can be determined analytically. For example, if the prior distri-

3. Planning

bution for K is:

$$K \sim \text{Beta-binomial}(N, \alpha, \beta) \quad K = 0, \dots, N \quad (3.6)$$

and the auditor has observed a sample of n items containing k misstatements, then the posterior distribution for K is:

$$K \sim \text{Beta-binomial}(N - n, \alpha + k, \beta + k - n) \quad K = k, k + 1, \dots, N - n + k. \quad (3.7)$$

3.2.2.1. No Expected Misstatements

Planning for no expected misstatements in the sample can be done by setting the value for the `expected` argument to zero. If we want to achieve an assurance level of 95% ($\alpha = 0.05$) for a performance materiality of $\theta_{max} = 0.1$ in a population of $N = 20$ items, then the required sample size under the assumption of zero expected misstatements in the sample is $n = 15$. The command below uses a default beta-binomial($N, 1, 1$) prior distribution to plan this sample, since `planning()` is given the hypergeometric likelihood.

```
plan <- planning(materiality = 0.1, expected = 0, conf.level = 0.95,
                   likelihood = "hypergeometric", N.units = 20,
                   prior = TRUE)
```

The `summary()` function can be used to obtain relevant information about the planning.

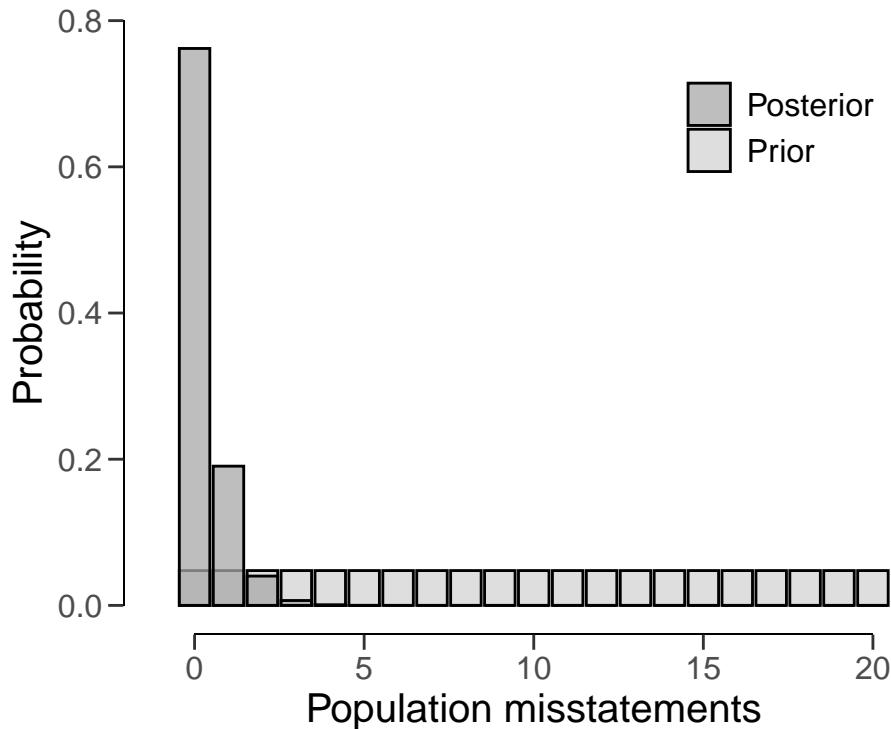
```
summary(plan)
#>
#> Bayesian Audit Sample Planning Summary
#>
#> Options:
#>   Confidence level:           0.95
#>   Population size:            20
#>   Materiality:                0.1
#>   Hypotheses:                 H : θ > 0.1 vs. H : θ < 0.1
#>   Expected:                     0
#>   Likelihood:                  hypergeometric
#>   Prior distribution:          beta-binomial(N = 20,    = 1,    = 1)
#>
#> Results:
#>   Minimum sample size:         15
```

3.2. The Hypergeometric Likelihood

```
#> Tolerable errors: 0
#> Posterior distribution: beta-binomial(N = 5,    = 1,    = 16)
#> Expected most likely error: 0
#> Expected upper bound: 0.05
#> Expected precision: 0.05
#> Expected BF : 190
```

You can inspect how the prior distribution compares to the expected posterior distribution by using the `plot()` function. The expected posterior distribution is the posterior distribution that would occur if you actually observed the planned sample containing the expected misstatements. Note that the posterior distribution is only defined in the range $[k; N - n + k]$, since a part of the population has already been seen.

```
plot(plan)
```



3.2.2.2. Expected Misstatements

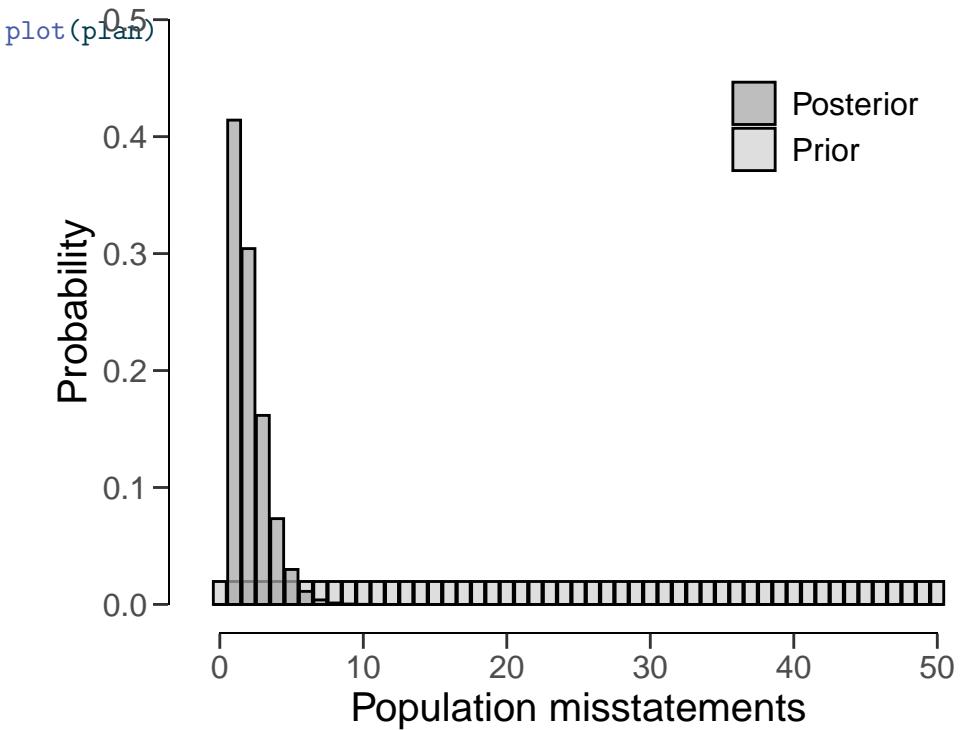
Planning for expected misstatements in the sample can be done by setting the value for the `expected` argument to a different value than zero. For example, the command below calculates the minimum sample size to achieve an assurance level of 95% ($\alpha = 0.05$) for

3. Planning

a performance materiality of $\theta_{max} = 0.1$ in a population of $N = 50$ items, given one expected misstatement in the sample. This sample size is $n = 32$.

```
plan <- planning(materiality = 0.1, expected = 1, conf.level = 0.95,
                   likelihood = "hypergeometric", N.units = 50,
                   prior = TRUE)
plan
#>
#> Bayesian Audit Sample Planning
#>
#> minimum sample size = 32
#> sample size obtained in 31 iterations via method 'hypergeometric' + 'prior'
```

Like before, you can inspect how the prior distribution compares to the expected posterior distribution by using the `plot()` function.



3.3. The Binomial Likelihood

The binomial distribution is a discrete probability distribution that is commonly used to model the number of events occurring in a fixed number of trials. It is similar to the hypergeometric distribution, however, it assumes that samples are drawn from the population with replacement. For our purpose, we can use the binomial distribution as a likelihood to model the number of misstatements that are expected to be found in the sample.

i Note

In audit sampling, the binomial likelihood is often used to approximate the hypergeometric likelihood since it is easier to work with (i.e., it only has two parameters: θ and n , while the hypergeometric has three: n , N , and K). However, the binomial likelihood is more conservative than the hypergeometric likelihood, meaning that resulting sample sizes will be higher.

The probability mass function (PMF) of the binomial distribution is given by:

$$p(k; n, \theta) = \binom{n}{k} \theta^k (1 - \theta)^{n-k}, \quad (3.8)$$

where k is the number of misstatements in the sample, n is the sample size and θ is the probability of misstatement in the population. Let's consider how to use the binomial likelihood to calculate the minimum sample size needed to achieve a certain level of assurance.

3.3.1. Classical Planning

In classical planning using the binomial likelihood, the following statistical model is specified:

$$k \sim \text{Binomial}(n, \theta_{max}) \quad (3.9)$$

3.3.1.1. No Expected Misstatements

If the auditor does not expect any misstatements in the sample, they can set $k = 0$, which consequently determines how the sample size can be calculated. Given a performance materiality θ_{max} , we can solve for the minimum sample size n needed to achieve a certain assurance level. A useful trick to utilize is that, if we do not expect any misstatements in the sample, the formula for the minimum required sample size reduces to:

3. Planning

$$n = \lceil \frac{\ln(\alpha)}{\ln(1 - \theta_{max})} \rceil. \quad (3.10)$$

i Note

[...] is the ceiling function, which means that $\lceil 1.2 \rceil = 2$.

For example, if we want to achieve an assurance level of 95% ($\alpha = 0.05$) for a performance materiality of $\theta_{max} = 0.03$, then the required sample size under the assumption of zero expected misstatements in the sample is $n = 99$.

```
ceiling(log(1 - 0.95) / log(1 - 0.03))
#> [1] 99
```

In **jfa**, this sample size can be replicated using the following code:

```
plan <- planning(materiality = 0.03, expected = 0, conf.level = 0.95,
                   likelihood = "binomial")
plan
#>
#> Classical Audit Sample Planning
#>
#> minimum sample size = 99
#> sample size obtained in 100 iterations via method 'binomial'
```

The sample size of 99 can be confirmed by checking that 99 is the minimum integer that results in less than 5% probability of finding 0 misstatements, given the assumption that the population misstatement is truly 3%. The `dbinom()` function calculates the probability of observing k misstatements in a sample of n items given an assumed misstatement probability θ_{max} . By calculating this probability for $n = 98$, we can show that this sample size is insufficient as the relevant probability is higher than the sampling risk α .

```
dbinom(x = 0, size = 98, prob = 0.03) < 0.05
#> [1] FALSE
```

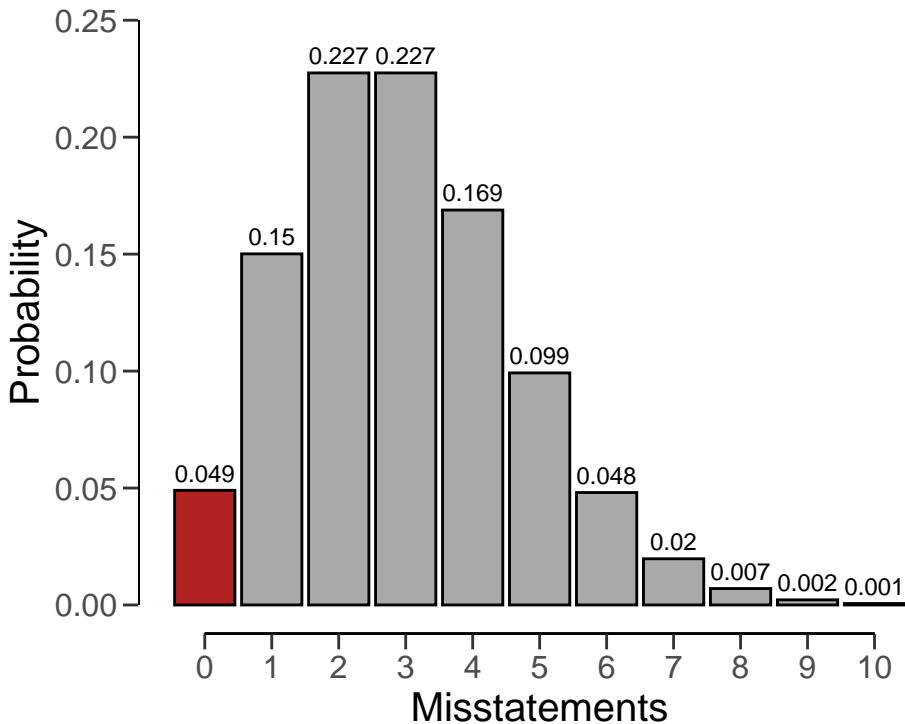
However, for $n = 99$ the relevant probability is lower than the sampling risk α and thus the sample size is considered to be sufficient.

```
dbinom(x = 0, size = 99, prob = 0.03) < 0.05
#> [1] TRUE
```

3.3. The Binomial Likelihood

We can make this sample size visually intuitive by showing the binomial($k \mid 99, 0.03$) distribution and highlighting the probability for $k = 0$. This probability is lower than the required sampling risk $\alpha = 0.05$.

```
plot(plan)
```



3.3.1.2. Expected Misstatements

However, if the number of expected misstatements in the sample is non-zero, it becomes more difficult to solve the formula for n . Hence, they will need to set k to a different integer value, which consequently determines how the sample size is calculated. Here, we can iteratively try every value of n and return the smallest integer that satisfies the sampling objectives.

In **jfa**, this can be done by adjusting the `expected` argument in the `planning()` function. For example, if we want to achieve an assurance level of 95% ($\alpha = 0.05$) for a performance materiality of $\theta_{max} = 0.03$, then the required sample size under the assumption of one expected misstatement in the sample is $n = 157$.

3. Planning

```
plan <- planning(materiality = 0.03, expected = 1, conf.level = 0.95,  
                   likelihood = "binomial")
```

Once again, the sample size of 157 can be confirmed by checking that 157 is the minimum integer that results in less than 5% probability of finding 0 or 1 misstatements, given the assumption that the population misstatement is truly 3%. By calculating this probability for $n = 156$, we can show that this sample size is insufficient as the relevant probability is higher than the sampling risk α .

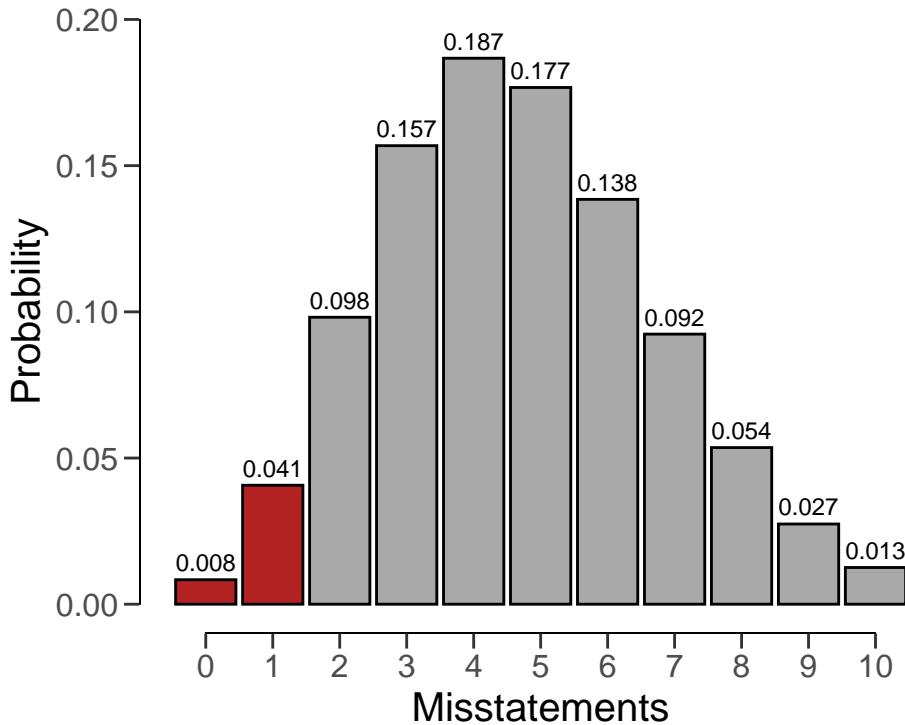
```
sum(dbinom(x = 0:1, size = 156, prob = 0.03)) < 0.05  
#> [1] FALSE
```

However, for $n = 157$ the relevant probability is lower than the sampling risk α and thus the sample size is considered to be sufficient.

```
sum(dbinom(x = 0:1, size = 157, prob = 0.03)) < 0.05  
#> [1] TRUE
```

Like before, we can make this sample size visually intuitive by showing the binomial($k | 157, 0.03$) distribution and highlighting the probabilities for $k = 0$ and $k = 1$. The sum of these probabilities is lower than the required sampling risk $\alpha = 0.05$.

```
plot(plan)
```



3.3.1.3. Expected Misstatement Rate

When the expected misstatement rate in the sample θ_{exp} is assessed, the value for k can be determined as $k = n\theta_{exp}$, which consequently determines how the sample size can be calculated.

To account for the fact that k can have non-integer values in this case, we can use a well-known similarity between the binomial distribution and the beta distribution to plan the sample size. The upper bound for any binomial($k; n, \theta$) distributed variable can also be obtained via percentiles of the beta($1 + k, n - k$) distribution.

For example, the upper bound for a sample of $n = 10$ items containing $k = 2$ misstatements, when calculated via the traditional `binom.test()` is:

```
ub_binom <- binom.test(x = 2, n = 10, p = 0.03, conf.level = 0.95,
                        alternative = "less")$conf.int[2]
ub_binom
#> [1] 0.5069013
```

When calculated via the beta relationship, the upper bound is:

3. Planning

```
ub_beta <- qbeta(p = 0.95, shape1 = 1 + 2, shape2 = 10 - 2)
ub_beta
#> [1] 0.5069013
```

It can be validated that the two approaches result in the same upper bound via:

```
ub_binom == ub_beta
#> [1] TRUE
```

This relationship between the binomial likelihood and the beta distribution is deliberately **not** used in **jfa**. That is because, in the case of the binomial distribution, the auditing standards round the tolerable misstatements upwards to a whole number. For example, if we try to call the **planning()** function with the argument **expected = 1.5**, **jfa** will internally convert this to **expected = 2** and base the sample size on this in order to stay compliant with American Institute of Certified Public Accountants (AICPA) (2016a). The resulting sample size is $n = 208$ in this case.

```
planning(materiality = 0.03, expected = 1.5, conf.level = 0.95,
          likelihood = "binomial")
#> Using 'expected = 2' since 'expected' must be a single integer >= 0
#>
#> Classical Audit Sample Planning
#>
#> minimum sample size = 208
#> sample size obtained in 206 iterations via method 'binomial'
```

3.3.2. Bayesian Planning

Performing Bayesian planning using the binomial likelihood requires that you specify a prior distribution for the parameter θ . Practically, this means that you should provide an input for the **prior** argument in the **planning()** function.

Setting **prior = TRUE** performs Bayesian planning using a default prior conjugate to the specified **likelihood** (i.e., a beta prior). Because this is a Bayesian analysis, the following statistical model is specified:

$$k \sim \text{Binomial}(n, \theta) \quad (3.11)$$

$$\theta \sim \text{Beta}(\alpha, \beta) \quad (3.12)$$

i Note

The beta prior distribution is the conjugate prior for the binomial likelihood (see this list of conjugate priors), which means that the posterior distribution of θ can be determined analytically. For example, if the prior distribution for θ is:

$$\theta \sim \text{Beta}(\alpha, \beta) \quad \theta \in [0, 1] \quad (3.13)$$

and the auditor has observed a sample of n items containing k misstatements, then the posterior distribution for θ is:

$$\theta \sim \text{Beta}(\alpha + k, \beta + n - k) \quad \theta \in [0, 1]. \quad (3.14)$$

For example, the command below uses a default beta($\alpha = 1, \beta = 1$) prior distribution to plan the sample, since `planning()` is given the binomial likelihood. If we want to achieve an assurance level of 95% ($\alpha = 0.05$) for a performance materiality of $\theta_{max} = 0.03$, then the required sample size under the assumption of zero expected misstatements in the sample is $n = 98$.

```
plan <- planning(materiality = 0.03, expected = 0, conf.level = 0.95,
                   likelihood = "binomial", prior = TRUE)
```

The `summary()` function can be used to obtain relevant information about the planning.

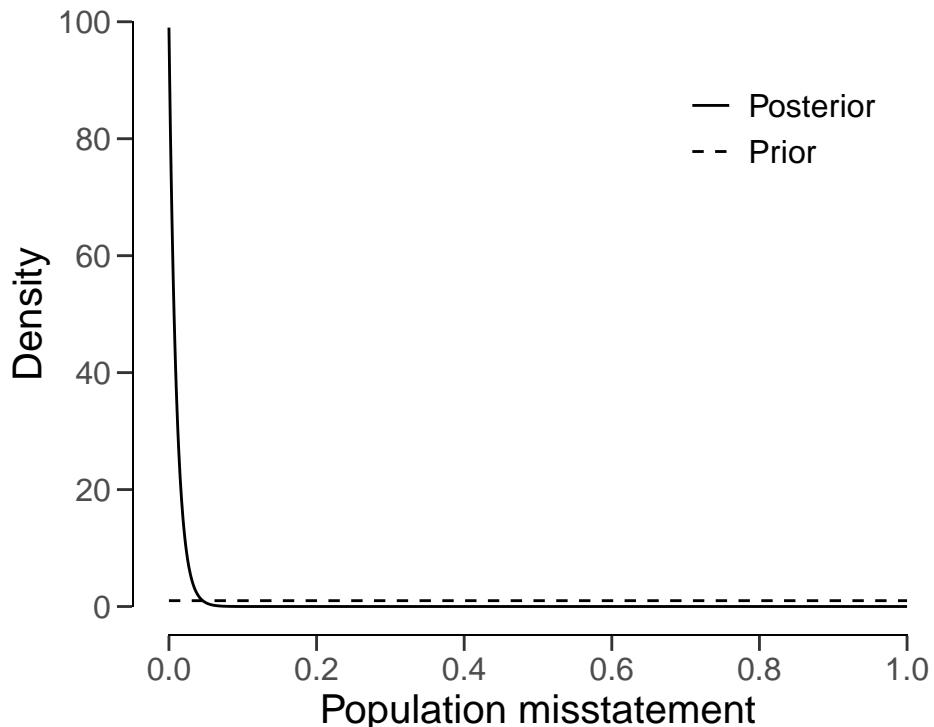
```
summary(plan)
#>
#> Bayesian Audit Sample Planning Summary
#>
#> Options:
#>   Confidence level:          0.95
#>   Materiality:              0.03
#>   Hypotheses:                H : θ > 0.03 vs. H : θ < 0.03
#>   Expected:                  0
#>   Likelihood:                binomial
#>   Prior distribution:        beta( = 1,    = 1)
#>
#> Results:
#>   Minimum sample size:       98
#>   Tolerable errors:          0
#>   Posterior distribution:    beta( = 1,    = 99)
#>   Expected most likely error: 0
#>   Expected upper bound:      0.029807
```

3. Planning

```
#> Expected precision: 0.029807  
#> Expected BF : 627.22
```

You can inspect how the prior distribution compares to the expected posterior distribution by using the `plot()` function. The expected posterior distribution is the posterior distribution that would occur if you actually observed the planned sample containing the expected misstatements.

```
plot(plan)
```



The input for the `prior` argument can also be an object created by the `auditPrior` function. If `planning()` receives a prior for which there is no conjugate likelihood available, it will numerically derive the posterior distribution. For example, the command below uses a $\text{Normal}(0, 0.05)$ prior distribution to plan the sample using the binomial likelihood. Because this is a Bayesian analysis, the following statistical model is specified:

$$k \sim \text{Binomial}(n, \theta) \quad (3.15)$$

$$\theta \sim \text{Normal}(\mu = 0, \sigma = 0.05) \quad (3.16)$$

3.3. The Binomial Likelihood

```
prior <- auditPrior(method = "param", likelihood = "normal",
                     alpha = 0, beta = 0.05)

plan <- planning(materiality = 0.03, expected = 0, conf.level = 0.95,
                  likelihood = "poisson", prior = prior)
```

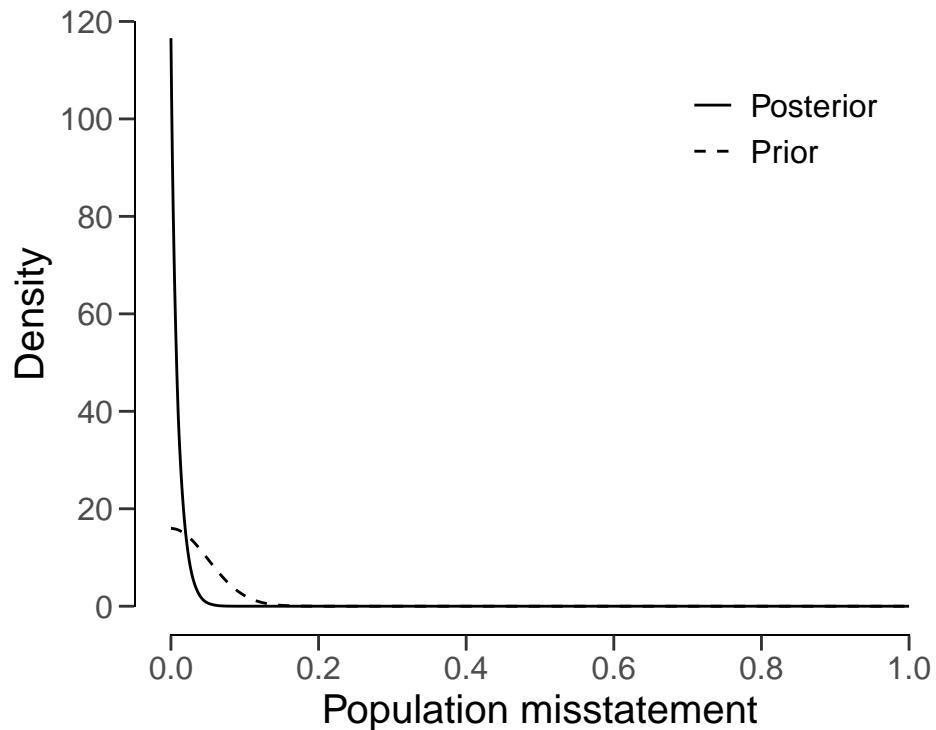
The `summary()` function can be used to obtain relevant information about the planning.

```
summary(plan)
#>
#> Bayesian Audit Sample Planning Summary
#>
#> Options:
#>   Confidence level:          0.95
#>   Materiality:              0.03
#>   Hypotheses:                H : θ > 0.03 vs. H : θ < 0.03
#>   Expected:                  0
#>   Likelihood:                poisson
#>   Prior distribution:        normal( = 0,  = 0.05)T[0,1]
#>
#> Results:
#>   Minimum sample size:       90
#>   Tolerable errors:          0
#>   Posterior distribution:    Nonparametric
#>   Expected most likely error: 0
#>   Expected upper bound:      0.029029
#>   Expected precision:        0.029029
#>   Expected BF :              19.08
```

The resulting sample size under this prior is $n = 90$, a reduction of 8 samples when compared to the default $\text{beta}(1, 1)$ prior distribution.

```
plot(plan)
```

3. Planning



3.4. The Poisson Likelihood

The Poisson distribution is a discrete probability distribution that is commonly used to model the number of events occurring in a fixed time or space. We can use the Poisson distribution as a likelihood to model the number of misstatements that are expected to be found in the sample.

i Note

In audit sampling, the Poisson likelihood is often used to approximate the binomial likelihood since it is easier to work with (i.e., it only has one parameter: λ , while the binomial has two parameters: θ and n). However, the Poisson likelihood is more conservative than the binomial likelihood, meaning that resulting sample sizes will be higher.

The probability mass function (PMF) of the Poisson distribution is given by:

$$p(k; \lambda) = \frac{\lambda^k e^{-\lambda}}{k!}, \quad (3.17)$$

where k is the number of misstatements in the sample, and λ is the average number of misstatements expected in the sample. The average number of misstatements is related to the misstatement rate in the population, denoted by θ , and the sample size, n , by the following equation:

$$\lambda = n\theta. \quad (3.18)$$

Let's consider how to use the Poisson likelihood to calculate the minimum sample size needed to achieve a certain level of assurance.

3.4.1. Classical planning

In classical planning using the Poisson likelihood, the following statistical model is specified:

$$k \sim \text{Poisson}(n\theta_{max}) \quad (3.19)$$

3.4.1.1. No Expected Misstatements

Given the performance materiality θ_{max} and the Poisson likelihood, we can solve for the minimum sample size n needed to achieve a certain assurance level. A useful trick to utilize is that, if we do not expect any misstatements in the sample, the formula for the required sample size reduces to:

$$n = \lceil -\frac{\ln(\alpha)}{\theta_{max}} \rceil. \quad (3.20)$$

For example, if we want to achieve an assurance level of 95% ($\alpha = 0.05$) for a performance materiality of $\theta_{max} = 0.03$, then the required sample size under the assumption of zero expected misstatements in the sample is $n = 100$.

```
ceiling(-log(1 - 0.95) / 0.03)
#> [1] 100
```

In **jfa**, this sample size can be replicated using the following code:

```
plan <- planning(materiality = 0.03, expected = 0, conf.level = 0.95,
                   likelihood = "poisson")
```

3. Planning

The sample size of 100 can be confirmed by checking that 100 is the minimum integer that results in less than 5% probability of finding 0 misstatements, given the assumption that the population misstatement is truly 3%. The `dpois()` function calculates the probability of observing k missatements in a sample of n items given an assumed misstatement probability θ_{max} . By calculating this probability for $n = 99$, we can show that this sample size is insufficient as the relevant probability is higher than the sampling risk α .

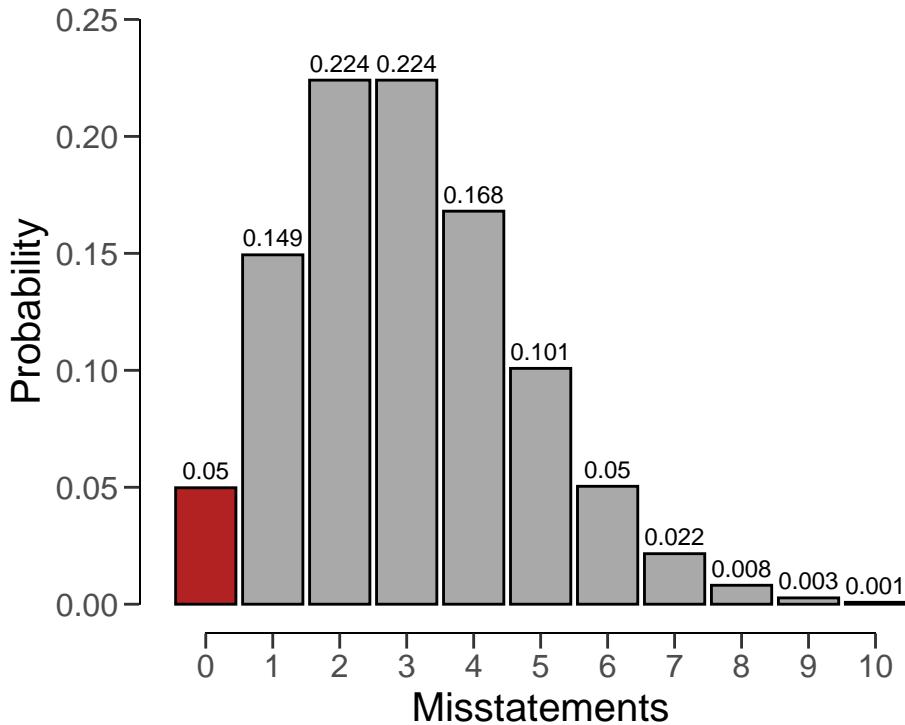
```
dpois(x = 0, lambda = 99 * 0.03) < 0.05  
#> [1] FALSE
```

However, for $n = 100$ the relevant probability is lower than the sampling risk α and thus the sample size is considered to be sufficient.

```
dpois(x = 0, lambda = 100 * 0.03) < 0.05  
#> [1] TRUE
```

We can make this visually intuitive by showing the $Poisson(k \mid 100 * 0.03)$ distribution and highlighting the probability for $k = 0$. This probability is lower than the required sampling risk $\alpha = 0.05$.

```
plot(plan)
```



3.4.1.2. Expected Misstatements

However, if the number of expected misstatements in the sample is non-zero, it becomes more difficult to solve the formula for n algebraically. Hence, they will need to set k to a different integer value. Next, we can iteratively try every value of n and return the smallest integer that satisfies the sampling objectives.

For example, if we want to achieve an assurance level of 95% ($\alpha = 0.05$) for a performance materiality of $\theta_{max} = 0.03$, then the required sample size under the assumption of one expected misstatement in the sample is $n = 159$.

```
plan <- planning(materiality = 0.03, expected = 1, conf.level = 0.95,
                    likelihood = "poisson")
```

Once again, the sample size of 159 can be confirmed by checking that 159 is the minimum integer that results in less than 5% probability of finding 0 or 1 misstatements, given the assumption that the population misstatement is truly 3%. By calculating this probability for $n = 158$, we can show that this sample size is insufficient as the relevant probability is higher than the sampling risk α .

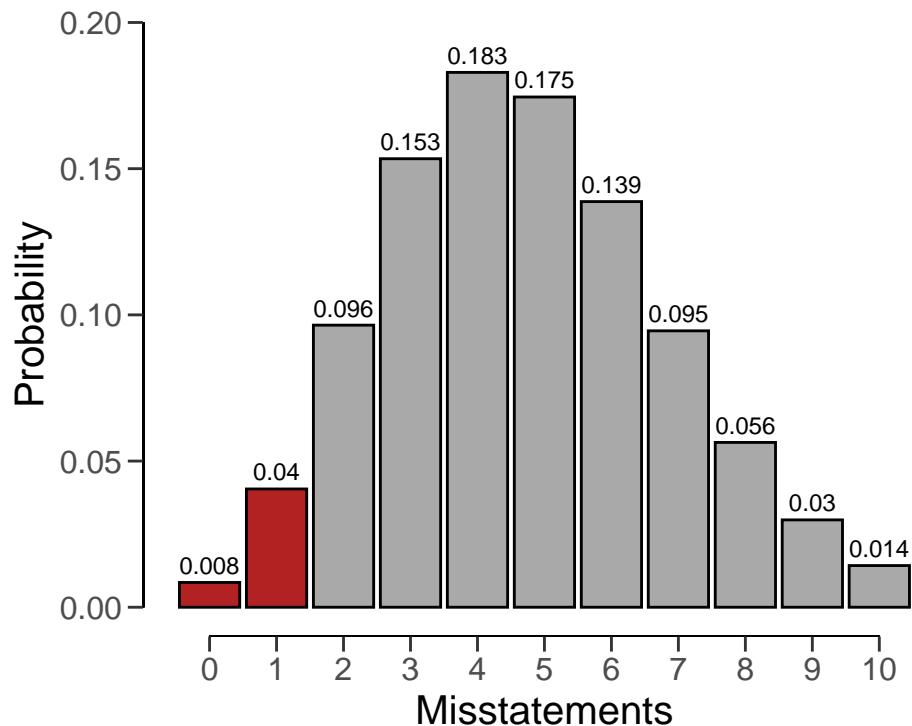
3. Planning

```
sum(dpois(x = 0:1, lambda = 158 * 0.03)) < 0.05
#> [1] FALSE
```

However, for $n = 159$ the relevant probability is lower than the sampling risk α and thus the sample size is considered to be sufficient.

```
sum(dpois(x = 0:1, lambda = 159 * 0.03)) < 0.05
#> [1] TRUE
```

Like before, we can make this visually intuitive by showing the $\text{Poisson}(k \mid 159 * 0.03)$ distribution and highlighting the probabilities for $k = 0$ and $k = 1$. The sum of these probabilities is lower than the required sampling risk $\alpha = 0.05$.



3.4.1.3. Expected Misstatement Rate

When the expected misstatements in the sample θ_{exp} is assessed, the value for k can be determined as $k = n\theta_{exp}$, which consequently determines how the sample size can be calculated.

To account for the fact that k can have non-integer values in this case, we use a well-known similarity between the Poisson distribution and the gamma distribution to plan

3.4. The Poisson Likelihood

the sample size. The upper bound for any $\text{Poisson}(k; n\theta)$ distributed variable can also be obtained via percentiles of the $\text{gamma}(1 + k, n)$ distribution.

For example, the upper bound for a sample of $n = 10$ items containing $k = 2$ misstatements, when calculated via the traditional `poisson.test()` is:

```
ub_pois <- poisson.test(x = 2, T = 10, r = 0.03, conf.level = 0.95,
                         alternative = "less")$conf.int[2]
ub_pois
#> [1] 0.6295794
```

When calculated via the relationship with the gamma distribution, the upper bound is:

```
ub_gamma <- qgamma(p = 0.95, shape = 1 + 2, rate = 10)
ub_gamma
#> [1] 0.6295794
```

It can be validated that the two approaches result in the same upper bound via:

```
ub_pois == ub_gamma
#> [1] TRUE
```

This relationship between the Poisson likelihood and the gamma distribution is used under the hood in `jfa`. For example, if we want to achieve an assurance level of 95% ($\alpha = 0.05$) for a performance materiality of $\theta_{max} = 0.03$, then the required sample size under the assumption of 1.5 expected misstatements in the sample is $n = 185$.

```
planning(materiality = 0.03, expected = 1.5, conf.level = 0.95,
          likelihood = "poisson")
#>
#> Classical Audit Sample Planning
#>
#> minimum sample size = 185
#> sample size obtained in 184 iterations via method 'poisson'
```

The sample size of 185 can be confirmed by checking that 185 is the minimum integer that results in less than 5% probability of finding a misstatement rate in the population equal to, or higher than, 3%. By calculating this probability for $n = 184$, we can show that this sample size is insufficient as the relevant upper bound is higher than the performance materiality θ_{max} .

```
qgamma(p = 0.95, shape = 1 + 1.5, rate = 184) < 0.03
#> [1] FALSE
```

3. Planning

However, for $n = 185$ the relevant upper bound is lower than the performance materiality θ_{max} and thus the sample size is sufficient.

```
qgamma(p = 0.95, shape = 1 + 1.5, rate = 185) < 0.03  
#> [1] TRUE
```

3.4.2. Bayesian Planning

Performing Bayesian planning with the Poisson likelihood requires that you specify a prior distribution for the parameter θ . Practically, this means that you should provide an input for the `prior` argument in the `planning()` function.

Setting `prior = TRUE` performs Bayesian planning using a default prior conjugate to the specified `likelihood` (i.e., a gamma prior). Because this is a Bayesian analysis, the following statistical model is specified:

$$k \sim \text{Poisson}(n\theta) \quad (3.21)$$

$$\theta \sim \text{Gamma}(\alpha, \beta) \quad (3.22)$$

i Note

The gamma prior distribution is the conjugate prior for the Poisson likelihood (see this list of conjugate priors), which means that the posterior distribution of θ can be determined analytically. For example, if the prior distribution for θ is:

$$\theta \sim \text{Gamma}(\alpha, \beta) \quad \theta \in [0, \infty] \quad (3.23)$$

and the auditor has observed a sample of n items containing k misstatements, then the posterior distribution for θ is:

$$\theta \sim \text{Gamma}(\alpha + k, \beta + n) \quad \theta \in [0, \infty]. \quad (3.24)$$

For example, the command below uses a default $\text{gamma}(\alpha = 1, \beta = 1)$ prior distribution to plan the sample, since `planning()` is given the Poisson likelihood. If we want to achieve an assurance level of 95% ($\alpha = 0.05$) for a performance materiality of $\theta_{max} = 0.03$, then the required sample size under the assumption of zero expected misstatements in the sample is $n = 99$.

```
plan <- planning(materiality = 0.03, expected = 0, conf.level = 0.95,  
                   likelihood = "poisson", prior = TRUE)
```

3.4. The Poisson Likelihood

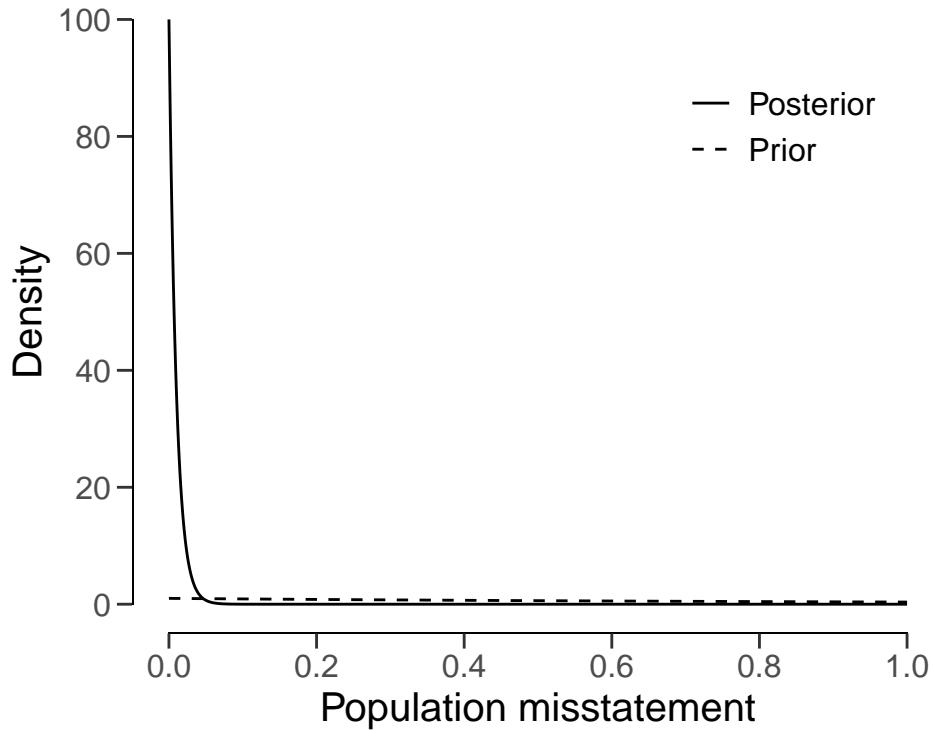
The `summary()` function can be used to obtain relevant information about the planning.

```
summary(plan)
#>
#> Bayesian Audit Sample Planning Summary
#>
#> Options:
#> Confidence level:          0.95
#> Materiality:              0.03
#> Hypotheses:                 H : θ > 0.03 vs. H : θ < 0.03
#> Expected:                  0
#> Likelihood:                poisson
#> Prior distribution:         gamma( = 1,   = 1)
#>
#> Results:
#> Minimum sample size:       99
#> Tolerable errors:          0
#> Posterior distribution:    gamma( = 1,   = 100)
#> Expected most likely error: 0
#> Expected upper bound:      0.029957
#> Expected precision:        0.029957
#> Expected BF :              626.69
```

You can inspect how the prior distribution compares to the expected posterior distribution by using the `plot()` function. The expected posterior distribution is the posterior distribution that would occur if you actually observed the planned sample containing the expected misstatements.

```
plot(plan)
```

3. Planning



The input for the `prior` argument can also be an object created by the `auditPrior` function. If `planning()` receives a prior for which there is no conjugate likelihood available, it will numerically derive the posterior distribution. For example, the command below uses a $\text{Normal}(0, 0.05)$ prior distribution to plan the sample using the Poisson likelihood. Concretely, this means that the following statistical model is specified:

$$k \sim \text{Poisson}(n\theta) \quad (3.25)$$

$$\theta \sim \text{Normal}(\mu = 0, \sigma = 0.05) \quad (3.26)$$

```
prior <- auditPrior(method = "param", likelihood = "normal",
                     alpha = 0, beta = 0.05)

plan <- planning(materiality = 0.03, expected = 0, conf.level = 0.95,
                  likelihood = "poisson", prior = prior)
```

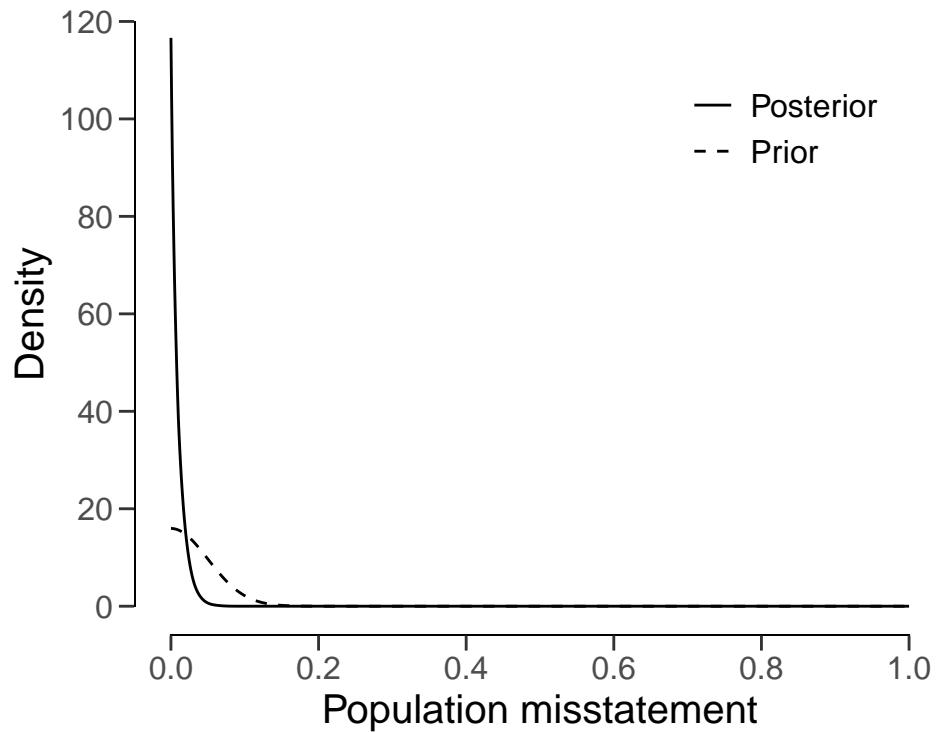
The `summary()` function can be used to obtain relevant information about the planning.

```
summary(plan)
#>
#> Bayesian Audit Sample Planning Summary
#>
#> Options:
#>   Confidence level:          0.95
#>   Materiality:              0.03
#>   Hypotheses:                H : θ > 0.03 vs. H : θ < 0.03
#>   Expected:                  0
#>   Likelihood:                poisson
#>   Prior distribution:        normal( = 0,  = 0.05)T[0,1]
#>
#> Results:
#>   Minimum sample size:       91
#>   Tolerable errors:          0
#>   Posterior distribution:    Nonparametric
#>   Expected most likely error: 0
#>   Expected upper bound:      0.029029
#>   Expected precision:        0.029029
#>   Expected BF :              19.296
```

The resulting sample size under this prior is $n = 91$, a reduction of 8 samples when compared to the default $\text{gamma}(1, 1)$ prior.

```
plot(plan)
```

3. Planning



3.5. Practical Examples

This section contains practical examples of how to conduct the planning of statistical audit samples and demonstrates how to set up a prior distribution based on various types of relevant audit information.

3.5.1. Audit Risk Model

In our first example, an auditor is performing tests of details on a population of the auditee. For instance, let's say an auditor is performing an audit on a company's accounts payable transactions. The company has a total of $N = 1000$ accounts payable transactions for the year. Rather than testing all 1000 transactions, the auditor can choose to test a sample of the transactions. The performance materiality for the payable transactions account is set to $\theta_{max} = 3\%$, and the audit risk is set to $\alpha = 0.05$, or 5%. Based on the results of last years audit, where the most likely estimate of the misstatement was 1%, the auditor wants to tolerate 1% misstatements in the sample before giving an unqualified opinion on the population.

```
ar      <- 0.05 # Audit risk
materiality <- 0.03 # Performance materiality
expected    <- 0.01 # Tolerable deviation rate
```

Before tests of details, the auditor has assessed risk of material misstatement via the audit risk model. In this example, the auditor has assessed the effectiveness of the company's internal controls, such as its segregation of duties and its risk management processes, and has determined that they are sufficient to prevent or detect material misstatements. Because the internal control systems were effective, the auditor assesses the control risk as medium. The auditor's firm defines the risk categories low, medium, and high respectively as 50%, 60%, and 100%. According to the Audit Risk Model, the detection risk can be calculated as a function of the audit risk, the inherent risk and the control risk.

```
ir <- 1          # Inherent risk
cr <- 0.6        # Control risk
dr <- ar / (ir * cr) # Detection risk
dr
#> [1] 0.08333333
```

By using the detection risk of 8.33% as the sampling risk for this population, the auditor can plan for a sample while taking into account the risk-reducing information. The required minimum sample size is 174 in this case.

```
planning(materiality = 0.03, expected = expected, conf.level = 1 - dr)
#>
#> Classical Audit Sample Planning
#>
#> minimum sample size = 174
#> sample size obtained in 175 iterations via method 'poisson'
```

The example above is a frequentist one. However, the auditor is free to apply a Bayesian philosophy in planning the sample. For example, the risk assessments from the ARM can be incorporated into a prior distribution. This can be done using `method = "arm"` in the `auditPrior()` function, which takes the values of the inherent risk probability `ir` and the control risk probability `cr`. Hence, the prior distribution in this example can be constructed using the following command:

```
prior <- auditPrior(method = "arm", materiality = 0.03, expected = expected,
                     ir = ir, cr = cr)
```

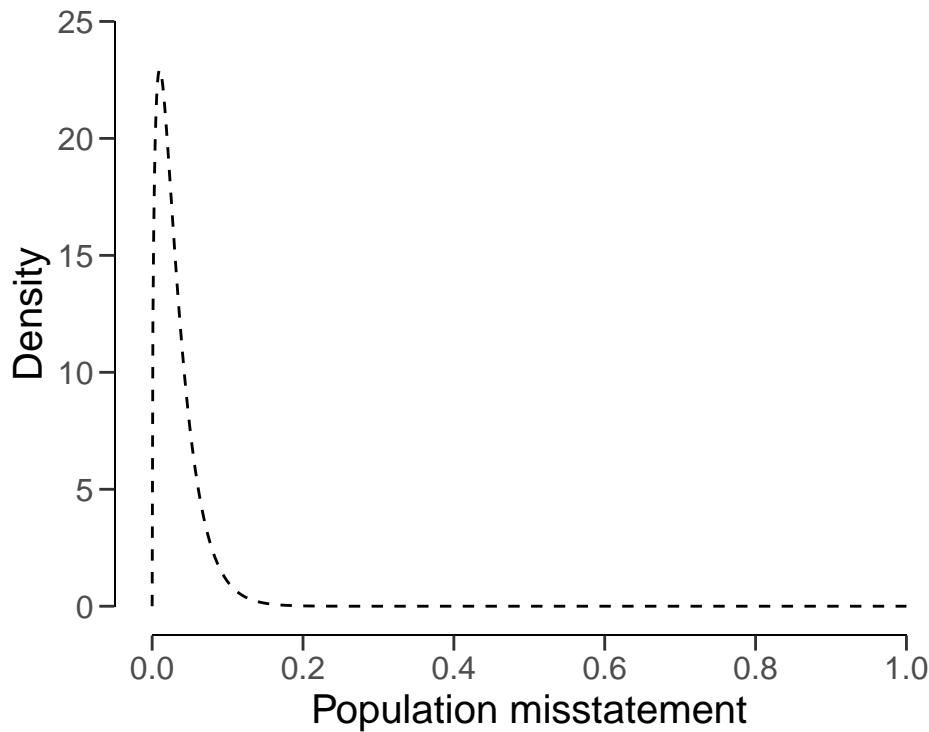
The `summary()` function can be used to obtain relevant information about the prior distribution.

3. Planning

```
summary(prior)
#>
#> Prior Distribution Summary
#>
#> Options:
#> Likelihood: poisson
#> Specifics: ir = 1; cr = 0.6; dr = 0.0833333
#>
#> Results:
#> Functional form: gamma( = 1.46, = 46)
#> Mode: 0.01
#> Mean: 0.031739
#> Median: 0.024859
#> Variance: 0.00069
#> Skewness: 1.6552
#> Information entropy (nat): -2.4894
#> 95 percent upper bound: 0.08343
#> Precision: 0.07343
```

Furthermore, the prior distribution can be visualized with a call to the `plot()` function.

```
plot(prior)
```



By using the prior distribution to incorporate the assessments of the inherent risk and the control risk, the auditor can plan a sample while taking into account the risk-reducing information. The required minimum sample size is also 174 in this case.

```
planning(materiality = 0.03, expected = expected, conf.level = 1 - ar,
          prior = prior)
#>
#> Bayesian Audit Sample Planning
#>
#> minimum sample size = 174
#> sample size obtained in 175 iterations via method 'poisson' + 'prior'
```

3.5.2. Benchmark Analysis

The auditor may incorporate information obtained through analytical procedures (Derkx, Swart, Batenburg, et al., 2021), such as a benchmark analysis, into the prior distribution for θ . While we have previously discussed methods for constructing a prior distribution based on existing knowledge, there is no set procedure for incorporating information obtained through analytical procedures, as these procedures can vary significantly depending on the type of information being incorporated into the prior distribution. Therefore,

3. Planning

it is important to thoroughly substantiate the data and assumptions used in this approach and to carefully consider how these assumptions are incorporated into the prior distribution.

One way to construct a prior distribution on the basis of data is through the use of regression models, such as benchmarking the relationship between sales and costs of sales within the auditee's specific industry sector. The **jfa** package includes a data set **benchmark** that can be used for this example.

```
data(benchmark)
head(benchmark)
#>   sales costofsales
#> 1 186273256 140755372
#> 2 336491541 248675452
#> 3 222693077 164299866
#> 4 364905221 285768790
#> 5 382140185 280187371
#> 6 113666950 101552955
```

The auditee's the sum of the sales is \$298,112,312 and the sum of the booked costs of sales is \$223,994,405, respectively. This is indicated by a blue dot in the figure below, which visualizes the industry sales versus the cost of sales.

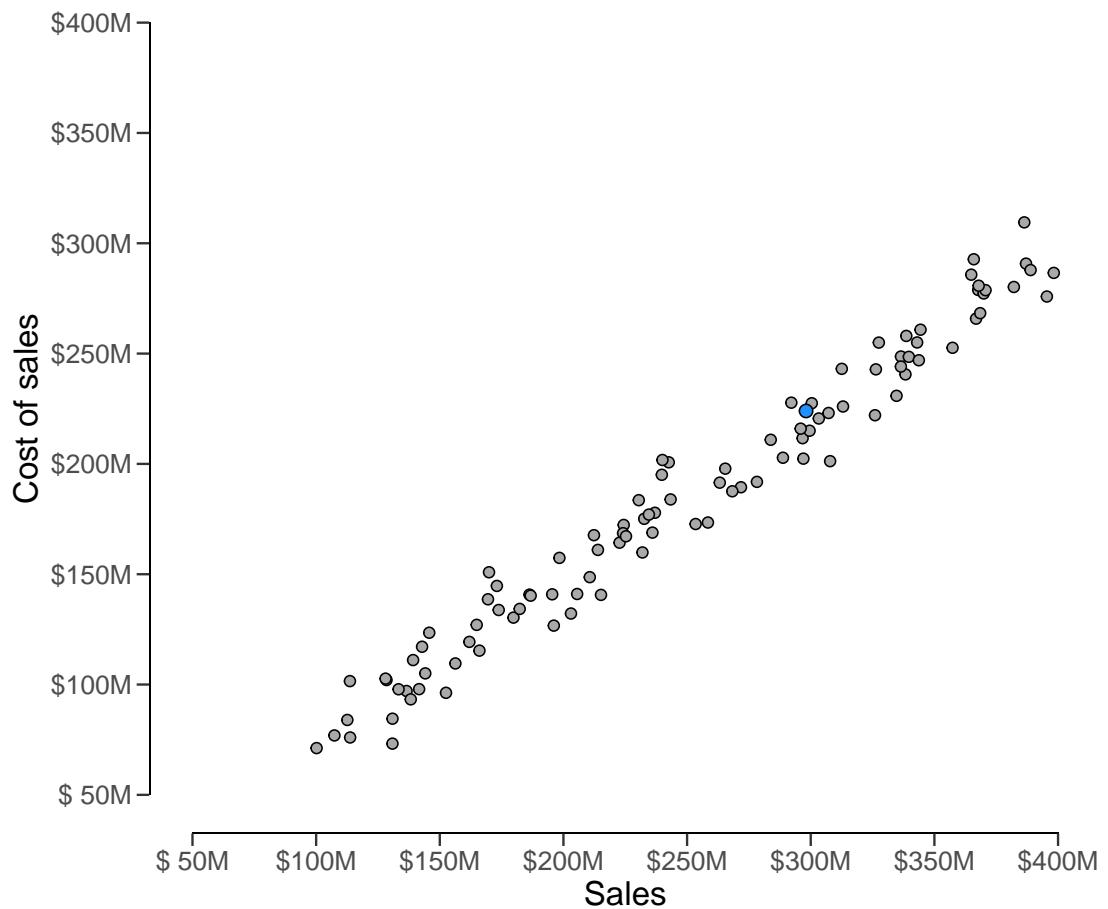
```
C_real <- 223994405
```

The relationship between the sales S and the cost of sales C can be modelled by a linear equation:

$$C = \beta_0 + \beta_1 \cdot S + \epsilon. \quad (3.27)$$

In practice, this relationship is often more complex than is presented above, and the auditor must carefully construct and evaluate the applied regression model. However, for ease of understanding we will continue our example with this simplified model. The auditor can estimate the regression model using the following command:

```
fit <- lm(costofsales ~ 1 + sales, data = benchmark)
summary(fit)
#>
#> Call:
#> lm(formula = costofsales ~ 1 + sales, data = benchmark)
#>
#> Residuals:
#>      Min       1Q   Median       3Q      Max
#>
```



3. Planning

```
#> -25736696 -7052141 -226945 6857840 25498106
#>
#> Coefficients:
#>             Estimate Std. Error t value Pr(>|t|)
#> (Intercept) 2.413e+05 3.455e+06    0.07    0.944
#> sales       7.366e-01 1.310e-02   56.21 <2e-16 ***
#> ---
#> Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
#>
#> Residual standard error: 11150000 on 98 degrees of freedom
#> Multiple R-squared:  0.9699, Adjusted R-squared:  0.9696
#> F-statistic:  3160 on 1 and 98 DF,  p-value: < 2.2e-16
```

The predicted cost of sales for the auditee, based on the industry benchmark, can be computed as follows:

```
C_pred <- predict(fit, newdata = data.frame(sales = 298112312),
                   interval = "prediction", level = 0.90)[1]
C_pred
#> [1] 219817866
```

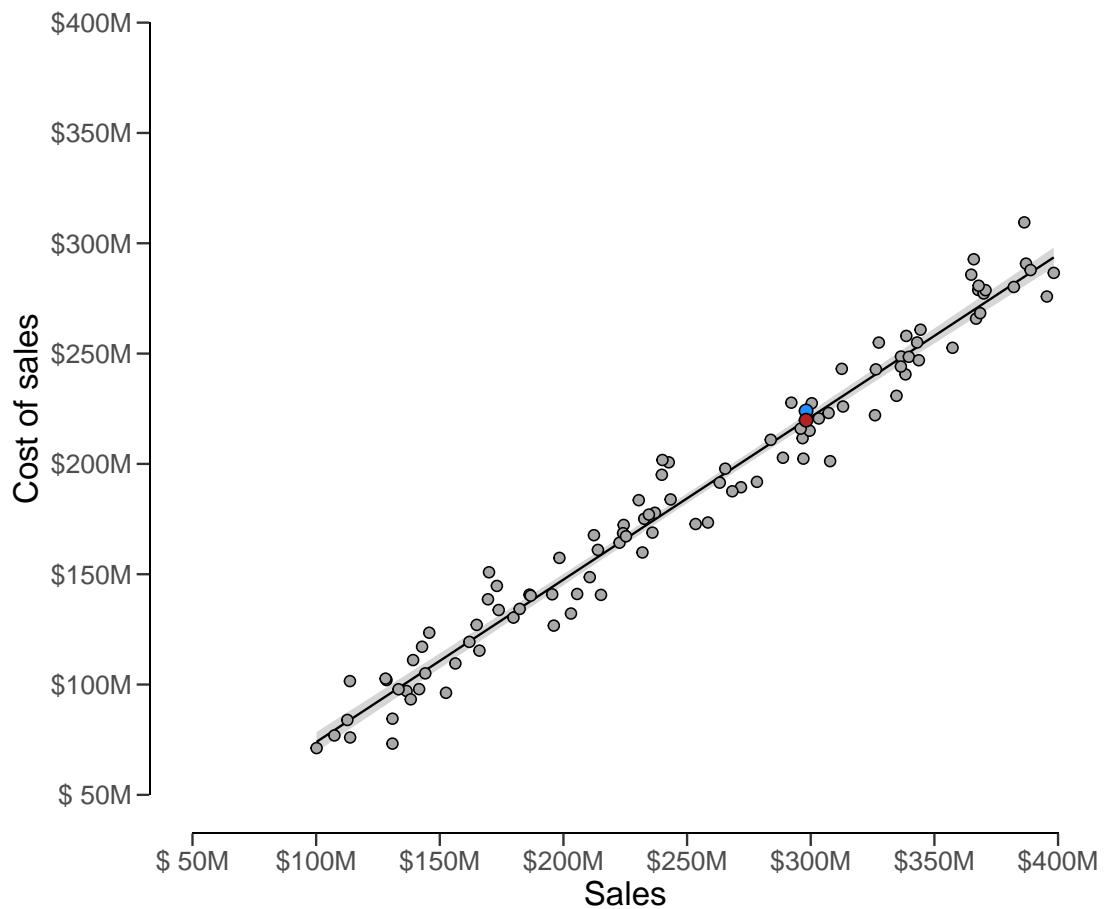
The fitted regression line and the predicted cost of sales (red dot) are visualized in the figure below:

The prior distribution can be justified by the data and the auditee's numerical prediction of the cost of sales. In this analytical procedure, the prior distribution on θ can utilize the relative error distribution from the linear regression. This relative error distribution, which is a $\text{Normal}(\mu, \sigma)$ distribution, captures the uncertainty of the prediction of the cost of sales through the use of linear regression, scaled to be a percentage of the total cost of sales. The mean μ of the prior distribution is determined by the relative deviation of the auditee's booked cost of sales when compared to the predicted cost of sales according to the benchmark data $\frac{C-\hat{C}}{C}$.

```
mu <- (C_real - C_pred) / C_real
mu
#> [1] 0.01864573
```

The standard deviation of the prior distribution is expressed through the standard deviation of the distribution of ϵ :

```
stdev <- sd(fit$residuals) / C_real
stdev
#> [1] 0.04951199
```



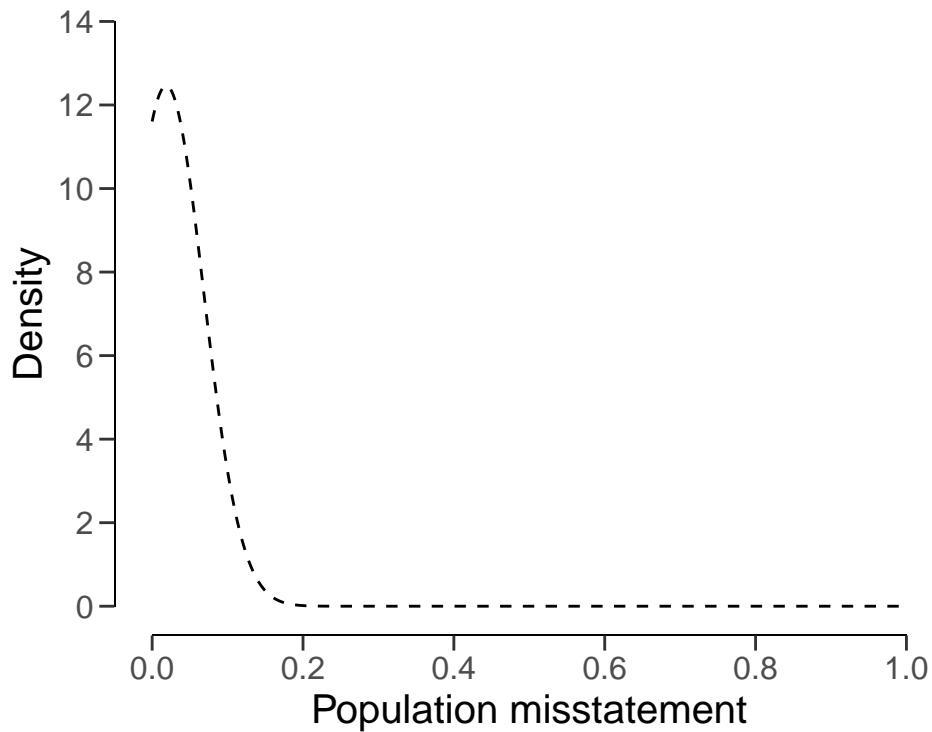
3. Planning

The $\text{Normal}(0.019, 0.05)$ prior distribution can be constructed through a call to `auditPrior()`, where the likelihood of the prior is specified as `normal`. We call the function with `method = "param"` to manually specify the parameters of the prior distribution.

```
prior <- auditPrior(method = "param", likelihood = "normal",
                      alpha = mu, beta = stdev)
summary(prior)
#>
#>  Prior Distribution Summary
#>
#>  Options:
#>    Likelihood:           normal
#>    Specifics:          = 0.0186457;   = 0.049512
#>
#>  Results:
#>    Functional form:      normal( = 0.019,   = 0.05)T[0,1]
#>    Mode:                 0.018646
#>    Mean:                 0.047096
#>    Median:               0.041335
#>    Variance:              0.0011116
#>    Skewness:              NA
#>    Information entropy (nat): -2.1306
#>    95 percent upper bound: 0.11012
#>    Precision:             0.091473
```

The specified prior distribution can be visualized using the `plot()` function.

```
plot(prior)
```



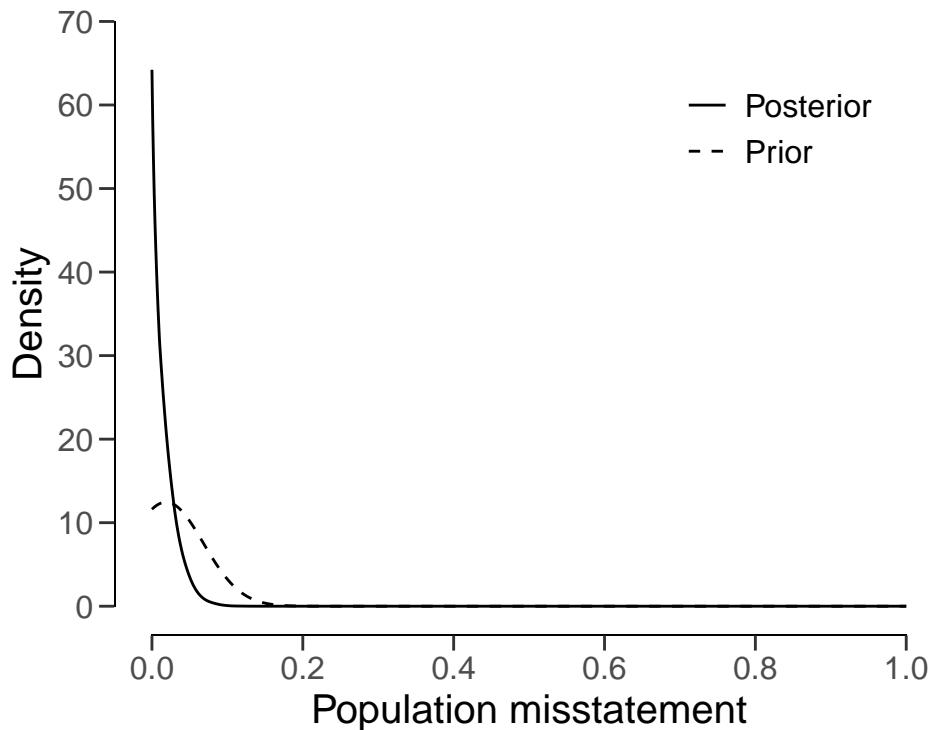
The performance materiality for this example is set to $\theta_{max} = 5\%$, and the audit risk is set to $\alpha = 0.05$, or 5%. By using this prior distribution, the required minimum sample size is 50.

```
plan <- planning(materiality = 0.05, conf.level = 0.95,
                   likelihood = "binomial", prior = prior)
plan
#>
#> Bayesian Audit Sample Planning
#>
#> minimum sample size = 50
#> sample size obtained in 51 iterations via method 'binomial' + 'prior'
```

You can inspect how the prior distribution compares to the expected posterior distribution by using the `plot()` function.

```
plot(plan)
```

3. Planning



By using a frequentist approach, the required minimum sample size is 50. Thus, by performing the analytical procedure and incorporating this information into the prior distribution, the auditor has achieved a reduction in sample size of 9 items.

```
plan <- planning(materiality = 0.05, conf.level = 0.95,
                   likelihood = "binomial")
plan
#>
#> Classical Audit Sample Planning
#>
#> minimum sample size = 59
#> sample size obtained in 60 iterations via method 'binomial'
```

3.5.3. Predictive Modeling

As a final example, we consider an example where the auditor incorporates information about the probability of misstatement obtained through a predictive analysis into the prior distribution for θ . In this example, the auditor is conducting an audit on a long-term client that they have been working with for fifteen years. Hence, the auditor has

3.5. Practical Examples

access to a historical data set called `history`, which contains the samples from all of the previous audits that have been conducted on this auditee over the past fifteen years.

```
history <- read.csv("https://github.com/koenderks/sasr/raw/master/data/ch3_history.csv",
                     colClasses = c("factor", "numeric", "numeric"))

head(history)
#>   k  ftes days
#> 1 0  2.09 1.83
#> 2 0 17.51 5.02
#> 3 0 17.29 8.47
#> 4 0 23.88 8.31
#> 5 0 15.31 8.25
#> 6 0 18.28 4.47
```

It should be noted that for all historical sample items, there are three known characteristics: whether they contained a misstatement (designated as `k`), the number of full-time equivalent employees (FTEs) who had access to that item within the internal computer systems of the auditee (designated as `ftes`), and the number of days that the item was outstanding (designated as `days`). Additionally, the `ftes` and `days` characteristics are also known for all items in the current year's population. Of course, it is unknown if any misstatements exist within the population of the current year.

```
population <- read.csv("https://github.com/koenderks/sasr/raw/master/data/ch3_population.csv")
head(population)
#>   ID bookValue ftes days
#> 1 82884    242.61  14    4
#> 2 25064    642.99  11    4
#> 3 81235    628.53  8     3
#> 4 71769    431.87  11    3
#> 5 55080    620.88  12    3
#> 6 93224    501.76  12    5
```

The objective of this analytical procedure is to forecast potential misstatements within the population of the current year. In order for the information obtained through this procedure to serve as prior knowledge in a Bayesian analysis, the procedure must yield a distribution of the probability of misstatement. Therefore, the auditor employs a machine learning technique known as Random Forest (Hastie et al., 2009) to learn the relationship between misstatements, the number of full-time equivalent employees, and the number of outstanding days in the historical data set.

```
set.seed(1)
fit <- randomForest::randomForest(formula = k ~ ftes + days, data = history)
```

The auditor specifically uses the random forest technique due to its ability to provide

3. Planning

a distribution of the misstatement probabilities. The probabilistic predictions for the unseen misstatements in the population data can be obtained by calling the `predict()` function with the argument `type = "prob"`.

```
predictions <- predict(object = fit, newdata = population, type = "prob")
```

These predictions come in a probabilistic format:

```
head(predictions)
#>      0     1
#> 1 0.974 0.026
#> 2 0.966 0.034
#> 3 0.992 0.008
#> 4 0.928 0.072
#> 5 0.776 0.224
#> 6 0.982 0.018
```

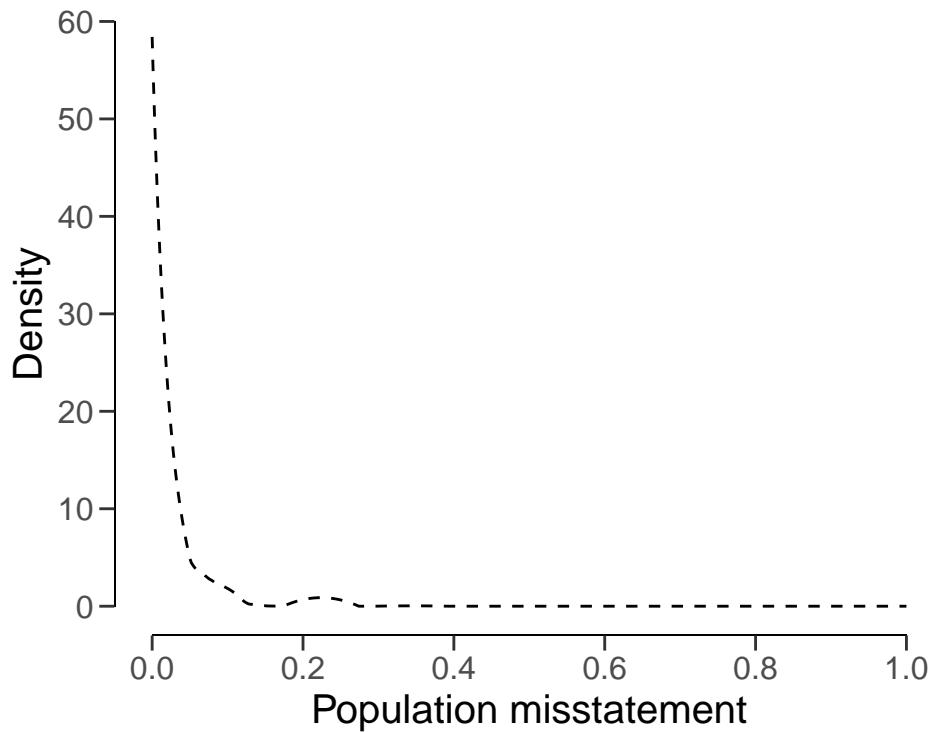
The prior distribution for θ will be based on the distribution of probabilities that each item in the population is misstated. In contrast to the previous example, this distribution is not a parametric distribution, which means we are unable to utilize any parametric priors from `jfa`. However, by providing samples of the prior distribution, `jfa` is able to construct a nonparametric prior distribution internally via the density of the samples.

The nonparametric prior distribution can be constructed through a call to `auditPrior()`, where the method to construct of the prior is specified as `nonparam`. The samples of the prior distribution can be provided through the `samples` argument. We use the second column of the `predictions` object for this.

```
prior <- auditPrior(method = "nonparam", samples = predictions[, 2])
```

The nonparametric prior distribution can be visualized using the `plot()` function.

```
plot(prior)
```



The performance materiality for this example is set to $\theta_{max} = 2\%$, and the audit risk is set to $\alpha = 0.05$, or 5%. The minimum sample size can be calculated with the command below and is 102.

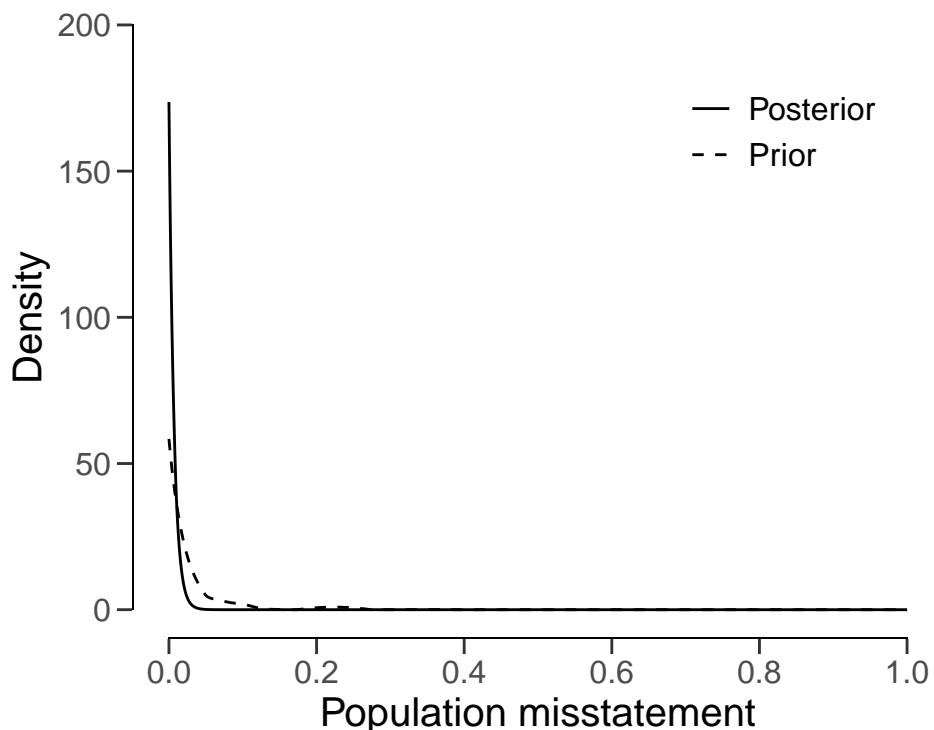
```
plan <- planning(materiality = 0.02, conf.level = 0.95,
                   likelihood = "binomial", prior = prior)
summary(plan)
#>
#> Bayesian Audit Sample Planning Summary
#>
#> Options:
#>   Confidence level:          0.95
#>   Materiality:              0.02
#>   Hypotheses:                H : θ > 0.02 vs. H : θ < 0.02
#>   Expected:                  0
#>   Likelihood:                binomial
#>   Prior distribution:        Nonparametric
#>
#> Results:
#>   Minimum sample size:       102
```

3. Planning

```
#> Tolerable errors: 0
#> Posterior distribution: Nonparametric
#> Expected most likely error: 0
#> Expected upper bound: 0.019
#> Expected precision: 0.019
#> Expected BF : 17.175
```

You can inspect how the prior distribution compares to the expected posterior distribution by using the `plot()` function.

```
plot(plan)
```



By using a frequentist approach, the required minimum sample size is 149. Thus, by performing the analytical procedure and incorporating this information into the prior distribution, the auditor has achieved a reduction in sample size of 47 items.

```
plan <- planning(materiality = 0.02, conf.level = 0.95,
                   likelihood = "binomial")
plan
```

3.5. Practical Examples

```
#>  
#> Classical Audit Sample Planning  
#>  
#> minimum sample size = 149  
#> sample size obtained in 150 iterations via method 'binomial'
```


4. Selection

i You are reading the work-in-progress first edition of Statistical Auditing with R. This chapter is currently a dumping ground of ideas, and it is incomplete.

Auditors must often evaluate balances or populations that include a large quantity of items. As it is not possible to individually examine all of these items, they must select a subset, or sample, from the total population to make a statement about a specific characteristic of the population. Several selection methodologies, which are widely accepted in the audit context, are available for this purpose. This chapter discusses the most frequently used sampling methodology for audit sampling and demonstrates how to select a sample using these methods in R.

4.1. Sampling Units

Selecting a subset from the population requires knowledge of the sampling units; physical representations of the population that needs to be audited. Generally, the auditor has to choose between two types of sampling units: individual items in the population or individual monetary units in the population. In order to perform statistical selection, the population must be divided into individual sampling units that can be assigned a probability to be included in the sample. The total collection of all sampling units which have been assigned a selection probability is called the sampling frame.

4.1.1. Items

A sampling unit for record (i.e., attributes) sampling is generally a characteristic of an item in the population. For example, suppose that you inspect a population of receipts. A possible sampling unit for record sampling can be the date of payment of the receipt. When a sampling unit (e.g., date of payment) is selected by the sampling method, the population item that corresponds to the sampled unit is included in the sample.

4. Selection

4.1.2. Monetary Units

A sampling unit for monetary unit sampling is different than a sampling unit for record sampling in that it is an individual monetary unit within an item or transaction, like an individual dollar. For example, a single sampling unit can be the 10th dollar from a specific receipt in the population. When a sampling unit (e.g., individual dollar) is selected by the sampling method, the population item that includes the sampling unit is included in the sample.

4.2. Sampling Methods

This section discusses four sampling methods that are commonly used in audit sampling. The methods that will be discussed are:

- Random sampling
- Fixed interval sampling
- Cell sampling
- Modified sieve sampling

First, let's get some notation out of the way. As discussed in Chapter 2, the population size N is defined as the total set of individual sampling units (denoted by x_i).

$$N = \{x_1, x_2, \dots, x_N\}. \quad (4.1)$$

In statistical sampling, every sampling unit x_i in the population should receive a selection probability $p(x_i)$. The purpose of the sampling method is to provide a framework to assign selection probabilities to each of the sampling units, and subsequently draw sampling units from the population until a set of size n has been created.

To illustrate how the resulting sample differs for various sampling methods, we will use the `BuildIt` data set included in the `jfa` package. These data can be loaded into R using the code below. For simplicity, we will use a sample size of $n = 10$ for all examples.

```
data(BuildIt)
n <- 10
```

4.2.1. Random Sampling

Random sampling is the most simple and straight-forward selection method. The random sampling method provides a method that allows every sampling unit in the population an equal chance of being selected, meaning that every combination of sampling units has the same probability of being selected as every other combination of the same number

of sampling units. Simply put, the algorithm draws a random selection of size n of the sampling units. Therefore, the selection probability for each sampling unit is defined as:

$$p(x) = \frac{1}{N}. \quad (4.2)$$

To make this procedure visually intuitive, the figure below provides an illustration of the random sampling method.



Figure 4.1.: Illustration of random sampling

- **Advantage(s):** The random sampling method yields an optimal random selection, with the additional advantage that the sample can be easily extended by applying the same method again.
- **Disadvantages:** Because the selection probabilities are equal for all sampling units there is no guarantee that items with a large monetary value in the population will be included in the sample.

4.2.1.1. Record Sampling

Random sampling can easily be coded in base R. First, we have to get a vector of of the possible items (rows) in the population that can be selected. When we are performing record sampling, we can simply use R's build in `sample()` function to draw a random sample from a vector `1:nrow(BuildIt)` representing the row indices of the items and store the result in a variable `items`.

```
set.seed(1)
items <- sample(1:nrow(BuildIt), size = n, replace = FALSE)
items
#> [1] 1017 679 2177 930 1533 471 2347 270 1211 3379
```

You can then select the sample from the population using the selected indices stored in `items`.

```
BuildIt[items, ]
#>           ID bookValue auditValue
#> 1017      50755     618.24     618.24
```

4. Selection

```
#> 679 20237 669.75 669.75
#> 2177 9517 454.02 454.02
#> 930 85674 257.82 257.82
#> 1533 31051 308.53 308.53
#> 471 84375 824.66 824.66
#> 2347 75616 623.70 623.70
#> 270 82033 352.75 352.75
#> 1211 12877 52.89 52.89
#> 3379 85322 330.24 330.24
```

The sample can be reproduced in **jfa** via the **selection()** function. This function takes as input the population data, the sample size, and the characteristics of the sampling method. The argument **units** allows you to specify that you want to use record sampling (**units = "items"**), while the **method** argument enables you to specify that you are performing random sampling (**method = 'random'**).

```
set.seed(1)
result <- selection(data = BuildIt, size = n, units = "items",
                      method = "random")
result$sample
#>      row times     ID bookValue auditValue
#> 1  1017    1 50755   618.24   618.24
#> 2   679    1 20237   669.75   669.75
#> 3  2177    1 9517    454.02   454.02
#> 4   930    1 85674   257.82   257.82
#> 5  1533    1 31051   308.53   308.53
#> 6   471    1 84375   824.66   824.66
#> 7  2347    1 75616   623.70   623.70
#> 8   270    1 82033   352.75   352.75
#> 9  1211    1 12877   52.89    52.89
#> 10 3379    1 85322   330.24   330.24
```

4.2.1.2. Monetary Unit Sampling

When we are performing record sampling, we have to consider that each item in the population consists of multiple smaller items (i.e., the monetary units), which means that items with a higher book value should get a higher probability of being selected. The **sample()** function facilitates weighted selection via the **prob** argument, which takes a vector of values and, using normalization, computes the weights for selection. The call below is similar to before, but in this case we use the book values in the column **bookValues** of the data set to weigh the items and store the result in a variable **items**.

```

set.seed(1)
items <- sample(1:nrow(BuildIt), size = n, replace = FALSE,
                 prob = BuildIt$bookValue)
items
#> [1] 2174 2928 1627 700 147 3056 3118 2045 1311 716

```

You can then select the sample from the population using the selected indices stored in `items`.

```

BuildIt[items, ]
#>      ID bookValue auditValue
#> 2174 90260    625.98    625.98
#> 2928 68595    548.21    548.21
#> 1627 98301    429.07    429.07
#> 700 29683    239.26    239.26
#> 147 72906    677.62    677.62
#> 3056 86317    246.22    246.22
#> 3118 14548    204.63    204.63
#> 2045 45416    381.05    381.05
#> 1311 91955    398.96    398.96
#> 716 12815    873.43    873.43

```

The sample can be reproduced in `jfa` via the `selection()` function. The argument `units` allows you to specify that you want to use monetary unit sampling (`units = "values"`), while the `method` argument enables you to specify that you are performing random sampling (`method = 'random'`). Note that you should provide the name of the column in the data that contains the monetary units via the `values` argument.

```

set.seed(1)
result <- selection(data = BuildIt, size = n, units = "values",
                      method = "random", values = "bookValue")
result$sample
#>      row times      ID bookValue auditValue
#> 1 2174     1 90260    625.98    625.98
#> 2 2928     1 68595    548.21    548.21
#> 3 1627     1 98301    429.07    429.07
#> 4 700      1 29683    239.26    239.26
#> 5 147      1 72906    677.62    677.62
#> 6 3056     1 86317    246.22    246.22
#> 7 3118     1 14548    204.63    204.63
#> 8 2045     1 45416    381.05    381.05
#> 9 1311     1 91955    398.96    398.96
#> 10 716     1 12815    873.43    873.43

```

4. Selection

4.2.2. Fixed Interval Sampling

Fixed interval sampling is a method designed for yielding representative samples from monetary populations. The algorithm determines a uniform interval on the (optionally ranked) sampling units. Next, a starting point is handpicked or randomly selected in the first interval and a sampling unit is selected throughout the population at each of the uniform intervals from the starting point. For example, if the interval has a width of 10 sampling units and sampling unit number 5 is chosen as the starting point, the sampling units 5, 15, 25, etc. are selected to be included in the sample.

The number of required intervals I can be determined by dividing the number of sampling units in the population by the required sample size:

$$I = \frac{N}{n}, \quad (4.3)$$

in which n is the required sample size and N is the total number of sampling units in the population.

If the space between the selected sampling units is equal, the selection probability for each sampling unit is theoretically defined as:

$$p(x) = \frac{1}{I}, \quad (4.4)$$

with the property that the space between selected units i is the same as the interval I , see the figure below. However, in practice the selection is deterministic and completely depends on the chosen starting points (using `start`).

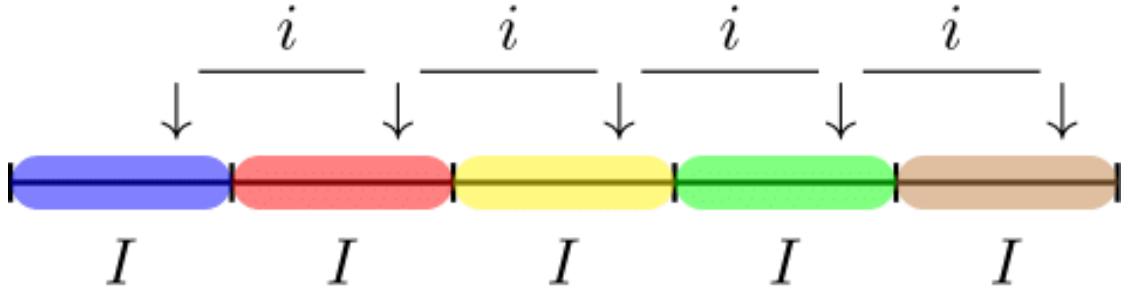


Figure 4.2.: Illustration of fixed interval sampling

The fixed interval method yields a sample that allows every sampling unit in the population an equal chance of being selected. However, the fixed interval method has the property that all items in the population with a monetary value larger than the interval I have an selection probability of one because one of these items' sampling units are always selected from the interval. Note that, if the population is arranged randomly

with respect to its deviation pattern, fixed interval sampling is equivalent to random selection.

- **Advantage(s):** The advantage of the fixed interval sampling method is that it is often simple to understand and fast to perform. Another advantage is that, in monetary unit sampling, all items that are greater than the calculated interval will be included in the sample. In record sampling, since units can be ranked on the basis of value, there is also a guarantee that some large items will be in the sample.
- **Disadvantage(s):** A pattern in the population can coincide with the selected interval, rendering the sample less representative. What is sometimes seen as an added complication for this method is that the sample is hard to extend after drawing the initial sample. This is due to the chance of selecting the same sampling unit. However, by removing the already selected sampling units from the population and redrawing the intervals this problem can be efficiently solved.

4.2.2.1. Record Sampling

To code fixed interval sampling in a record sampling context, we first have to compute the size of the interval we are working with. This is computed by dividing the number of items in the population by the desired sample size n . Suppose the auditor wants to select a sample of 10 items, then the interval is computed by:

```
interval <- nrow(BuildIt) / n
```

Next, we have to determine the starting point. We are going to take the fifth unit in each interval in this case.

```
start <- 5
```

To find which rows are part of the sample, we execute the following code:

```
items <- floor(start + interval * 0:(n - 1))
```

You can then select the sample from the population using the selected indices stored in `items`.

```
BuildIt[items, ]
#>      ID bookValue auditValue
#> 5   55080    620.88   620.88
#> 355 27934    749.38   749.38
#> 705 21900    919.00   919.00
#> 1055 66675   384.27   384.27
#> 1405 13472   360.05   360.05
```

4. Selection

```
#> 1755 61607    389.75    389.75  
#> 2105 68519    354.71    354.71  
#> 2455 91983    467.72    467.72  
#> 2805 25646    420.80    420.80  
#> 3155 94955    248.77    248.77
```

The sample can be reproduced in **jfa** via the **selection()** function. The argument **units** allows you to specify that you want to use record sampling (**units = "items"**), while the **method** argument enables you to specify that you are performing fixed interval sampling (**method = 'interval'**). Note that, by default, the first sampling unit from each interval is selected. However, this can be changed by setting the argument **start** to a different value.

```
result <- selection(data = BuildIt, size = n, units = "items",  
                      method = "interval", start = start)  
result$sample  
#>   row times     ID bookValue auditValue  
#> 1    5    1 55080    620.88    620.88  
#> 2   355    1 27934    749.38    749.38  
#> 3   705    1 21900    919.00    919.00  
#> 4  1055    1 66675    384.27    384.27  
#> 5  1405    1 13472    360.05    360.05  
#> 6  1755    1 61607    389.75    389.75  
#> 7  2105    1 68519    354.71    354.71  
#> 8  2455    1 91983    467.72    467.72  
#> 9  2805    1 25646    420.80    420.80  
#> 10 3155    1 94955    248.77    248.77
```

4.2.2.2. Monetary Unit Sampling

In monetary unit sampling, the only difference is that we are computing the interval on the basis of the booked values in the column **bookValue** of the data set. In this case, the starting point **start = 5** determines which monetary unit from each interval is selected.

```
interval <- sum(BuildIt$bookValue) / n
```

To find which units are part of the sample, we execute the following code:

```
units <- floor(start + interval * 0:(n - 1))
```

To obtain which items are part of the sample, we can run the following for loop. Note that this does not take into account whether the book values contain negative values,

4.2. Sampling Methods

which should not be included in the cumulative sum below.

```
all_units <- ifelse(BuildIt$bookValue < 0, 0, BuildIt$bookValue)
all_items <- 1:nrow(BuildIt)
items <- numeric(n)
for (i in 1:n) {
  item <- which(units[i] <= cumsum(all_units))[1]
  items[i] <- all_items[item]
}
```

You can then select the sample from the population using the selected indices stored in `items`.

```
BuildIt[items, ]
#>           ID bookValue auditValue
#> 1    82884    242.61    242.61
#> 358 20711    610.88    610.88
#> 715 99012    313.75    313.75
#> 1081 65319    502.54    201.02
#> 1421 88454    856.28    856.28
#> 1774 87258    157.68    157.68
#> 2103 48652    497.21    497.21
#> 2435 37248    1041.44   1041.44
#> 2787 10925    377.10    377.10
#> 3152 71832    1001.82   1001.82
```

The sample can be reproduced in `jfa` via the `selection()` function. The argument `units` allows you to specify that you want to use monetary unit sampling (`units = "values"`), while the `method` argument enables you to specify that you are performing fixed interval sampling (`method = 'interval'`). Note that you should provide the name of the column in the data that contains the monetary units via the `values` argument.

```
result <- selection(data = BuildIt, size = n, units = "values",
                      method = "interval", values = "bookValue", start = start)
result$sample
#>      row times     ID bookValue auditValue
#> 1      1     1 82884    242.61    242.61
#> 2     358     1 20711    610.88    610.88
#> 3      715     1 99012    313.75    313.75
#> 4     1081     1 65319    502.54    201.02
#> 5     1421     1 88454    856.28    856.28
#> 6     1774     1 87258    157.68    157.68
#> 7     2103     1 48652    497.21    497.21
```

4. Selection

```
#> 8 2435      1 37248    1041.44    1041.44
#> 9 2787      1 10925     377.10     377.10
#> 10 3152     1 71832    1001.82    1001.82
```

4.2.3. Cell Sampling

The cell sampling method divides the (optionally ranked) population into a set of intervals I that are computed through the previously given equations. Within each interval, a sampling unit is selected by randomly drawing a number between 1 and the interval range I . This causes the space i between the sampling units to vary.

Like in the fixed interval sampling method, the selection probability for each sampling unit is defined as:

$$p(x) = \frac{1}{I}. \quad (4.5)$$

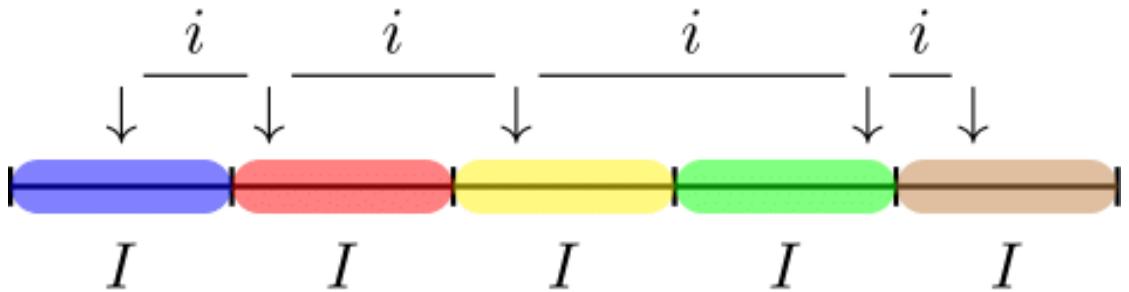


Figure 4.3.: Illustration of cell sampling

The cell sampling method has the property that all items in the population with a monetary value larger than twice the interval I have a selection probability of one.

- **Advantage(s):** More sets of samples are possible than in fixed interval sampling, as there is no systematic interval i to determine the selections. It is argued that the cell sampling algorithm offers a solution to the pattern problem in fixed interval sampling.
- **Disadvantage(s):** A disadvantage of this sampling method is that not all items in the population with a monetary value larger than the interval have a selection probability of one. Besides, population items can be in two adjacent cells, thereby creating the possibility that an item is included in the sample twice.

4.2.3.1. Record Sampling

To code cell sampling in a record sampling context, we again have to compute the size of the interval we are working with:

```
interval <- nrow(BuildIt) / n
```

Next, we have to randomly determine which items are going to be selected in each interval.

```
set.seed(1)
starts <- floor(runif(n, 0, interval))
```

To find which rows are part of the sample, we execute the following code:

```
items <- floor(starts + interval * 0:(n - 1))
```

You can then select the sample from the population using the selected indices stored in `items`.

```
BuildIt[items, ]
#>      ID bookValue auditValue
#> 92    75133    355.16    355.16
#> 480   81037    456.27    456.27
#> 900   1730     449.87    449.87
#> 1367  36587    282.32    282.32
#> 1470  10305    648.70    648.70
#> 2064  96344    268.94    268.94
#> 2430  60885    493.77    493.77
#> 2681  60935    312.98    312.98
#> 3020  8716     450.76    450.76
#> 3171  61036    387.67    387.67
```

The sample can be reproduced in `jfa` via the `selection()` function. The argument `units` allows you to specify that you want to use record sampling (`units = "items"`), while the `method` argument enables you to specify that you are performing cell sampling (`method = 'cell'`).

```
set.seed(1)
result <- selection(data = BuildIt, size = n, units = "items",
                      method = "cell")
result$sample
#>      row times      ID bookValue auditValue
```

4. Selection

```
#> 1    92     1 75133   355.16   355.16
#> 2    480    1 81037   456.27   456.27
#> 3    900    1 1730    449.87   449.87
#> 4   1367    1 36587   282.32   282.32
#> 5   1470    1 10305   648.70   648.70
#> 6   2064    1 96344   268.94   268.94
#> 7   2430    1 60885   493.77   493.77
#> 8   2681    1 60935   312.98   312.98
#> 9   3020    1 8716    450.76   450.76
#> 10  3171    1 61036   387.67   387.67
```

4.2.3.2. Monetary Unit Sampling

In monetary unit sampling, the only difference is that we are computing the interval on the basis of the booked values in the column `bookValue` of the data set. In this case, the starting points `start` determines which monetary unit from each interval is selected.

```
interval <- sum(BuildIt$bookValue) / n
```

To obtain which items are part of the sample, we can run the following for loop. Note that this does not take into account whether the book values contain negative values, which should not be included in the cumulative sum below.

```
set.seed(1)
all_units <- ifelse(BuildIt$bookValue < 0, 0, BuildIt$bookValue)
all_items <- 1:nrow(BuildIt)
intervals <- 0:n * interval
items <- numeric(n)
for (i in 1:n) {
  unit <- stats::runif(1, intervals[i], intervals[i + 1])
  item <- which(unit <= cumsum(all_units))[1]
  items[i] <- all_items[item]
}
```

You can then select the sample from the population using the selected indices stored in `items`.

```
BuildIt[items, ]
#>           ID bookValue auditValue
#> 95     15009   415.60   415.60
#> 486    79093   635.85   635.85
#> 931    28025   429.14   429.14
```

```
#> 1387 56444    296.37    296.37
#> 1492 81443    543.80    543.80
#> 2074 14196    270.45    270.45
#> 2418 87743    347.99    347.99
#> 2660 23927    454.81    454.81
#> 3024 78925    251.44    251.44
#> 3172 18286    450.57    450.57
```

The sample can be reproduced in **jfa** via the `selection()` function. The argument `units` allows you to specify that you want to use monetary unit sampling (`units = "values"`), while the `method` argument enables you to specify that you are performing cell sampling (`method = 'cell'`). Note that you should provide the name of the column in the data that contains the monetary units via the `values` argument.

```
set.seed(1)
result <- selection(data = BuildIt, size = n, units = "values",
                      method = "cell", values = "bookValue")

result$sample
#>      row times     ID bookValue auditValue
#> 1      95     1 15009    415.60    415.60
#> 2     486     1 79093    635.85    635.85
#> 3     931     1 28025    429.14    429.14
#> 4    1387     1 56444    296.37    296.37
#> 5    1492     1 81443    543.80    543.80
#> 6    2074     1 14196    270.45    270.45
#> 7    2418     1 87743    347.99    347.99
#> 8    2660     1 23927    454.81    454.81
#> 9    3024     1 78925    251.44    251.44
#> 10   3172     1 18286    450.57    450.57
```

4.2.4. Modified Sieve Sampling

The fourth option for the sampling method is modified sieve sampling (Hoogduin, Hall, & Tsay, 2010). The algorithm starts by selecting a standard uniform random number R_i between 0 and 1 for each item in the population. Next, the sieve ratio:

$$S_i = \frac{Y_i}{R_i} \quad (4.6)$$

is computed for each item by dividing the book value of that item by the random number. Lastly, the items in the population are sorted by their sieve ratio S (in decreasing order) and the top n items are selected for inspection. In contrast to the classical sieve sampling

4. Selection

method (Rietveld, 1978), the modified sieve sampling method provides precise control over sample sizes.

4.2.4.1. Monetary Unit Sampling

```
set.seed(1)
all_units <- ifelse(BuildIt$bookValue < 0, 0, BuildIt$bookValue)
all_items <- 1:nrow(BuildIt)
ri <- all_units / stats::runif(length(all_items), 0, 1)
items <- all_items[order(-ri)]
items <- items[1:n]
```

You can then select the sample from the population using the selected indices stored in `items`.

```
BuildIt[items, ]
#>      ID bookValue auditValue
#> 2329 29919    681.10    681.10
#> 2883 59402    279.29    279.29
#> 1949 56012    581.22    581.22
#> 3065 47482    621.73    621.73
#> 1072 79901    789.97    789.97
#> 488  50811    651.35    651.35
#> 1916 53565    266.37    266.37
#> 463   65768    480.89    480.89
#> 1311 91955    398.96    398.96
#> 2895 8688     492.02    492.02
```

The sample can be reproduced in `jfa` via the `selection()` function. The argument `units` allows you to specify that you want to use monetary unit sampling (`units = "values"`), while the `method` argument enables you to specify that you are performing modified sieve sampling (`method = 'sieve'`). Note that you should provide the name of the column in the data that contains the monetary units via the `values` argument.

```
set.seed(1)
result <- selection(data = BuildIt, size = n, units = "values",
                      method = "sieve", values = "bookValue")
result$sample
#>      row times      ID bookValue auditValue
#> 1  2329     1 29919    681.10    681.10
#> 2  2883     1 59402    279.29    279.29
#> 3  1949     1 56012    581.22    581.22
```

4.3. Ordering or Randomizing the Population

```
#> 4 3065     1 47482    621.73    621.73
#> 5 1072     1 79901    789.97    789.97
#> 6 488      1 50811    651.35    651.35
#> 7 1916     1 53565    266.37    266.37
#> 8 463      1 65768    480.89    480.89
#> 9 1311     1 91955    398.96    398.96
#> 10 2895    1 8688     492.02    492.02
```

4.3. Ordering or Randomizing the Population

The `selection()` function from the `jfa` package has three additional arguments which you can use to preprocess your population before selection. These arguments are `order`, `decreasing` and `randomize`.

First, the `order` argument takes as input a column name in `data` which determines the order of the population. For example, you can order the population from lowest book value to highest book value before engaging in the selection. In this case, you should use the `decreasing = FALSE` (its default value) argument.

```
set.seed(1)
result <- selection(data = BuildIt, size = n, units = "values",
                      values = "bookValue", order = "bookValue")
result$sample
#>   row times     ID bookValue auditValue
#> 1 2662     1 30568    14.47    14.47
#> 2 2307     1 95785   244.49   244.49
#> 3 1639     1 39570   307.54   307.54
#> 4 2318     1 11857   360.57   360.57
#> 5 3409     1 18796   414.75   414.75
#> 6 1158     1 28985   468.57   468.57
#> 7 2892     1 27448   530.73   530.73
#> 8 2759     1 16442   596.34   596.34
#> 9 3112     1 62440   678.36   678.36
#> 10 1820    1 82309   796.31   796.31
```

Next, the `randomize` argument can be used to randomly shuffle the items in the population before selection. For example, you can randomly shuffle the population before engaging in the selection using `randomize = TRUE`.

```
set.seed(1)
result <- selection(data = BuildIt, size = n, units = "values",
                      values = "bookValue", randomize = TRUE)
```

4. Selection

```
result$sample
#>      row times     ID bookValue auditValue
#> 1  1017     1 50755    618.24    618.24
#> 2  2097     1 53186    642.34    642.34
#> 3  2508     1 55666    426.33    426.33
#> 4   779     1 82046    569.21    227.68
#> 5  2255     1 63601    234.12    234.12
#> 6  2072     1 92569    517.14    517.14
#> 7  1938     1 35525    366.94    366.94
#> 8   595     1 30750    338.36    338.36
#> 9  3207     1 74858    884.51    884.51
#> 10 2981     1 24682    302.23    302.23
```

5. Evaluation

i You are reading the work-in-progress first edition of Statistical Auditing with R. This chapter is currently a dumping ground of ideas, and it is incomplete.

Evaluating a sample is the final stage in the audit sampling workflow. In this stage, the auditor weighs the audit evidence obtained from the sample together with the audit evidence obtained in earlier stages of the audit to determine their opinion about the population.

Like planning a sample, evaluating a sample requires knowledge of the conditions that lead to acceptance or rejection of the population (i.e., the sampling objectives). A quick reminder: Sampling objectives can be classified into one of the following two categories:

- **Hypothesis testing:** The goal of the sample is to obtain evidence for or against the claim that the misstatement in the population is lower than a given value (i.e., the performance materiality).
- **Estimation:** The goal of the sample is to obtain an accurate estimate of the misstatement in the population (with a minimum precision).

A non-stratified audit sample does not involve dividing the population into subgroups. Here, the auditor selects a random sample from the entire population without considering any specific characteristics of the population.

For example, in an audit of a company's inventory, the auditor may simply select a random sample of items from the entire inventory without dividing it into subgroups based on characteristics such as location or type of item. Another example of such a situation would be where the auditor is auditing the general ledger of a small business.

Naturally, non-stratified sampling is simpler than stratified sampling and can be used when the population is considered homogenous or the auditor does not need to consider differences between subgroups (e.g., strata).

5.1. Classical Evaluation

Classical hypothesis testing uses the p -value to make a decision about whether to reject the hypothesis H_0 or not. As an example, consider that an auditor wants to verify

5. Evaluation



Figure 5.1.: When evaluating a sample with respect to a specific hypothesis, the auditor must consider the evidence in favour as well as the evidence against that hypothesis. Image available under a CC-BY-NC 4.0 license.

whether the population contains less than 5 percent misstatement, implying the hypotheses $H_1 : \theta < 0.05$ and $H_0 : \theta \geq 0.05$. They have taken a sample of 100 items, of which 1 contained an error. They set the significance level for the p -value to 0.05, implying that $p < 0.05$ will be enough to reject the hypothesis H_0 . The call below evaluates the sample using a classical non-stratified evaluation procedure.

```
evaluation(materiality = 0.05, x = 1, n = 100)
#>
#> Classical Audit Sample Evaluation
#>
#> data: 1 and 100
#> number of errors = 1, number of samples = 100, taint = 1, p-value =
#> 0.040428
#> alternative hypothesis: true misstatement rate is less than 0.05
#> 95 percent confidence interval:
#> 0.00000000 0.04743865
#> most likely estimate:
#> 0.01
#> results obtained via method 'poisson'
```

The output shows that the most likely error in the population is estimated to be $1 / 100 = 1\%$ and that the 95% (one-sided) confidence interval ranges from 0% to 4.74%. The output also shows that the p -value is lower than 0.05 implying that the hypothesis H_0

can be rejected. Hence, the auditor is able to conclude that the sample provides sufficient evidence to state with reasonable assurance that the population does not contain material misstatement.

5.2. Bayesian Evaluation

Bayesian hypothesis testing uses the Bayes factor, BF_{10} or BF_{01} , to make a statement about the evidence provided by the sample in support for one of the two hypotheses H_1 or H_0 . As an example of how to interpret the Bayes factor, the value of $BF_{10} = 10$ (provided by the `evaluation()` function) can be interpreted as: *the data are 10 times more likely to have occurred under the hypothesis H_1 than under the hypothesis H_0 .* $BF_{10} > 1$ indicates evidence in favor of H_1 and against H_0 , while $BF_{10} < 1$ indicates evidence in favor of H_0 and against H_1 . The `evaluation()` function returns the value for BF_{10} , but BF_{01} can be computed as $\frac{1}{BF_{10}}$.

Consider the previous example of an auditor who wants to verify whether the population contains less than 5 percent misstatement, implying the hypotheses $H_1 : \theta < 0.05$ and $H_0 : \theta \geq 0.05$. They have taken a sample of 100 items, of which 1 was found to contain a misstatement. The prior distribution is assumed to be a default `beta(1,1)` prior. The call below evaluates the sample using a Bayesian non-stratified evaluation procedure.

```
prior <- auditPrior(materiality = 0.05, method = "default", likelihood = "binomial")
evaluation(materiality = 0.05, x = 1, n = 100, prior = prior)
#>
#> Bayesian Audit Sample Evaluation
#>
#> data: 1 and 100
#> number of errors = 1, number of samples = 100, taint = 1, BF =
#> 515.86
#> alternative hypothesis: true misstatement rate is less than 0.05
#> 95 percent credible interval:
#> 0.00000000 0.04610735
#> most likely estimate:
#> 0.01
#> results obtained via method 'binomial' + 'prior'
```

The output shows that the most likely error in the population is estimated to be $1 / 100 = 1\%$ and that the 95% (one-sided) credible interval ranges from 0% to 4.61%. The small difference between the classical and default Bayesian results is due to the prior distribution, which must be proper in order to calculate a Bayes factor (classical results can be emulated by constructing a prior with `method = "strict"` in the `auditPrior()` function). The Bayes factor in this case is shown to be $BF_{10} = 515$, meaning that the

5. Evaluation

data from the sample are about 515 times more likely to occur under the hypothesis of tolerable misstatement than under the hypothesis of material misstatement.

Note that this is a very high Bayes factor for the little data that is observed. That is because the Bayes factor is dependent on the prior distribution for θ . As a rule of thumb, when the prior distribution is highly conservative (as with `method = 'default'`) with respect to the hypothesis of tolerable misstatement, the Bayes factor tends to over quantify the evidence in favor of this hypothesis. You can mitigate this dependency by using a prior distribution that is impartial with respect to the hypotheses via `method = "impartial"` in the `auditPrior()` function (Derks et al., 2022).

```
prior <- auditPrior(materiality = 0.05, method = "impartial", likelihood = "binom"
evaluation(materiality = 0.05, x = 1, n = 100, prior = prior)
#>
#> Bayesian Audit Sample Evaluation
#>
#> data: 1 and 100
#> number of errors = 1, number of samples = 100, taint = 1, BF =
#> 47.435
#> alternative hypothesis: true misstatement rate is less than 0.05
#> 95 percent credible interval:
#> 0.00000000 0.04110834
#> most likely estimate:
#> 0.0088878
#> results obtained via method 'binomial' + 'prior'
```

The output shows that $BF_{10} = 47$, implying that under the assumption of impartiality there is strong evidence for H_1 , the hypothesis that the population contains misstatements lower than 5 percent of the population (tolerable misstatement). Since the two prior distributions both resulted in convincing Bayes factors, the results can be considered robust to the choice of prior distribution. Hence, the auditor is able to conclude that the sample provides convincing evidence to state that the population does not contain material misstatement.

5.3. Evaluation using Data

For this example, we take the `allowances` data set that comes with the package. This data set contains 3500 financial statement line items, each with a booked value `bookValue` and, for illustrative purposes, and audited (true) value `auditValue`. Since the focus of this vignette is the evaluation stage in the audit, the sample is already indicated in the data set. The performance materiality in this example is set to 5%.

```
data(allowances)
head(allowances)
#>   item branch bookValue auditValue times
#> 1     1      12     1600     1600     1
#> 2     2      12     1625       NA     0
#> 3     3      12     1775       NA     0
#> 4     4      12     1250     1250     1
#> 5     5      12     1400       NA     0
#> 6     6      12     1190       NA     0
```

Evaluating a non-stratified sample using data requires specification of the `data`, `values` and `values.audit` arguments. The input for these arguments is the name of the specific column in `data`.

5.3.1. Classical Evaluation

The call below evaluates the `allowances` sample using a classical non-stratified evaluation procedure.

```
x <- evaluation(
  materiality = 0.05, data = allowances,
  values = "bookValue", values.audit = "auditValue", times = "times"
)
summary(x)
#>
#>   Classical Audit Sample Evaluation Summary
#>
#> Options:
#>   Confidence level:                 0.95
#>   Materiality:                     0.05
#>   Hypotheses:                      H : θ >= 0.05 vs. H : θ < 0.05
#>   Method:                          poisson
#>
#> Data:
#>   Sample size:                     1604
#>   Number of errors:                401
#>   Sum of taints:                  252.9281046
#>
#> Results:
#>   Most likely error:              0.15769
#>   95 percent confidence interval: [0, 0.175]
#>   Precision:                      0.017311
```

5. Evaluation

```
#>   p-value:          1
```

In this case, the output shows that the estimate of the misstatement in the population is 15.77%, with the 95% (one-sided) confidence interval ranging from 0% to 17.5%.

5.3.2. Bayesian evaluation

The call below evaluates the `allowances` sample using a Bayesian non-stratified evaluation procedure.

```
x <- evaluation(
  materiality = 0.05, data = allowances, prior = TRUE,
  values = "bookValue", values.audit = "auditValue", times = "times"
)
summary(x)
#>
#> Bayesian Audit Sample Evaluation Summary
#>
#> Options:
#>   Confidence level:           0.95
#>   Materiality:              0.05
#>   Hypotheses:                H : θ > 0.05 vs. H : θ < 0.05
#>   Method:                    poisson
#>   Prior distribution:        gamma(  = 1,  = 1)
#>
#> Data:
#>   Sample size:               1604
#>   Number of errors:         401
#>   Sum of taints:            252.9281046
#>
#> Results:
#>   Posterior distribution:    gamma(  = 253.928,  = 1605)
#>   Most likely error:        0.15759
#>   95 percent credible interval: [0, 0.17489]
#>   Precision:                 0.0173
#>   BF :                      0
```

The output shows that the estimate of the misstatement in the population is 15.76%, with the 95% (one-sided) credible interval ranging from 0% to 17.49%.

6. Stratified Evaluation

i You are reading the work-in-progress first edition of Statistical Auditing with R. This chapter is currently a dumping ground of ideas, and it is incomplete.

In an audit context, stratified sampling can be used to select a sample of transactions from different departments, locations, or business units to ensure that the sample is representative of the population. This means that there is a hierarchical structure in the data (strata are nested in the population) and the statistical model should ideally reflect this.

In general, there are three approaches to evaluating a stratified sample: no pooling, complete pooling, and partial pooling. No pooling assumes no similarities between strata, which means that all strata are analyzed independently. Complete pooling assumes no difference between strata, which means that all data is aggregated and analyzed as a whole. Finally, partial pooling assumes differences and similarities between strata, which means that information can be shared between strata. Partial pooling (i.e., hierarchical modeling) is a powerful technique that can result in more efficient population and stratum estimates.

For example, if you want to audit the expense claims of a large organization, you can stratify the population based on the department, location, or business unit to ensure that all departments are represented accordingly in the sample. Another example of such a situation would be a group audit where the audited organization consists of different components or branches. Stratification is relevant for the group auditor if they must form an opinion on the group as a whole because they must aggregate the samples taken by the component auditors.

As a data example, consider the `retailer` data set that comes with the package. The organization in question consists of 20 branches across the country. In each of the 20 strata, a component auditor has taken a statistical sample and reported the outcomes to the group auditor.

```
data(retailer)
head(retailer)
#>   stratum items samples errors
#> 1       1  5000     300      21
#> 2       2  5000     300      16
```

6. Stratified Evaluation



Figure 6.1.: There are three approaches to evaluating a stratified audit sample: no pooling, complete pooling and partial pooling. Image available under a CC-BY-NC 4.0 license.

```
#> 3      3  5000    300    15  
#> 4      4  5000    300    14  
#> 5      5  5000    300    16  
#> 6      6  5000    150     5
```

The number of units per stratum in the population can be provided with `N.units` to weigh the stratum estimates to determine population estimate. This is called poststratification. If `N.units` is not specified, each stratum is assumed to be equally represented in the population.

6.1. No pooling

No pooling (`pooling = "none"`, default) assumes no similarities between strata. This means that the prior distribution specified through `prior` is applied independently for each stratum. This allows for independent estimates for the misstatement in each stratum but also results in a relatively high uncertainty in the population estimate. The call below evaluates the sample using a Bayesian stratified evaluation procedure, in which the stratum estimates are poststratified to arrive at the population estimate.

```

set.seed(1) # Important because the posterior distribution is determined via sampling
result_np <- evaluation(
  materiality = 0.05, method = "binomial", prior = TRUE,
  n = retailer$samples, x = retailer$errors, N.units = retailer$items,
  alternative = "two.sided", pooling = "none"
)
summary(result_np)
#>
#> Bayesian Audit Sample Evaluation Summary
#>
#> Options:
#>   Confidence level:          0.95
#>   Population size:          144000
#>   Materiality:              0.05
#>   Hypotheses:                H : θ = 0.05 vs. H : θ ≠ 0.05
#>   Method:                   binomial
#>   Prior distribution:       Nonparametric
#>
#> Data:
#>   Sample size:               2575
#>   Number of errors:          115
#>   Sum of taints:             115
#>
#> Results:
#>   Posterior distribution:    Nonparametric
#>   Most likely error:        0.0598
#>   95 percent credible interval: [0.042763, 0.082201]
#>   Precision:                 0.022401
#>   BF :                       0
#>
#> Strata (20):
#>   N  n  x  t      mle      lb      ub precision
#> 1 5000 300 21 21 0.07000 0.04637 0.10467  0.03467
#> 2 5000 300 16 16 0.05333 0.03324 0.08489  0.03156
#> 3 5000 300 15 15 0.05000 0.03069 0.08086  0.03086
#> 4 5000 300 14 14 0.04667 0.02816 0.07681  0.03014
#> 5 5000 300 16 16 0.05333 0.03324 0.08489  0.03156
#> 6 5000 150  5  5 0.03333 0.01472 0.07558  0.04225
#> 7 5000 150  4  4 0.02667 0.01084 0.06643  0.03977
#> 8 5000 150  3  3 0.02000 0.00726 0.05696  0.03696
#> 9 5000 150  4  4 0.02667 0.01084 0.06643  0.03977
#> 10 5000 150  5  5 0.03333 0.01472 0.07558  0.04225

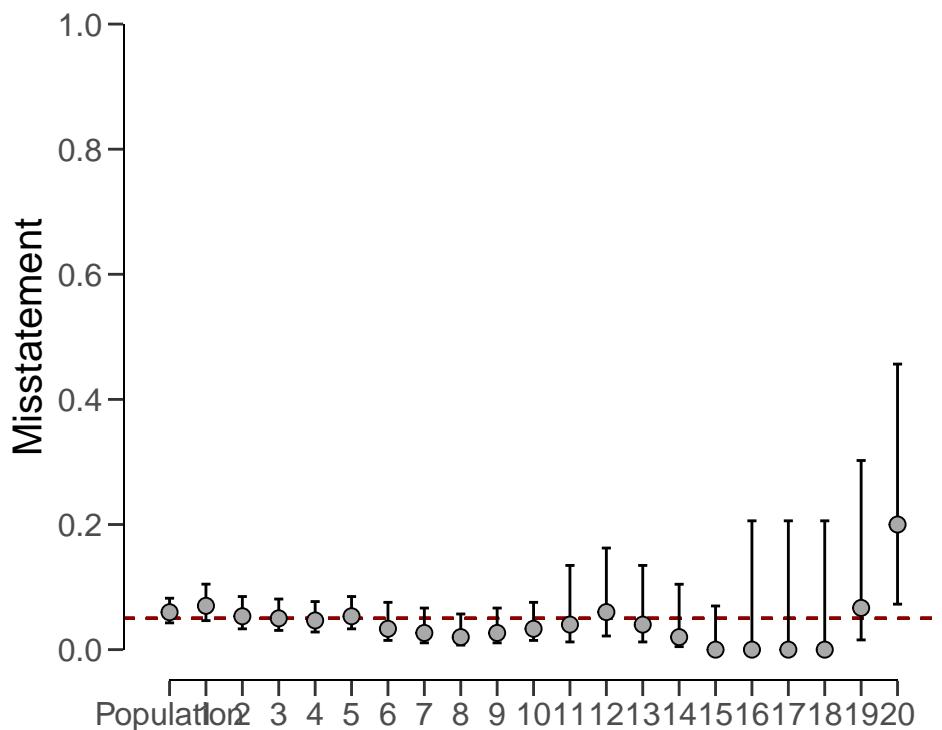
```

6. Stratified Evaluation

```
#> 11 10000 50 2 2 0.04000 0.01230 0.13459 0.09459
#> 12 10000 50 3 3 0.06000 0.02178 0.16242 0.10242
#> 13 10000 50 2 2 0.04000 0.01230 0.13459 0.09459
#> 14 10000 50 1 1 0.02000 0.00478 0.10447 0.08447
#> 15 10000 50 0 0 0.00000 0.00050 0.06978 0.06978
#> 16 10000 15 0 0 0.00000 0.00158 0.20591 0.20591
#> 17 10000 15 0 0 0.00000 0.00158 0.20591 0.20591
#> 18 10000 15 0 0 0.00000 0.00158 0.20591 0.20591
#> 19 10000 15 1 1 0.06667 0.01551 0.30232 0.23565
#> 20 4000 15 3 3 0.20000 0.07266 0.45646 0.25646
```

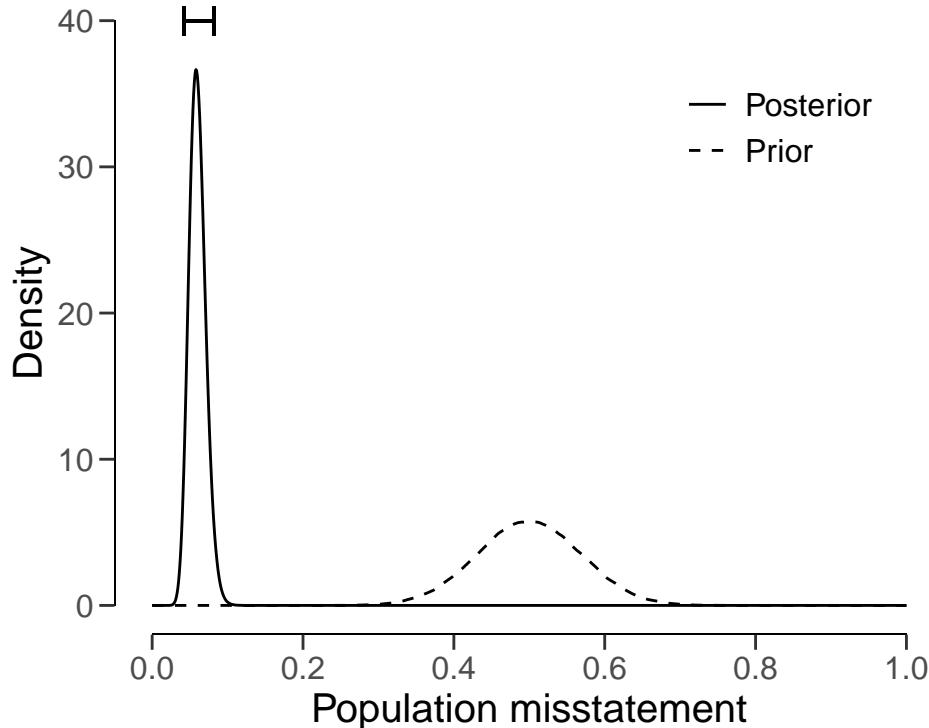
In this case, the output of the `summary()` function shows that the estimate of the misstatement in the population is 5.85%, with the 95% credible interval ranging from 4.28% to 8.22%. The stratum estimates differ substantially from each other but are relatively uncertain.

```
plot(result_np, type = "estimates")
```



The prior and posterior distribution for the population misstatement can be requested via the `plot()` function.

```
plot(result_np, type = "posterior")
```



6.2. Complete pooling

Complete pooling (`pooling = "complete"`) assumes no differences between strata. This has the advantages that data from all strata can be aggregated, which decreases the uncertainty in the population estimate compared to the no pooling approach. However, the disadvantage of this approach is that it does not facilitate the distinction between between strata, as every stratum receives the same estimate equal to that of the population. The call below evaluates the sample using a Bayesian stratified evaluation procedure, in which the strata are assumed to be the same.

```
result_cp <- evaluation(
  materiality = 0.05, method = "binomial", prior = TRUE,
  n = retailer$samples, x = retailer$errors, N.units = retailer$items,
  alternative = "two.sided", pooling = "complete"
)
summary(result_cp)
```

6. Stratified Evaluation

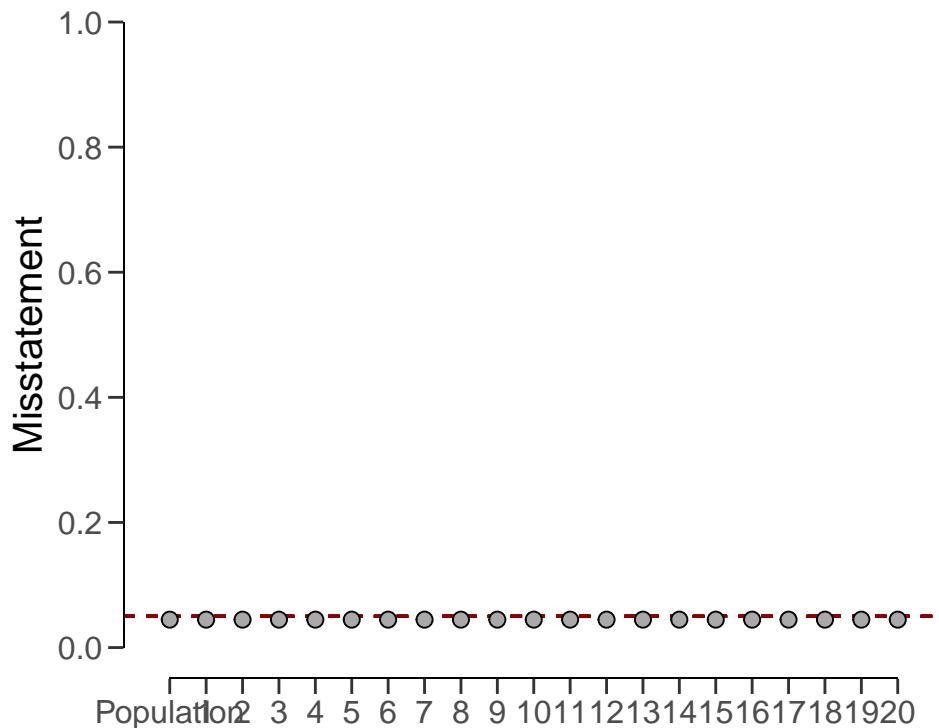
```
#>
#> Bayesian Audit Sample Evaluation Summary
#>
#> Options:
#>   Confidence level:          0.95
#>   Population size:          144000
#>   Materiality:              0.05
#>   Hypotheses:                H : θ = 0.05 vs. H : θ < 0.05
#>   Method:                   binomial
#>   Prior distribution:       beta( = 1,   = 1)
#>
#> Data:
#>   Sample size:               2575
#>   Number of errors:          115
#>   Sum of taints:             115
#>
#> Results:
#>   Posterior distribution:    beta( = 116,   = 2461)
#>   Most likely error:        0.04466
#>   95 percent credible interval: [0.03735, 0.053345]
#>   Precision:                 0.0086852
#>   BF :                      0.022725
#>
#> Strata (20):
#>      N  n  x  t      mle      lb      ub precision     bf10
#>  1  5000 300 21 21 0.04466 0.03735 0.05335  0.00869 0.02273
#>  2  5000 300 16 16 0.04466 0.03735 0.05335  0.00869 0.02273
#>  3  5000 300 15 15 0.04466 0.03735 0.05335  0.00869 0.02273
#>  4  5000 300 14 14 0.04466 0.03735 0.05335  0.00869 0.02273
#>  5  5000 300 16 16 0.04466 0.03735 0.05335  0.00869 0.02273
#>  6  5000 150  5  5 0.04466 0.03735 0.05335  0.00869 0.02273
#>  7  5000 150  4  4 0.04466 0.03735 0.05335  0.00869 0.02273
#>  8  5000 150  3  3 0.04466 0.03735 0.05335  0.00869 0.02273
#>  9  5000 150  4  4 0.04466 0.03735 0.05335  0.00869 0.02273
#> 10 5000 150  5  5 0.04466 0.03735 0.05335  0.00869 0.02273
#> 11 10000 50  2  2 0.04466 0.03735 0.05335  0.00869 0.02273
#> 12 10000 50  3  3 0.04466 0.03735 0.05335  0.00869 0.02273
#> 13 10000 50  2  2 0.04466 0.03735 0.05335  0.00869 0.02273
#> 14 10000 50  1  1 0.04466 0.03735 0.05335  0.00869 0.02273
#> 15 10000 50  0  0 0.04466 0.03735 0.05335  0.00869 0.02273
#> 16 10000 15  0  0 0.04466 0.03735 0.05335  0.00869 0.02273
#> 17 10000 15  0  0 0.04466 0.03735 0.05335  0.00869 0.02273
#> 18 10000 15  0  0 0.04466 0.03735 0.05335  0.00869 0.02273
```

6.2. Complete pooling

```
#> 19 10000 15 1 1 0.04466 0.03735 0.05335 0.00869 0.02273
#> 20 4000 15 3 3 0.04466 0.03735 0.05335 0.00869 0.02273
```

For example, the output of the `summary()` function shows that the estimate of the misstatement in the population is 4.47%, with the 95% credible interval ranging from 3.74% to 5.33%. Since the data is aggregated, the stratum estimates contain relatively little uncertainty. However, the probability of misstatement in stratum 20 (many misstatements) under this assumption is the same as that of stratum 15 (few misstatements).

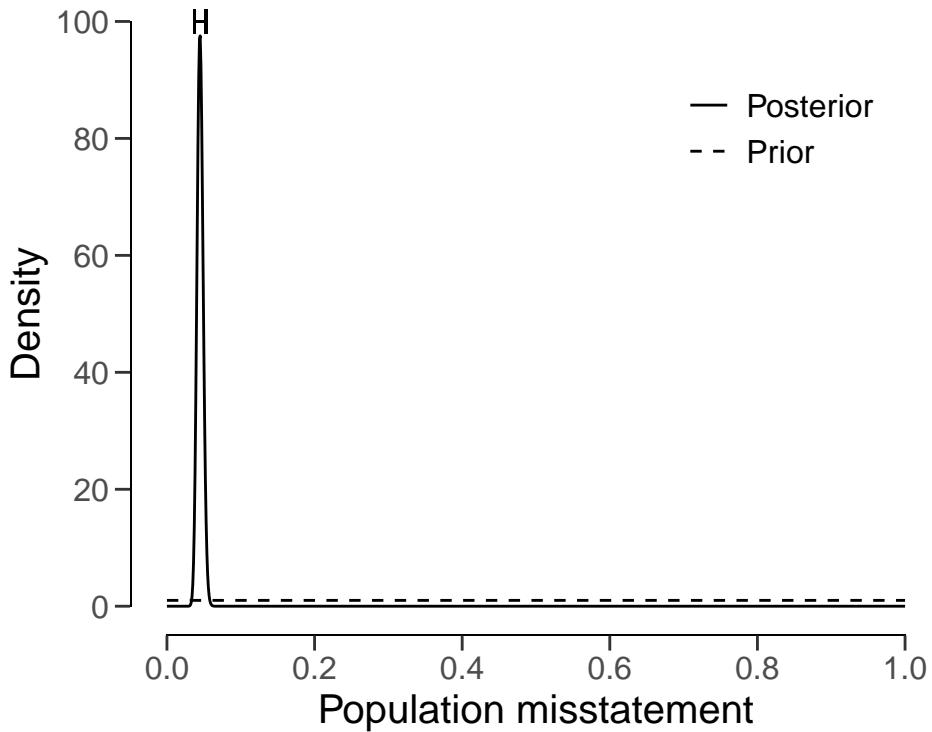
```
plot(result_cp, type = "estimates")
```



The prior and posterior distribution for the population misstatement can be requested via the `plot()` function.

```
plot(result_cp, type = "posterior")
```

6. Stratified Evaluation



6.3. Partial pooling

Finally, partial pooling (`pooling = "partial"`) assumes differences and similarities between strata. This allows the auditor to differentiate between strata, while also sharing information between the strata to reduce uncertainty in the population estimate.

The call below evaluates the sample using a Bayesian stratified evaluation procedure, in which the stratum estimates are poststratified to arrive at the population estimate.

```
set.seed(1) # Important because the posterior distribution is determined via sample
result_pp <- evaluation(
  materiality = 0.05, method = "binomial", prior = TRUE,
  n = retailer$samples, x = retailer$errors, N.units = retailer$items,
  alternative = "two.sided", pooling = "partial"
)
summary(result_pp)
#>
#> Bayesian Audit Sample Evaluation Summary
#>
```



Figure 6.2.: Partial pooling takes into account the hierarchical structure in the data.
Image available under a CC-BY-NC 4.0 license.

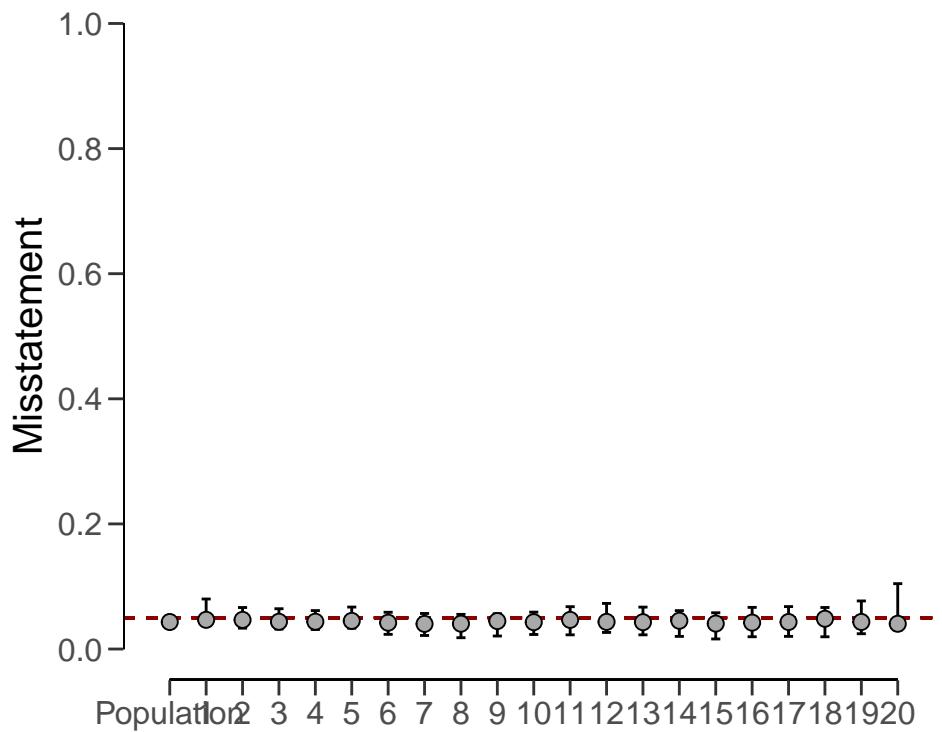
```
#> Options:
#>   Confidence level:          0.95
#>   Population size:          144000
#>   Materiality:              0.05
#>   Hypotheses:                H : θ = 0.05 vs. H : θ < 0.05
#>   Method:                   binomial
#>   Prior distribution:       beta( = 1,   = 1)
#>
#> Data:
#>   Sample size:               2575
#>   Number of errors:          115
#>   Sum of taints:             115
#>
#> Results:
#>   Posterior distribution:    beta( = 116,   = 2461)
#>   Most likely error:         0.0432
#>   95 percent credible interval: [0.034575, 0.052994]
#>   Precision:                 0.0097941
#>   BF :                       0.022725
#>
#> Strata (20):
#>   N   n   x   t   mle      lb      ub precision
#> 1  5000 300 21 21 0.0467 0.03778 0.07996  0.03326
```

6. Stratified Evaluation

```
#> 2 5000 300 16 16 0.0466 0.03335 0.06634 0.01974
#> 3 5000 300 15 15 0.0436 0.03184 0.06439 0.02079
#> 4 5000 300 14 14 0.0437 0.03121 0.06147 0.01777
#> 5 5000 300 16 16 0.0450 0.03328 0.06707 0.02207
#> 6 5000 150 5 5 0.0421 0.02346 0.05876 0.01666
#> 7 5000 150 4 4 0.0402 0.02165 0.05686 0.01666
#> 8 5000 150 3 3 0.0405 0.01807 0.05541 0.01491
#> 9 5000 150 4 4 0.0447 0.02082 0.05651 0.01181
#> 10 5000 150 5 5 0.0429 0.02326 0.05899 0.01609
#> 11 10000 50 2 2 0.0466 0.02271 0.06766 0.02106
#> 12 10000 50 3 3 0.0434 0.02654 0.07291 0.02951
#> 13 10000 50 2 2 0.0429 0.02270 0.06701 0.02411
#> 14 10000 50 1 1 0.0452 0.02022 0.06137 0.01617
#> 15 10000 50 0 0 0.0407 0.01608 0.05795 0.01725
#> 16 10000 15 0 0 0.0422 0.01964 0.06657 0.02437
#> 17 10000 15 0 0 0.0432 0.02023 0.06786 0.02466
#> 18 10000 15 0 0 0.0484 0.01956 0.06639 0.01799
#> 19 10000 15 1 1 0.0433 0.02447 0.07685 0.03355
#> 20 4000 15 3 3 0.0404 0.03031 0.10450 0.06410
```

In this case, the output shows that the estimate of the misstatement in the population is 4.34%, with the 95% credible interval ranging from 3.45% to 5.33%. Note that this population estimate is substantially less uncertain than that of the no pooling approach. Note that, like in the no pooling approach, the stratum estimates are different from each other but lie closer together and are less uncertain.

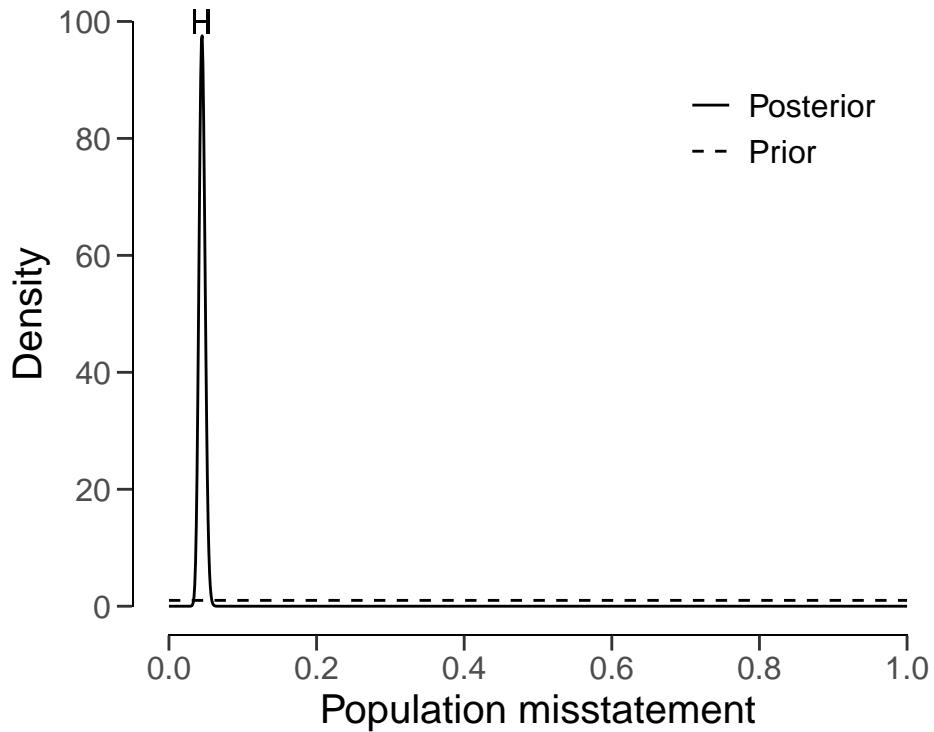
```
plot(result_pp, type = "estimates")
```



The prior and posterior distribution for the population misstatement can be requested via the `plot()` function.

```
plot(result_pp, type = "posterior")
```

6. Stratified Evaluation



6.4. Evaluation using data

For this example, we take the `allowances` data set that comes with the package. This data set contains 3500 financial statement line items, each with a booked value `bookValue` and, for illustrative purposes, and audited (true) value `auditValue`. Since the focus of this vignette is the evaluation stage in the audit, the sample is already indicated in the data set. The performance materiality in this example is set to 5%.

```
data(allowances)
head(allowances)
#>   item branch bookValue auditValue times
#> 1    1     12      1600      1600     1
#> 2    2     12      1625       NA     0
#> 3    3     12      1775       NA     0
#> 4    4     12      1250      1250     1
#> 5    5     12      1400       NA     0
#> 6    6     12      1190       NA     0
```

Evaluating a stratified sample using data requires specification of the `data`, `values`, `values.audit` and `strata` arguments in the `evaluation()` function. In this case, the

units are monetary and calculated by aggregating the book values of the items in each stratum.

```
N.units <- aggregate(allowances$bookValue, list(allowances$branch), sum)$x
```

6.4.1. Classical Evaluation

The call below evaluates the `allowances` sample using a classical stratified evaluation procedure, in which the stratum estimates are poststratified to arrive at the population estimate.

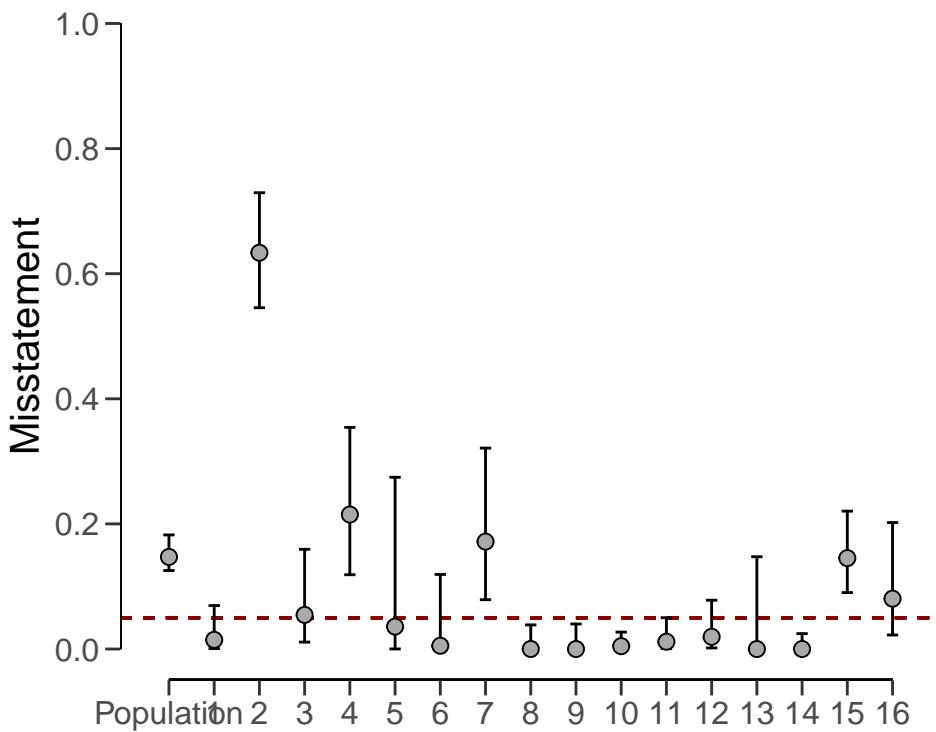
```
x <- evaluation(
  materiality = 0.05, data = allowances,
  values = "bookValue", values.audit = "auditValue", strata = "branch", times = "times",
  alternative = "two.sided", N.units = N.units
)
summary(x)
#>
#> Classical Audit Sample Evaluation Summary
#>
#> Options:
#>   Confidence level:          0.95
#>   Population size:          16772249
#>   Materiality:              0.05
#>   Hypotheses:                H : θ = 0.05 vs. H : θ ≠ 0.05
#>   Method:                   poisson
#>
#> Data:
#>   Sample size:                 1604
#>   Number of errors:            401
#>   Sum of taints:               252.9281046
#>
#> Results:
#>   Most likely error:           0.14723
#>   95 percent confidence interval: [0.12549, 0.18239]
#>   Precision:                  0.03516
#>   p-value:                     NA
#>
#> Strata (16):
#>   N   n   x       t     mle      lb      ub precision p.value
#> 1 317200.09 87    6  1.27814 0.01469 0.00073 0.06950  0.05481 0.46285
#> 2 2792814.33 305  233 193.23313 0.63355 0.54558 0.72945  0.09590 0.00000
```

6. Stratified Evaluation

```
#> 3 1144231.69 55 3 3.00000 0.05455 0.01105 0.15940 0.10486 0.75827
#> 4 414202.89 70 45 15.05094 0.21501 0.11878 0.35434 0.13933 0.00000
#> 5 96660.53 18 1 0.64537 0.03585 0.00015 0.27456 0.23871 0.59343
#> 6 348006.13 34 1 0.17866 0.00525 0.00000 0.11926 0.11401 1.00000
#> 7 2384079.33 55 14 9.44448 0.17172 0.07885 0.32122 0.14950 0.00058
#> 8 1840399.33 96 1 0.00813 0.00008 0.00000 0.03860 0.03852 0.10355
#> 9 563957.70 92 0 0.00000 0.00000 0.00000 0.04010 0.04010 0.01783
#> 10 3198877.73 201 7 0.92023 0.00458 0.00009 0.02703 0.02245 0.00122
#> 11 1983299.06 128 7 1.50034 0.01172 0.00084 0.05013 0.03841 0.10773
#> 12 319144.13 86 5 1.68141 0.01955 0.00174 0.07806 0.05851 0.46069
#> 13 148905.79 25 0 0.00000 0.00000 0.00000 0.14756 0.14756 0.64187
#> 14 513058.76 150 0 0.00000 0.00000 0.00000 0.02459 0.02459 0.00134
#> 15 432007.61 150 39 21.80000 0.14533 0.09026 0.22045 0.07511 0.00001
#> 16 275403.70 52 39 4.18726 0.08052 0.02237 0.20215 0.12163 0.12258
```

In this case, the output shows that the estimate of the misstatement in the population is 14.72%, with the 95% confidence interval ranging from 12.55% to 18.26%. The precision of the population estimate is 3.54%. The stratum estimates can be seen in the output of the `summary()` function and are visualized below.

```
plot(x, type = "estimates")
```



6.4.2. Bayesian Evaluation

Bayesian inference can improve upon the estimates of the classical approach by pooling information between strata where possible. The call below evaluates the `allowances` sample using a Bayesian multilevel stratified evaluation procedure, in which the stratum estimates are poststratified to arrive at the population estimate.

```
x <- evaluation(
  materiality = 0.05, data = allowances, prior = TRUE,
  values = "bookValue", values.audit = "auditValue", strata = "branch", times = "times",
  alternative = "two.sided", N.units = N.units, pooling = "partial"
)
summary(x)
#>
#>   Bayesian Audit Sample Evaluation Summary
#>
#> Options:
#>   Confidence level:          0.95
#>   Population size:           16772249
#>   Materiality:              0.05
```

6. Stratified Evaluation

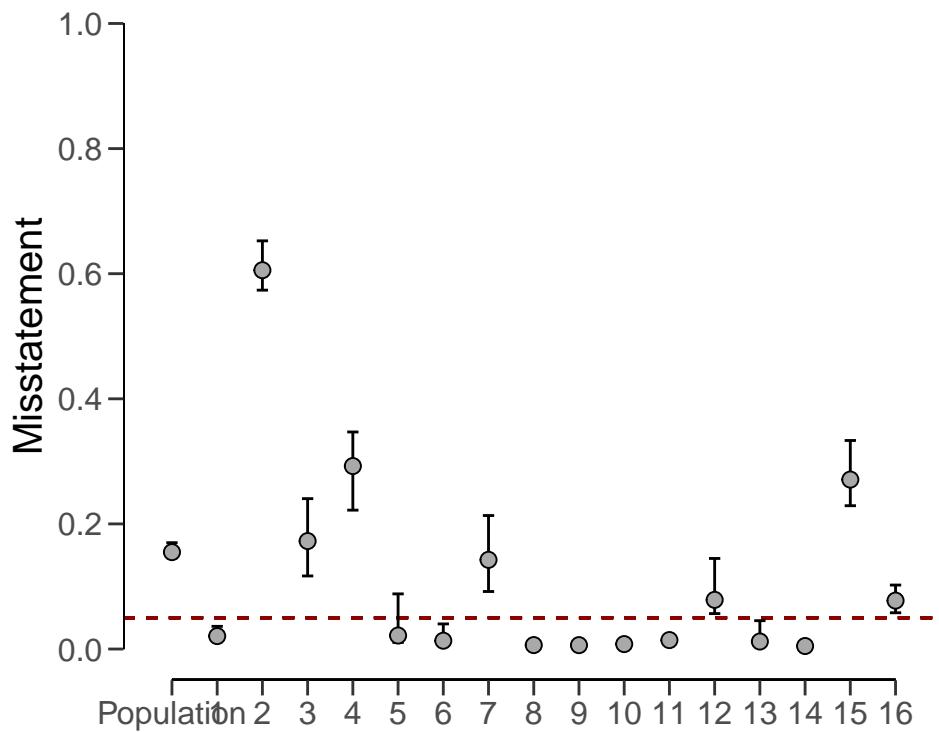
```

#>   Hypotheses:          H : θ = 0.05 vs. H : θ < 0.05
#>   Method:             poisson
#>   Prior distribution: gamma(α = 1, β = 1)
#>
#> Data:
#>   Sample size:         1350
#>   Number of errors:   401
#>   Sum of taints:      224.657517
#>
#> Results:
#>   Posterior distribution: gamma(α = 225.658, β = 1351)
#>   Most likely error:    0.1549
#>   95 percent credible interval: [0.14616, 0.17]
#>   Precision:            0.015096
#>   BF :                  2.751e+47
#>
#> Strata (16):
#>   N   n   x       t   mle     lb     ub precision
#> 1 317200.09 62   6   1.07814 0.0207 0.01365 0.03625  0.01555
#> 2 2792814.33 283 233 176.87807 0.6056 0.57370 0.65251  0.04691
#> 3 1144231.69 55   3   3.00000 0.1726 0.11680 0.24039  0.06779
#> 4 414202.89 48   45  11.65094 0.2925 0.22203 0.34698  0.05448
#> 5 96660.53   9   1   0.21512 0.0216 0.01058 0.08814  0.06654
#> 6 348006.13 14   1   0.04467 0.0133 0.00707 0.04013  0.02683
#> 7 2384079.33 44   14  6.65401 0.1426 0.09194 0.21343  0.07083
#> 8 1840399.33 78   1   0.00813 0.0063 0.00467 0.00939  0.00309
#> 9 563957.70  74   0   0.00000 0.0062 0.00460 0.00972  0.00352
#> 10 3198877.73 184  7   0.89982 0.0079 0.00618 0.01055  0.00265
#> 11 1983299.06 111  7   1.35066 0.0143 0.01100 0.02263  0.00833
#> 12 319144.13  56   5   1.19069 0.0786 0.05665 0.14477  0.06617
#> 13 148905.79  10   0   0.00000 0.0120 0.00624 0.04535  0.03335
#> 14 513058.76 150  0   0.00000 0.0048 0.00381 0.00635  0.00155
#> 15 432007.61 128  39  18.30000 0.2709 0.22911 0.33339  0.06249
#> 16 275403.70  44  39  3.38726 0.0774 0.05796 0.10226  0.02486

```

The output shows that the estimate of the misstatement in the population is 15.66%, with the 95% credible interval ranging from 14.59% to 17%. The precision of the population estimate is 1.34%, which is substantially lower than that of the classical approach. The stratum estimates can be seen in the output of the `summary()` function and are visualized below.

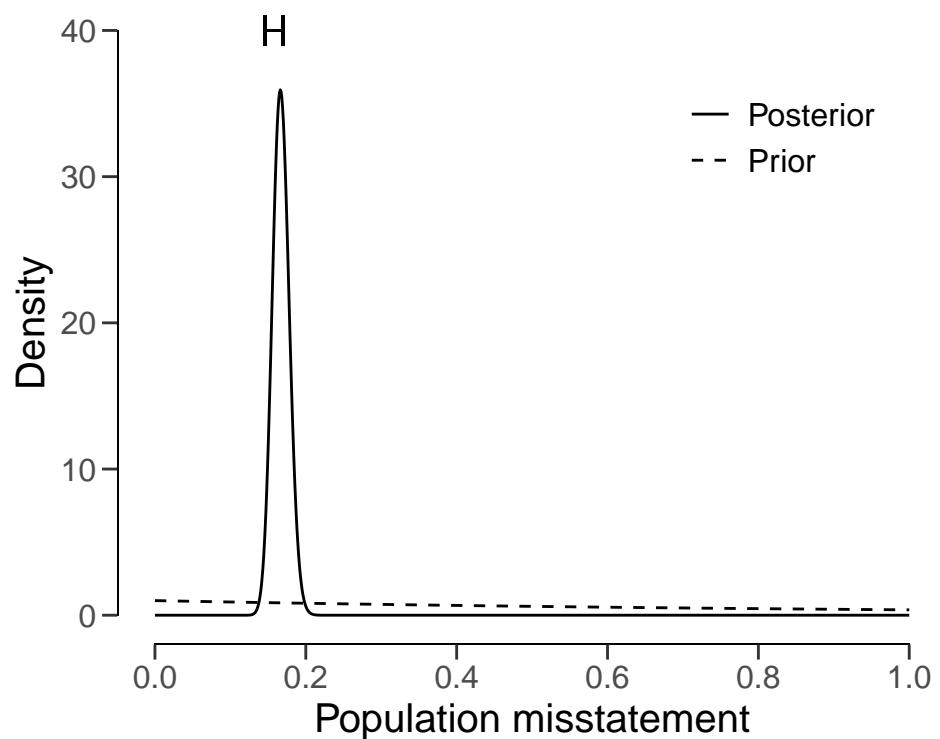
```
plot(x, type = "estimates")
```



The prior and posterior distribution for the population misstatement can be requested via the `plot()` function.

```
plot(x, type = "posterior")
```

6. Stratified Evaluation



Part III.

Other Software

In addition to using the **jfa** R package for statistical audit sampling, there are several other software options that can be incorporated into the workflow. These options may be preferred for specific tasks or may simply provide an alternative approach to achieving the same results.

In this part of the book, we will explore some of these other software options and how they can be integrated into the audit sampling workflow. Chapter 7 discusses JASP for Audit, an open-source and user-friendly graphical user interface around the **jfa** R package. In Chapter 8, we will discuss other R packages implementing statistical techniques for audit sampling.

7. JASP for Audit

i You are reading the work-in-progress first edition of Statistical Auditing with R. This chapter is currently a dumping ground of ideas, and it is incomplete.

JASP for Audit (Derks, Swart, Wagenmakers, et al., 2021) is an add-on module for the open-source statistical software JASP (JASP Team, 2022), based on the **jfa** package, that facilitates statistical audit sampling. The module provides graphical a user interface (GUI) for calculating sample sizes, selecting items according to standard audit sampling techniques, and performing inference about the population misstatement on the basis of a data sample or summary statistics of a sample. The module also features Bayesian equivalents of these analyses that enable the user to easily incorporate prior information into the statistical procedure. In all analyses, the Audit module offers explanatory text that helps the auditor in interpreting, explaining, and reporting the analysis. Since JASP for Audit is an R-based GUI around **jfa**, its functionality can be mapped almost one-on-one to that of the package.

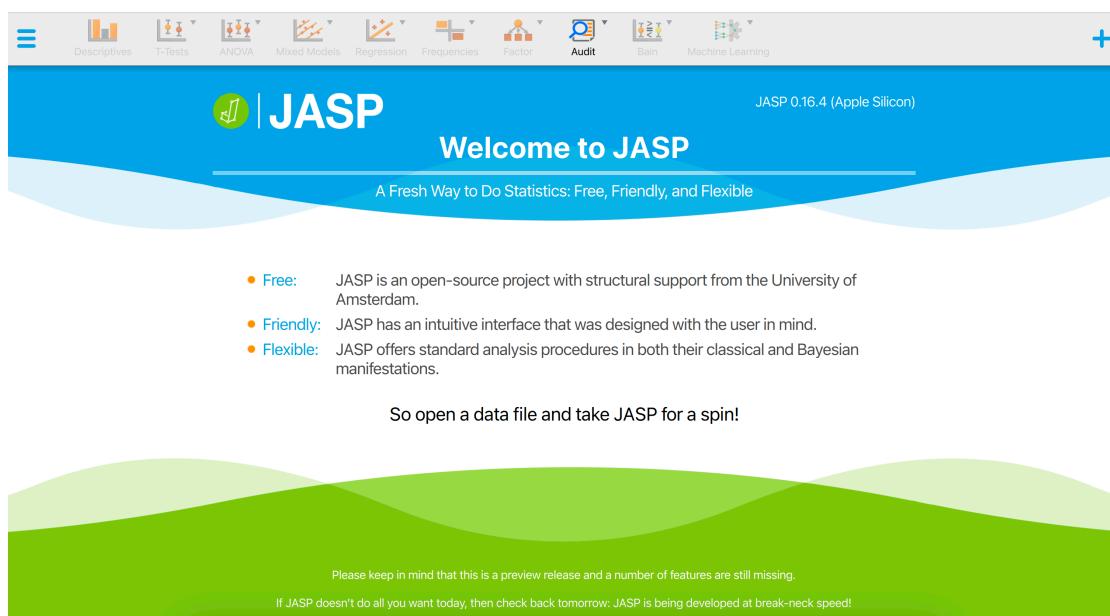


Figure 7.1.: The JASP welcome screen.

7. JASP for Audit

7.1. Planning

In order to plan a sample for an audit using JASP, the process is similar to using **jfa**. For example, the screenshot below shows how to plan a minimum sample size for a performance materiality of 5% using a beta(1, 1) prior distribution, while expecting zero misstatements in the sample.

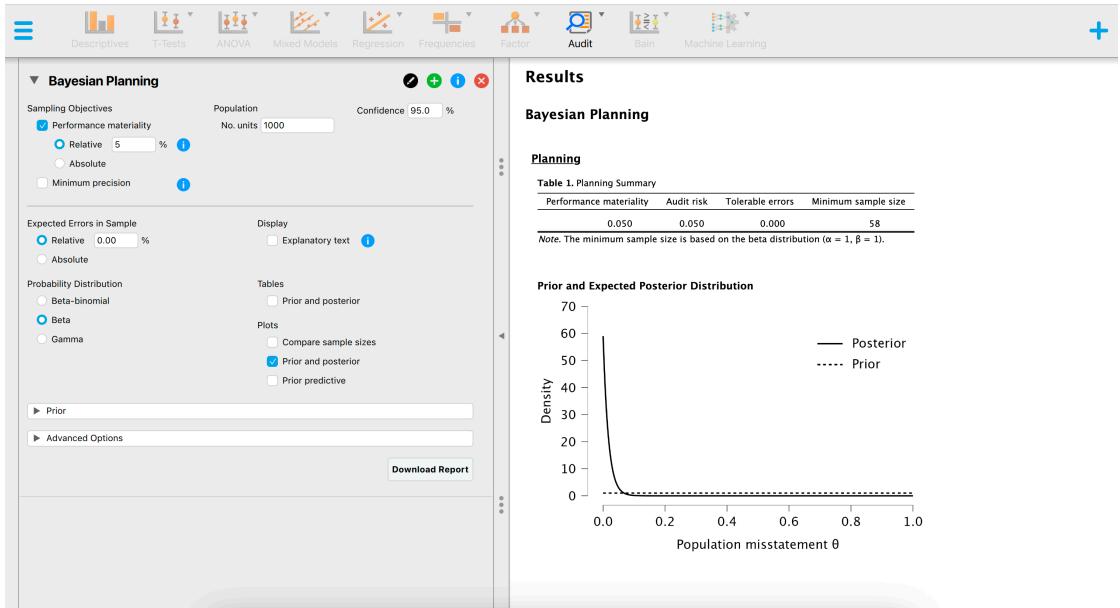


Figure 7.2.: The JASP interface for planning a sample.

In **jfa**, these results can be reproduced via the following command:

```
planning(materiality = 0.05, likelihood = "binomial", prior = TRUE)
#>
#> Bayesian Audit Sample Planning
#>
#> minimum sample size = 58
#> sample size obtained in 59 iterations via method 'binomial' + 'prior'
```

7.2. Selection

Selecting a sample in JASP for Audit works similar to how you would do it in **jfa**. For example, the screenshot below shows how to select a sample of 60 monetary units from the **BuildIt** data set that is included in the package using a fixed interval sampling method with a starting point of 1.

```

set.seed(1)
data(BuildIt)
result <- selection(data = BuildIt, size = 60, units = "values",
                      method = "interval", start = 1, values = "bookValue")
head(result$sample)
#>   row times     ID bookValue auditValue
#> 1    1      1 82884    242.61    242.61
#> 2    63      1 51272    248.40    248.40
#> 3   123      1 37985    562.09    562.09
#> 4   183      1 96080    449.07    449.07
#> 5   240      1 92819    690.08    690.08
#> 6   302      1 94296    198.59    198.59

```

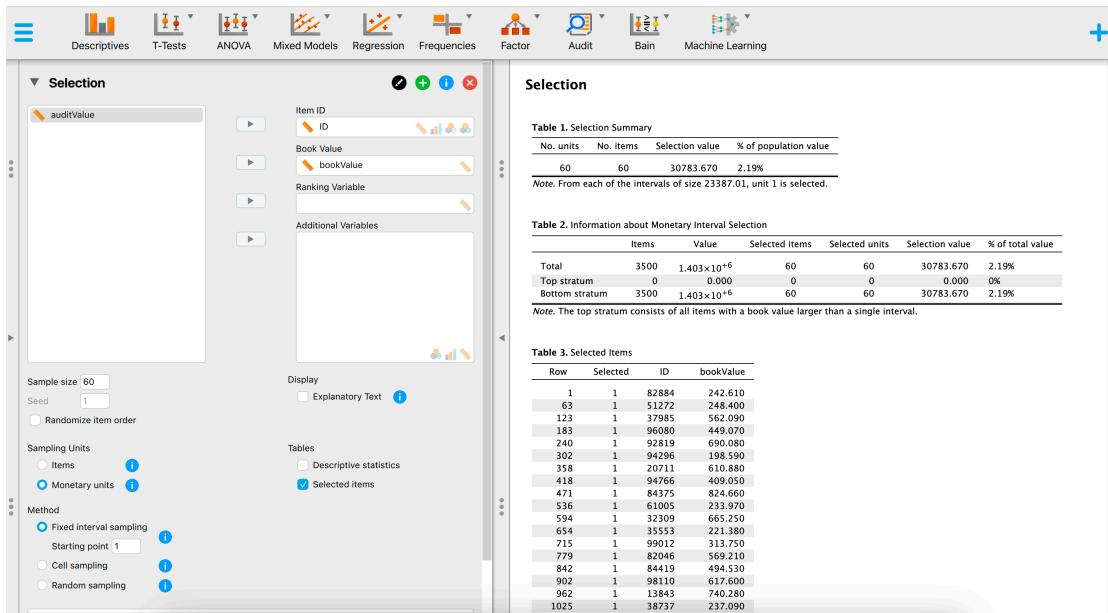


Figure 7.3.: The JASP interface for selecting a sample.

7.3. Evaluation

Evaluating a sample in JASP for Audit works similar to how you would do it in **jfa**. For example, the screenshot below shows how to evaluating a sample of $n = 60$ containing $x = 0$ misstatements against a performance materiality of 5% using a $\text{beta}(1, 1)$ prior distribution.

In **jfa**, these results can be reproduced via the following command:

7. JASP for Audit

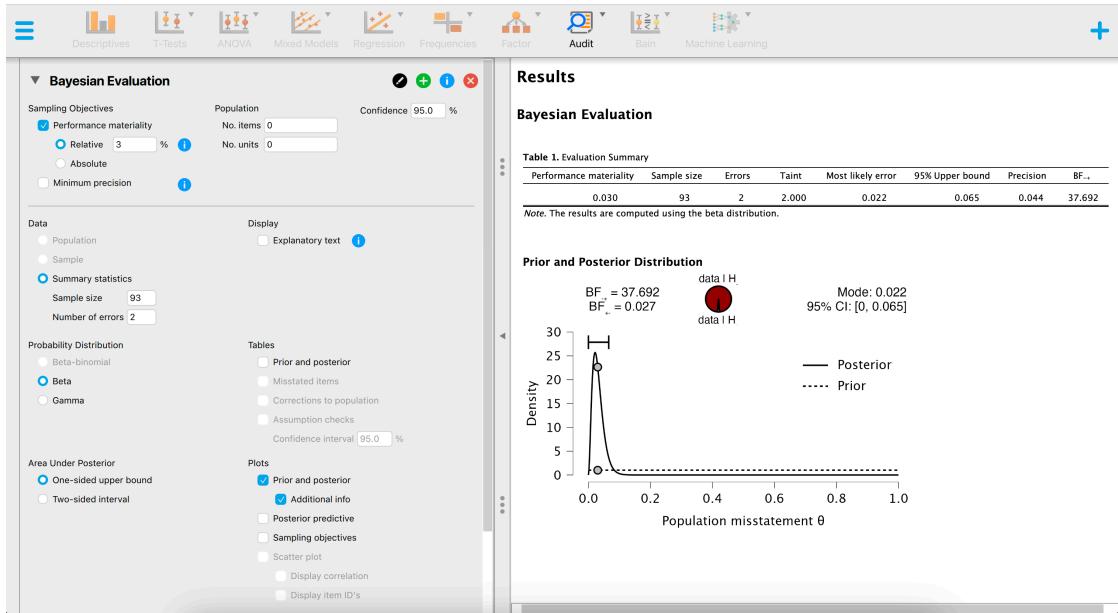


Figure 7.4.: The JASP interface for evaluating a sample.

```
evaluation(materiality = 0.03, method = "binomial", x = 2, n = 93,
           prior = TRUE)
#>
#> Bayesian Audit Sample Evaluation
#>
#> data: 2 and 93
#> number of errors = 2, number of samples = 93, taint = 2, BF  =
#> 37.692
#> alternative hypothesis: true misstatement rate is less than 0.03
#> 95 percent credible interval:
#> 0.0000000 0.0654624
#> most likely estimate:
#> 0.021505
#> results obtained via method 'binomial' + 'prior'
```

8. R Packages

i You are reading the work-in-progress first edition of Statistical Auditing with R. This chapter is currently a dumping ground of ideas, and it is incomplete.

This chapter discusses other R packages implementing statistical techniques for audit sampling.

8.1. MUS

MUS (Prömpers & Guimarães, 2019) is an R package providing sampling and evaluation methods to apply Monetary Unit Sampling during an audit of financial statements. The package is available via CRAN and can be downloaded by running the following code:

```
install.packages("MUS")
```

Note that the **MUS** package provides no functionality for Bayesian audit sampling.

8.2. audit

audit (Meeden, 2021) is an R package based on the paper by Meeden & Sargent (2007) that can be used to find an upper bound for the total amount of overstatement of assets in a set of accounts and estimating the amount of sales tax owed on a collection of transactions. The package is available via CRAN and can be downloaded by running the following code:

```
install.packages("audit")
```

Note that the **audit** package provides limited functionality for evaluating audit samples and has no functionality for planning and selecting audit samples.

8. R Packages

8.3. samplingbook

samplingbook (Manitz et al., 2021) is an R package based on the book **Stichproben: Methoden und praktische Umsetzung mit R** (Kauermann & Kuechenhoff, 2010) that focuses on survey sampling and statistical analysis of these samples. The package is available via CRAN and can be downloaded by running the following code:

```
install.packages("samplingbook")
```

Note that the **samplingbook** package provides limited functionality for auditing as it mainly focuses on survey sampling.

References

- American Institute of Certified Public Accountants (AICPA). (2016a). Appendix A: Attributes statistical sampling tables. In *Audit guide: Audit sampling*. <https://doi.org/10.1002/9781119448617.app1>
- American Institute of Certified Public Accountants (AICPA). (2016b). Appendix C: Monetary unit sampling tables. In *Audit guide: Audit sampling*. <https://doi.org/10.1002/9781119448617.app3>
- Chang, W. (2022). *R Graphics Cookbook*. O'Reilly. <https://r-graphics.org/>
- Derks, K., Swart, J. de, Batenburg, P. van, Wagenmakers, E.-J., & Wetzel, R. (2021). Priors in a Bayesian audit: How integration of existing information into the prior distribution can improve audit transparency and efficiency. *International Journal of Auditing*, 25(3), 621–636. <https://doi.org/10.1111/ijau.12240>
- Derks, K., Swart, J. de, Wagenmakers, E., & Wetzel, R. (2022). The Bayesian approach to audit evidence: Quantifying statistical evidence using the Bayes factor. *PsyArXiv*. <https://doi.org/10.31234/osf.io/kzqp5>
- Derks, K., Swart, J. de, Wagenmakers, E., Wille, J., & Wetzel, R. (2021). JASP for Audit: Bayesian tools for the auditing practice. *Journal of Open Source Software*, 6(68), 2733.
- Dyer, D., & Pierce, R. L. (1993). On the choice of the prior distribution in hypergeometric sampling. *Communications in Statistics - Theory and Methods*, 22(8), 2125–2146. <https://doi.org/10.1080/03610929308831139>
- Grolemund, G. (2014). *Hands-On Programming with R*. <https://rstudio-education.github.io/hopr/>
- Hastie, T., Tibshirani, R., & Friedman, J. H. (2009). *The elements of statistical learning: Data mining, inference, and prediction* (Vol. 2). Springer.
- International Auditing and Assurance Standards Board (IAASB). (2018). ISA 530: Audit sampling. In *International standards on auditing (ISA)*.
- JASP Team. (2022). *JASP (Version 0.16.4)[Computer software]*. <https://jasp-stats.org/>
- Kauermann, G., & Kuechenhoff, H. (2010). *Stichproben: Methoden und praktische umsetzung mit r*. Springer-Verlag.
- Lin, J. (2021). *Audit Analytics with R*. <https://auditanalytics.jonlin.ca/>
- Manitz, J., Hempelmann, M., Kauermann, G., Kuechenhoff, H., Shao, S., Oberhauser, C., Westerheide, N., & Wiesenfarth, M. (2021). *Samplingbook: Survey sampling procedures*. <https://CRAN.R-project.org/package=samplingbook>
- Meeden, G. (2021). *Audit: Bounds for accounting populations*. <https://CRAN.R-project.org/package=audit>

References

- Meeden, G., & Sargent, D. (2007). Some Bayesian methods for two auditing problems. *Communications in Statistics — Theory and Methods*, 36(15), 2727–2740. <https://doi.org/10.1080/03610920701386802>
- Prömpers, H., & Guimarães, A. (2019). *MUS: Monetary unit sampling and estimation methods, widely used in auditing*. <https://CRAN.R-project.org/package=MUS>
- R Core Team. (2022). *R: A language and environment for statistical computing*. R Foundation for Statistical Computing. <https://www.R-project.org/>
- Stewart, T. (2012). *Technical notes on the AICPA audit guide audit sampling*. 5–8.
- Wickam, H., & Brian, J. (2022). *R Packages*. <https://r-pkgs.org/>
- Wickam, H., & Grolemund, G. (2017). *R for Data Science*. <https://r4ds.had.co.nz/>
- Wilke, C. O. (2022). *Fundamentals of Data Visualization*. <https://clauswilke.com/dataviz/>